



中国科学技术大学

University of Science and Technology of China

图像颜色变换

Color Transformation of Images

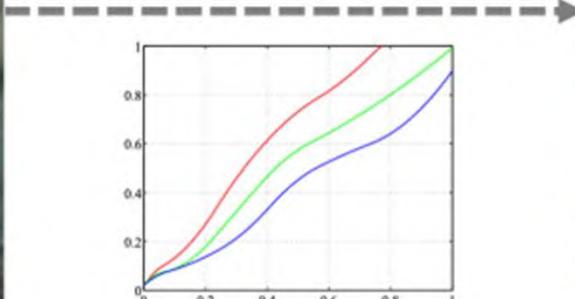
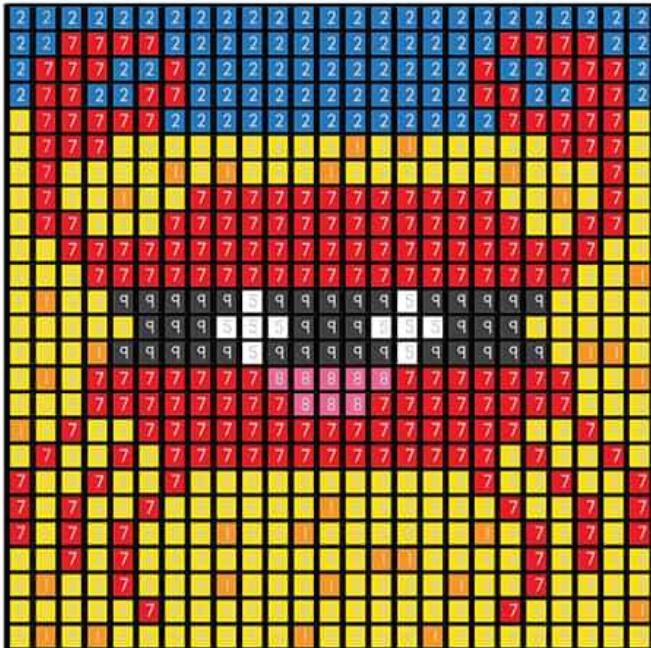


Image as function of Coordinates



$$f(x, y) = c, R^2 \rightarrow R^d, d \in \{1, 3\}$$

$$x \in [0, w), y \in [0, h)$$

Pixel Perfect Pictures

Summer I

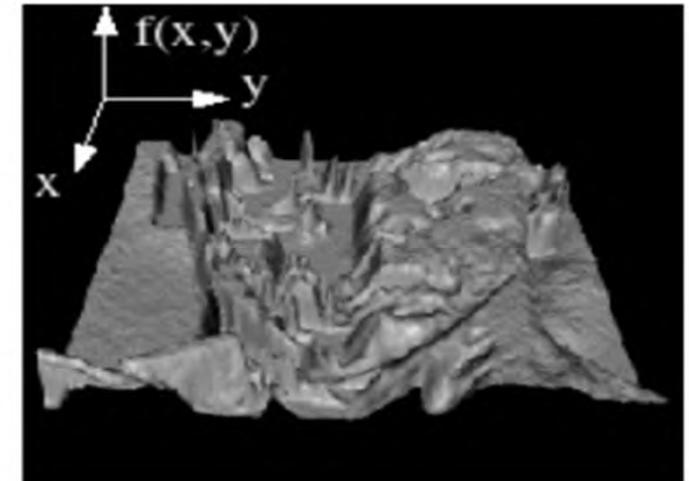
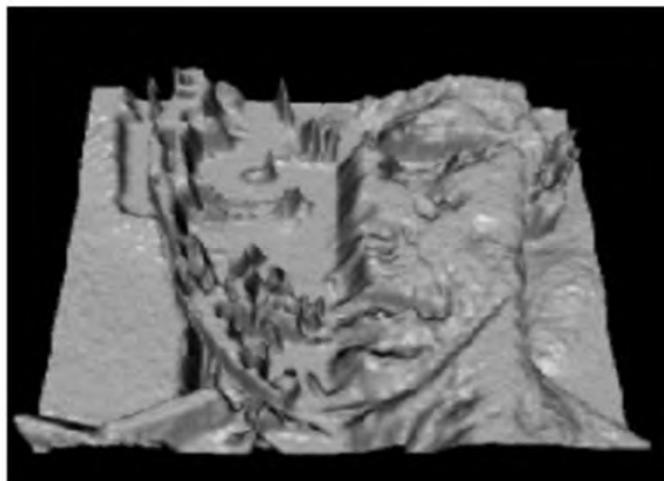
Use the code below to color the grid and reveal the picture.

1 = Orange 3 = Peach 5 = White 7 = Red 9 = Black
2 = Blue 4 = Brown 6 = Yellow 8 = Pink



Images ©2012 SuperStarWorksheets.com

Image as function of Coordinates



Visualization of single channel grayscale image $f(x, y)$

Color Transform v.s. Geometry Transform

image warping: change **domain** of image

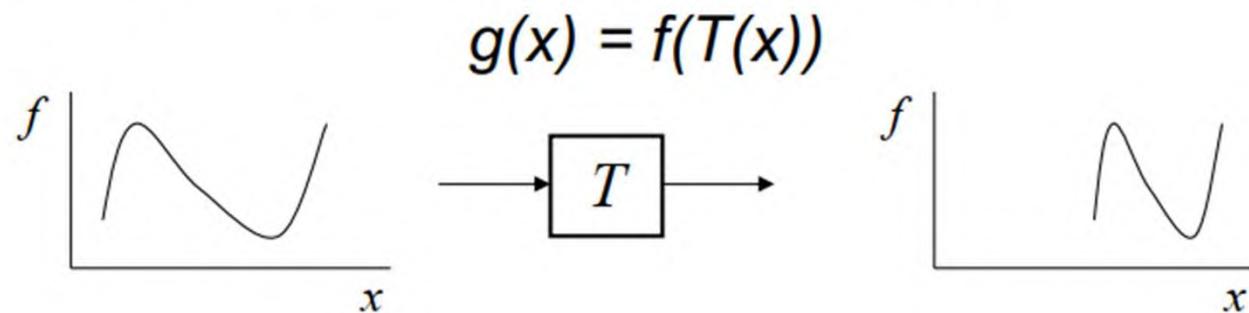
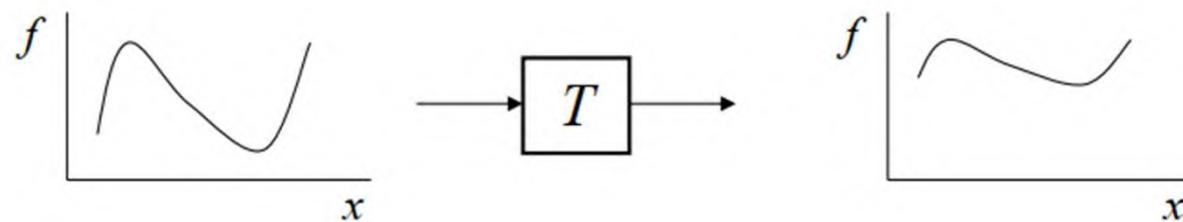


image filtering: change **range** of image

$$g(x) = T(f(x))$$



Color Transform v.s. Geometry Transform

image warping: change **domain** of image

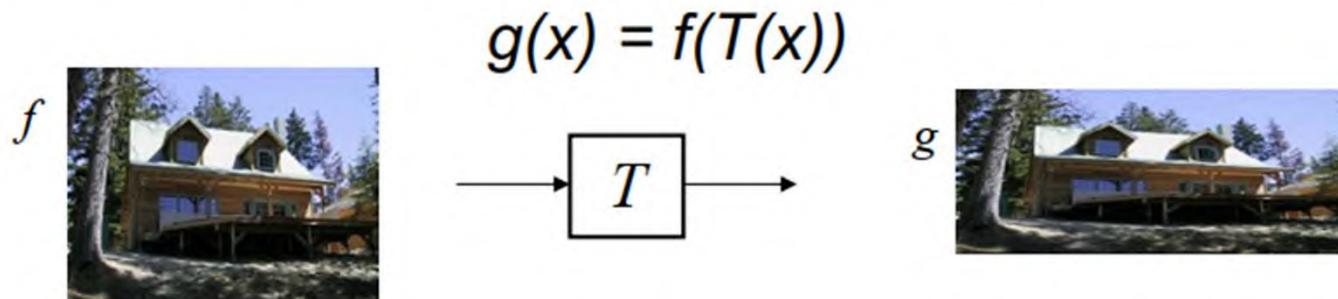
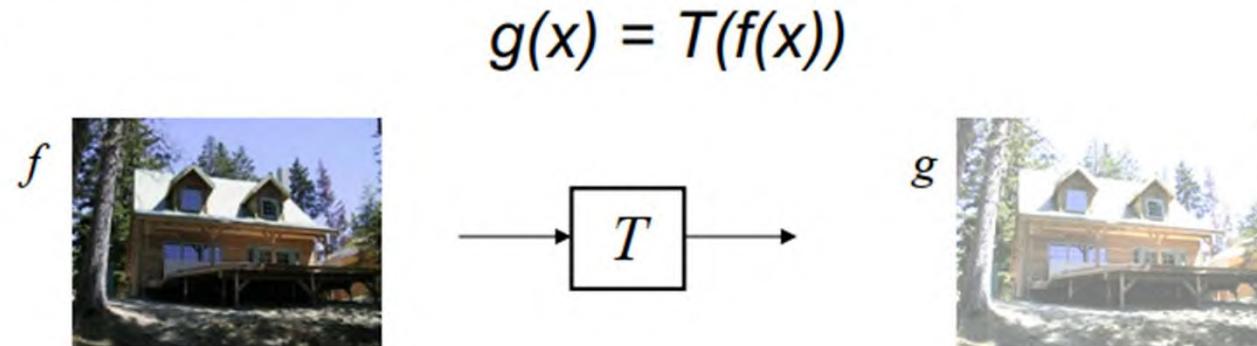


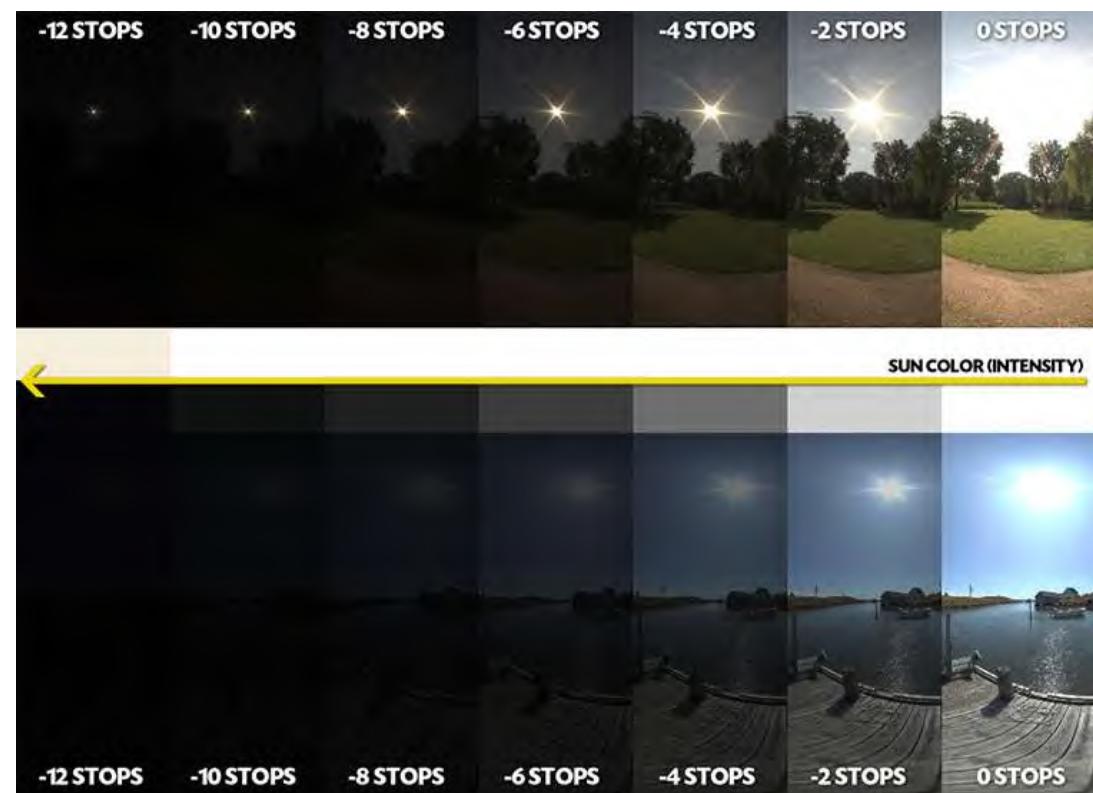
image filtering: change **range** of image



An example: HDR Image

What Is Dynamic Range in Photography?

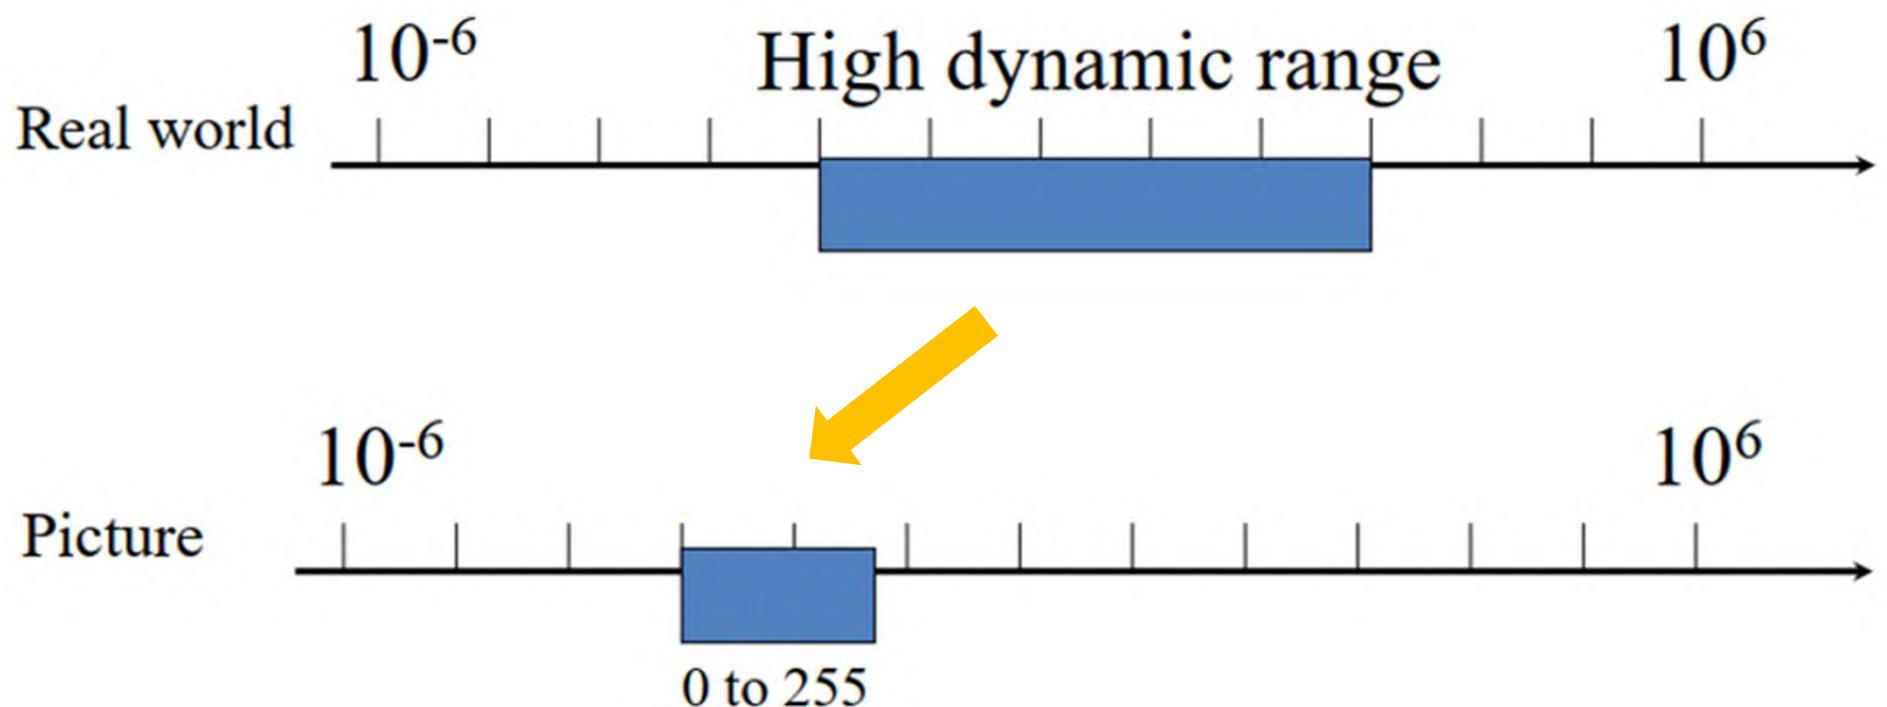
In photography, dynamic range is the contrast ratio between the darkest and brightest color tones that a camera can capture in a single exposure. Maximum dynamic range is the greatest range of light a digital camera sensor or strip of film can capture. Dynamic range is measured in stops. Each stop indicates a doubling of the level of brightness captured. While the human eye can see up to 20 stops of dynamic range, even high-end DSLR and mirrorless cameras can only reach around 14 stops.



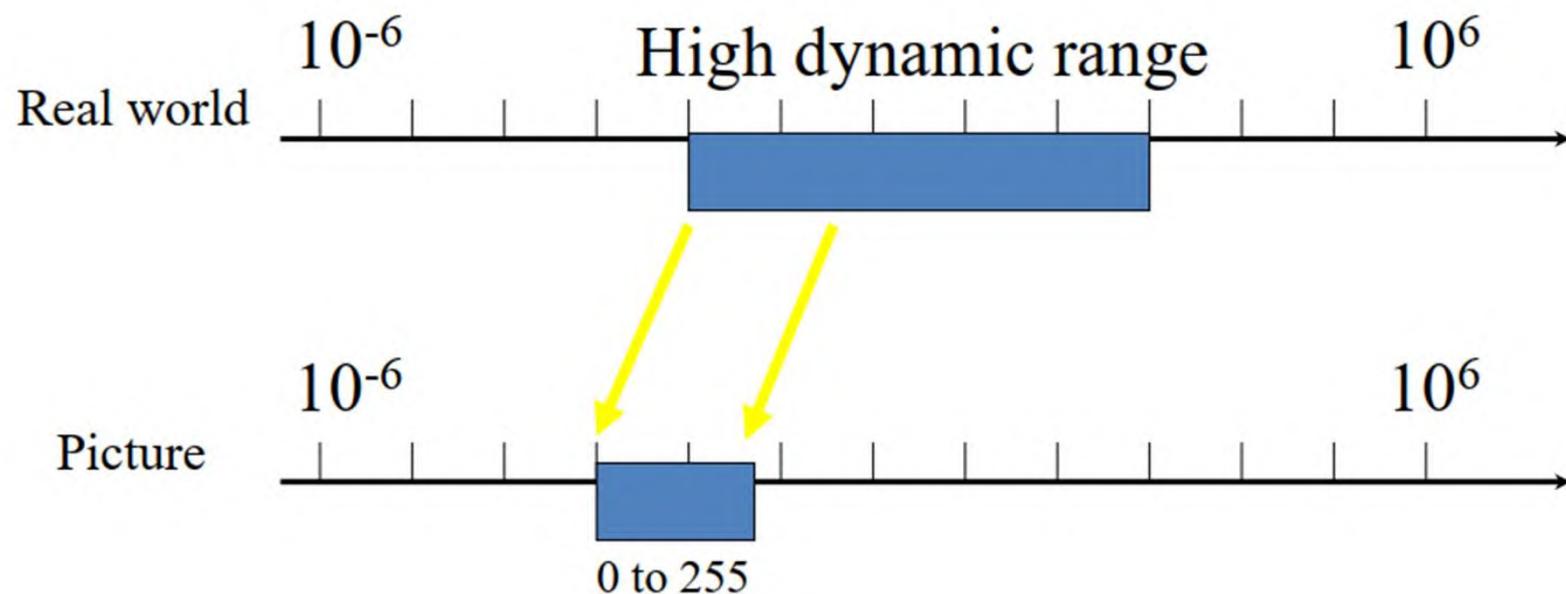
An example: HDR Image



An example: HDR Image

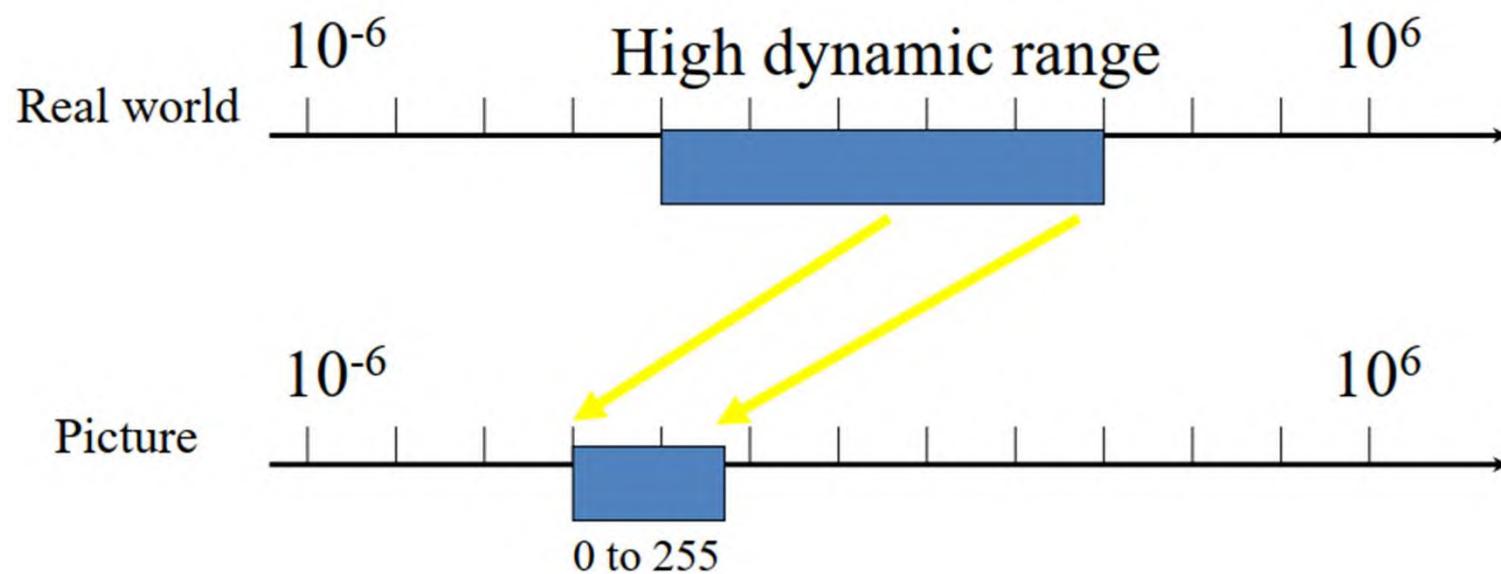


Capture Dark Light



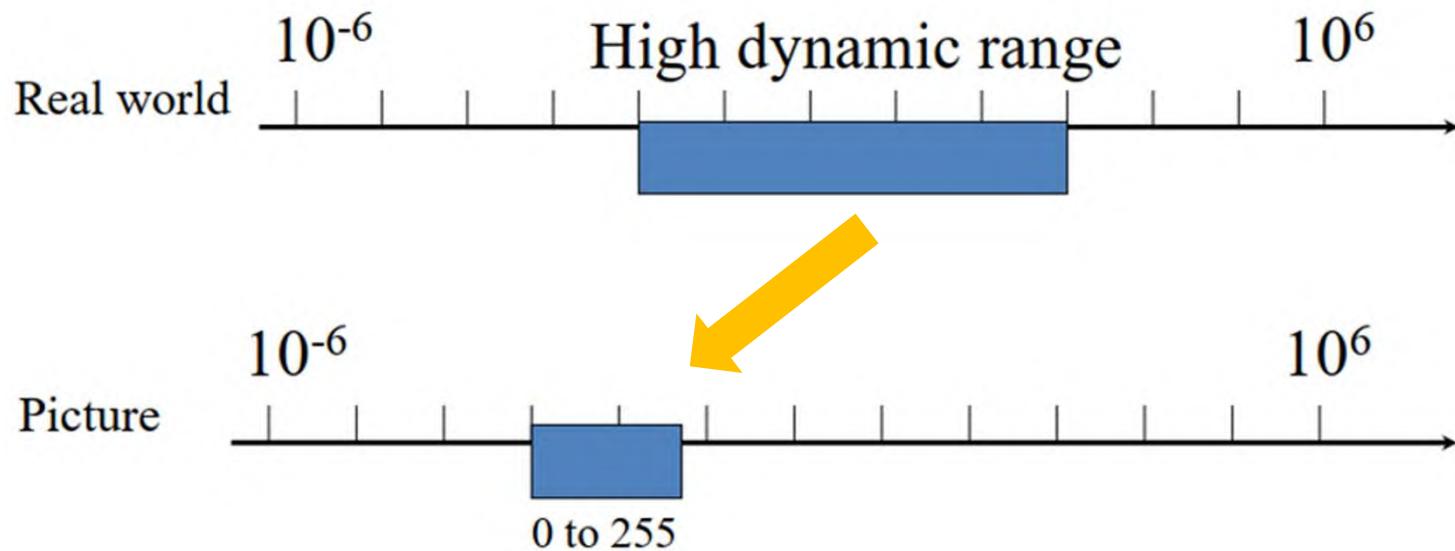
长时间曝光 Long Exposure

Capture Bright Light



短时间曝光 Short Exposure

An example: HDR Image



Basic Transform: Enhancement

a
b
c
d

FIGURE 3.9
(a) Aerial image.
(b)–(d) Results of applying the transformation in Eq. (3.2-3) with $c = 1$ and $\gamma = 3.0, 4.0$, and 5.0 , respectively. (Original image for this example courtesy of NASA.)

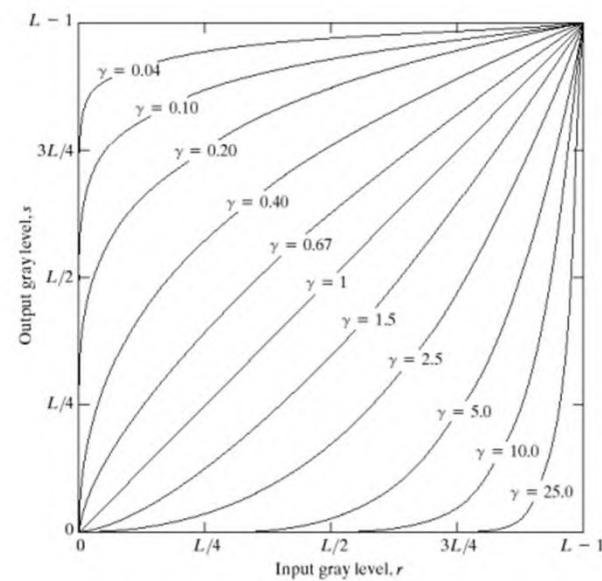
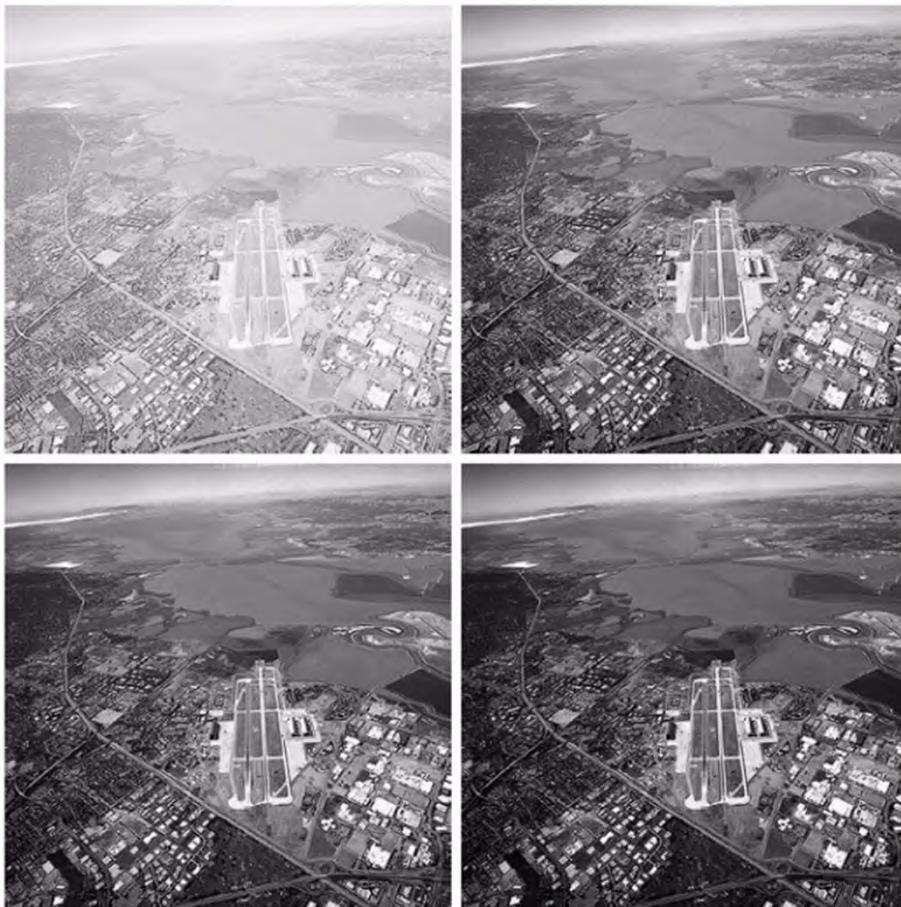


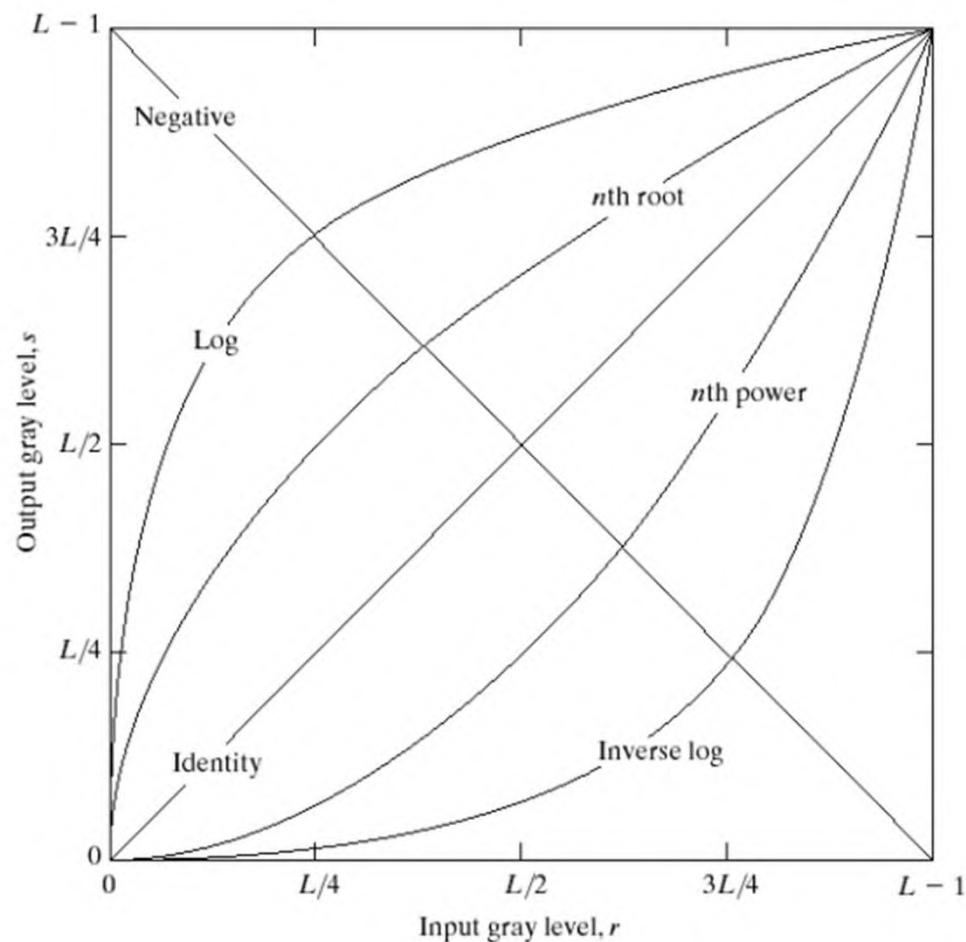
FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

$$c \cdot f(x, y)^\gamma$$

Power-law transformations

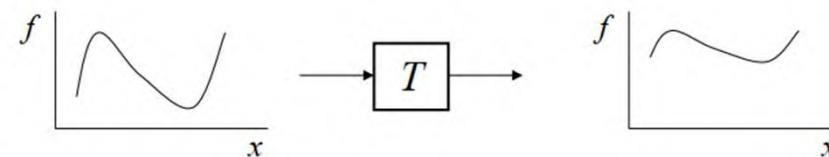
Basic Transform: Enhancement

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.



Different T for different tasks

image filtering: change **range** of image
$$g(x) = T(f(x))$$



Gamma Correction

$$T(c; \gamma) = c^\gamma$$



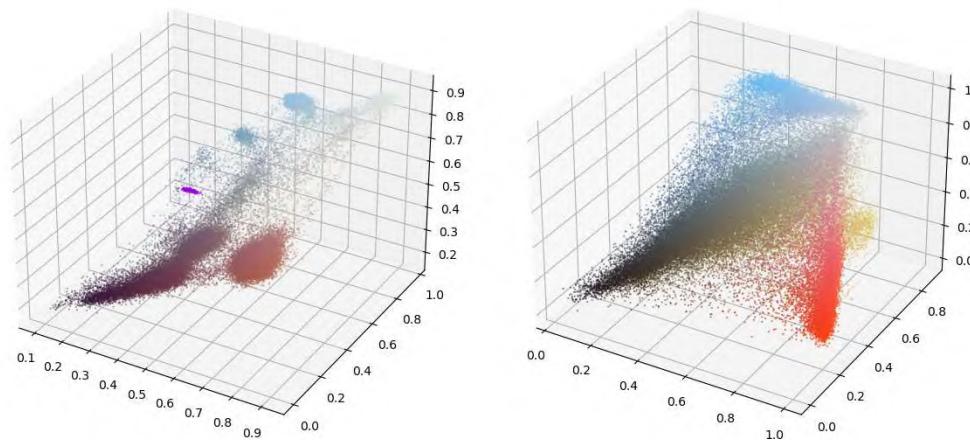
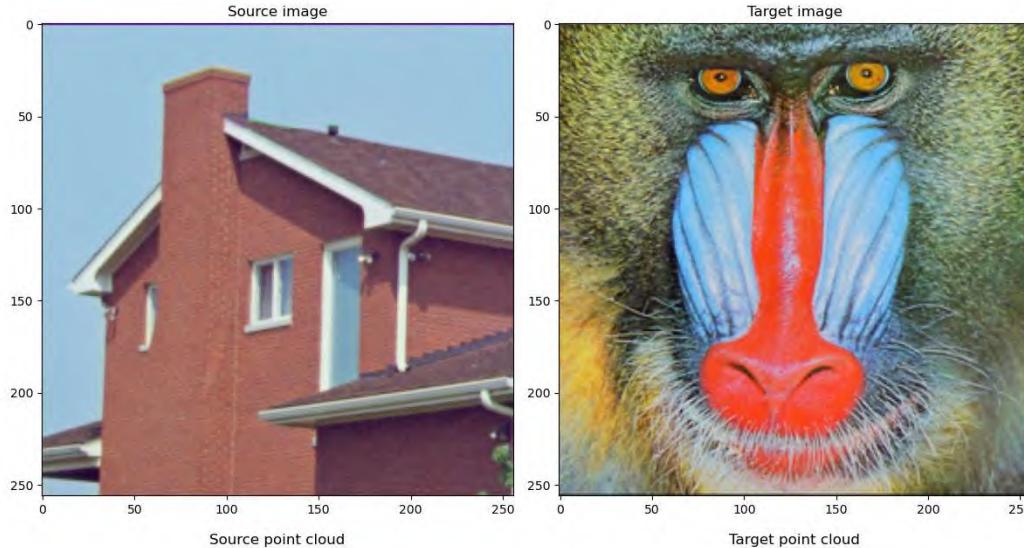
上述变换的特点

Shared transformation for all pixels

- In a digital image, point = pixel.
- Point processing transforms a pixel's value as function of its value alone;
- it does not depend on the values of the pixel's neighbors.

Point Processing of Images
类比：三维几何点云的刚性变换

图像看作3D点云

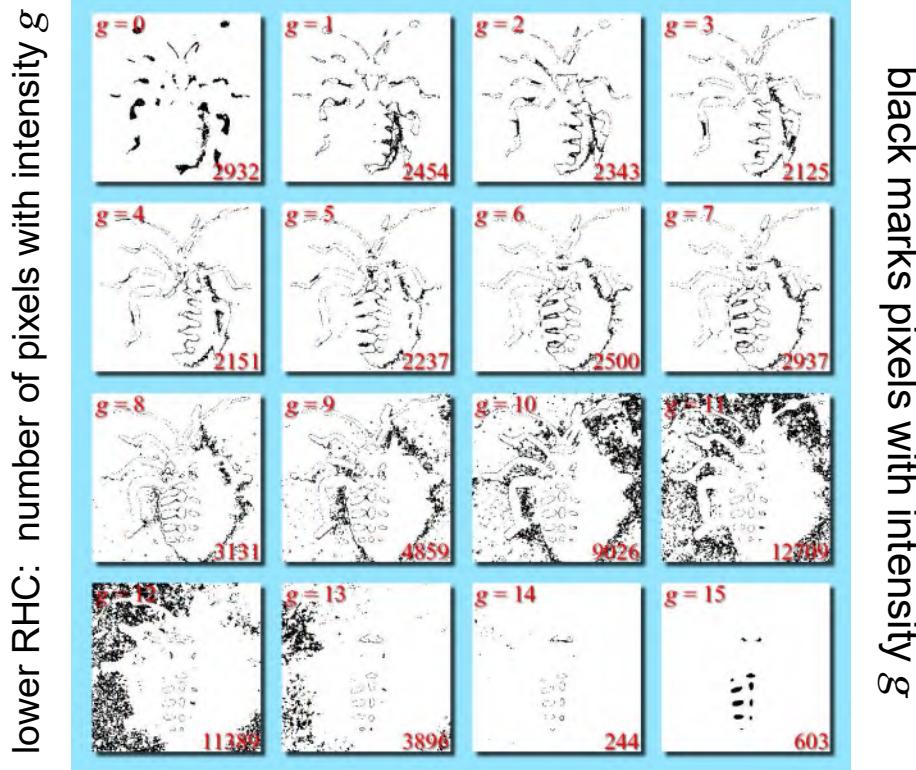


每个像素的颜色
都是三维空间的
一个样本

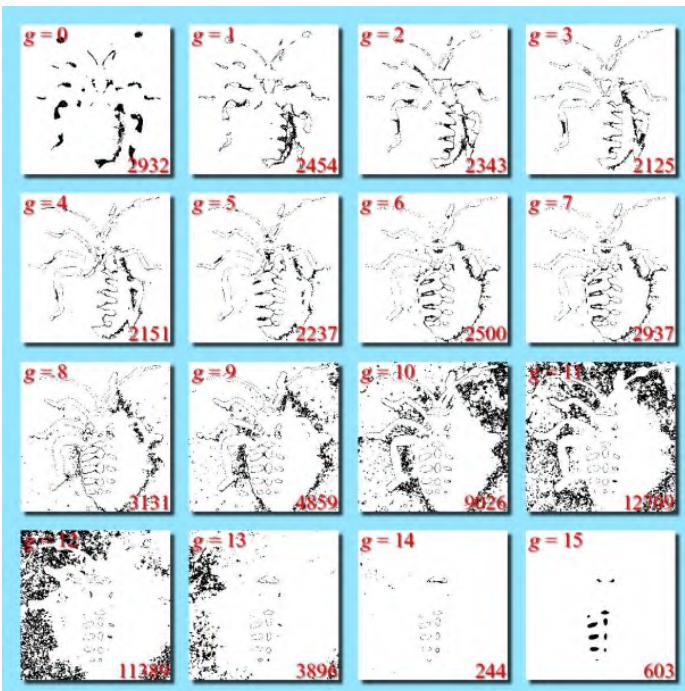
Histogram of (Grayscale) Image



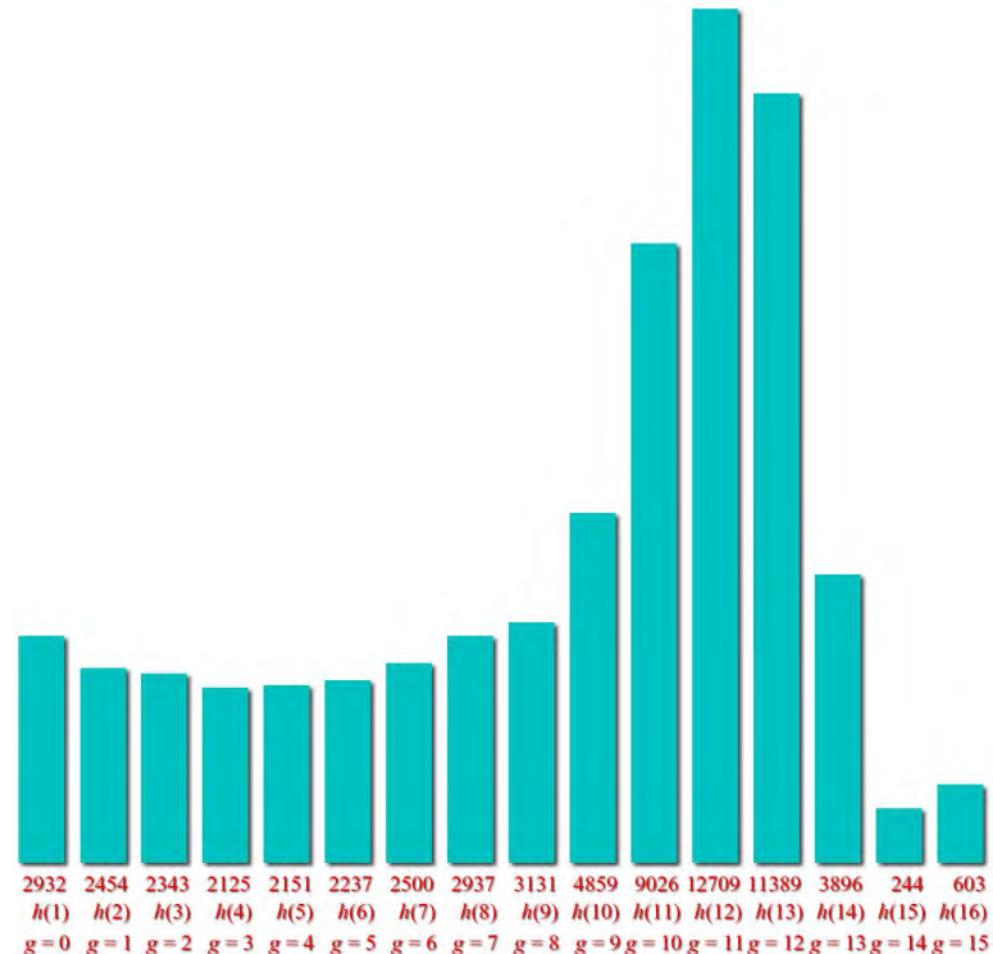
16-level (4-bit) image



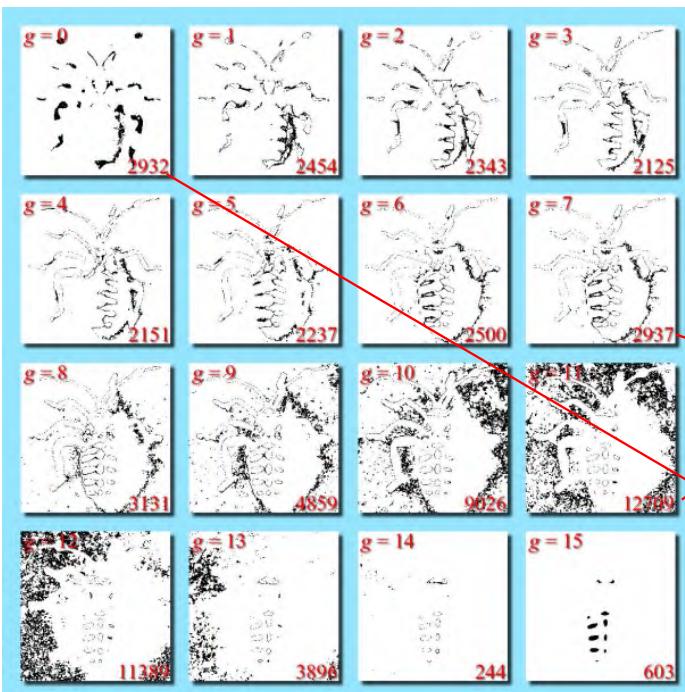
Histogram of Grayscale Image



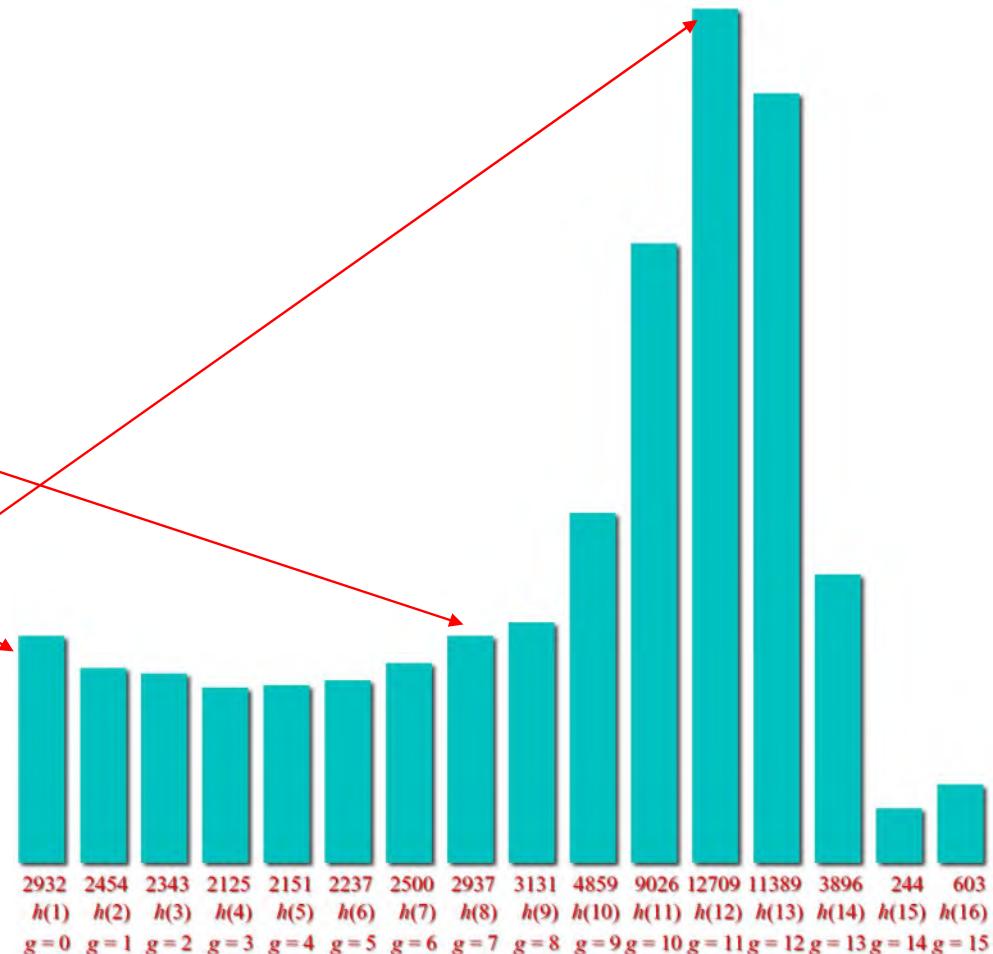
Plot of histogram:
number of pixels with intensity g



Histogram of Grayscale Image



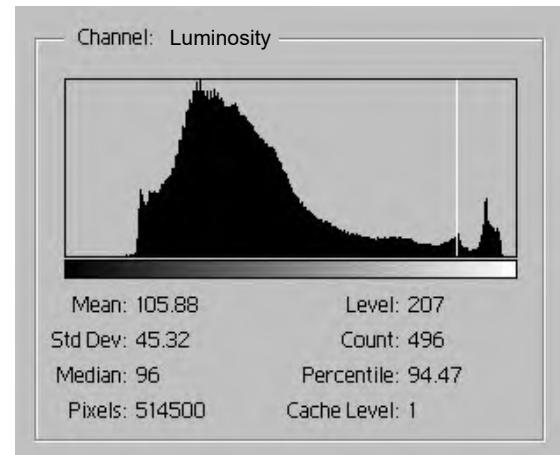
Plot of histogram:
number of pixels with intensity g



Histogram of Grayscale Image

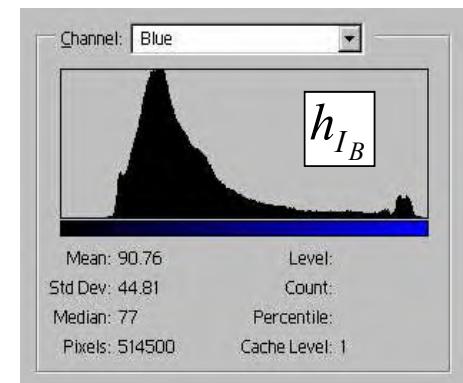
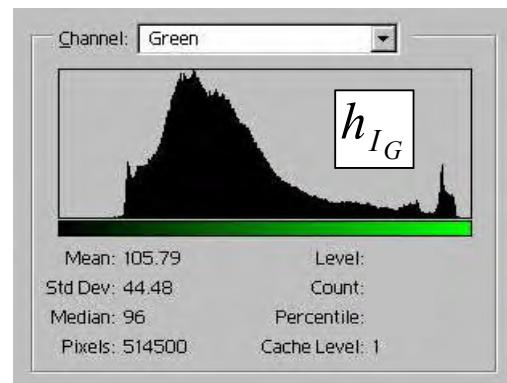
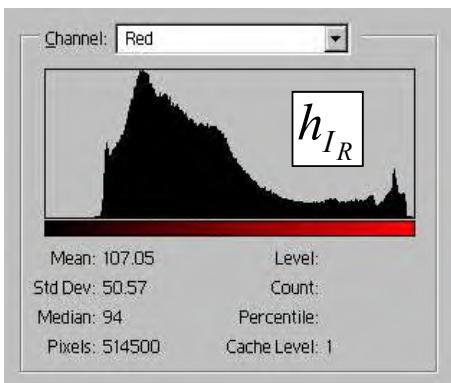
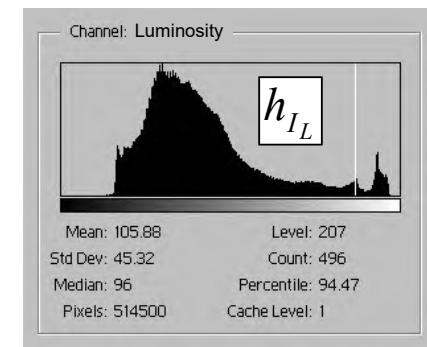


$h_{\mathbf{I}}(g+1) =$ the number of pixels in \mathbf{I} with graylevel g .



Histogram of Color Image

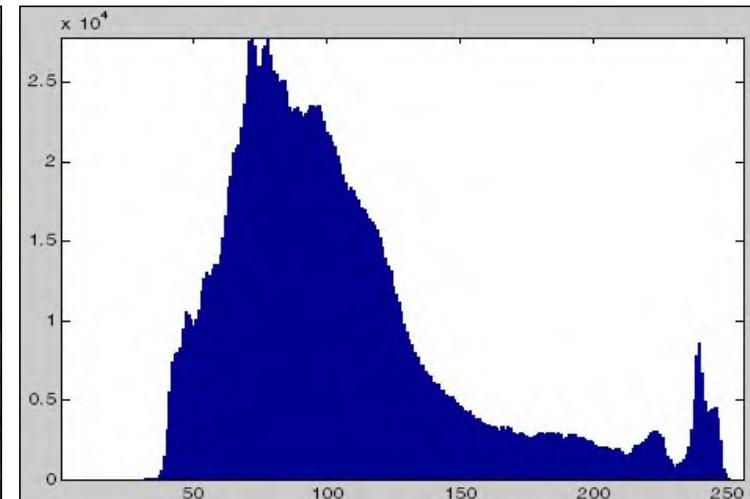
There is one histogram per color band R, G, & B.
Gray histogram
 $0.299*R + 0.587*G + 0.114*B$



Transformations in Histogram



Kinkaku-ji (Temple of the Golden Pavilion), also known as Rokuon-ji (Deer Garden Temple), is a Zen Buddhist temple in Kyoto, Japan.



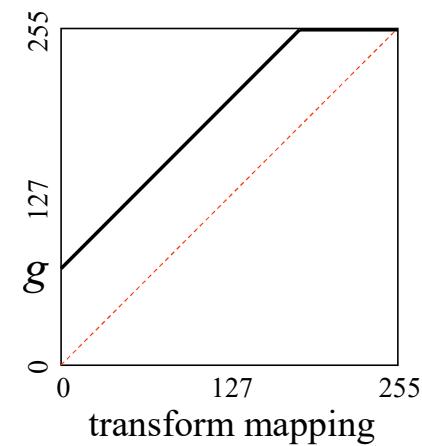
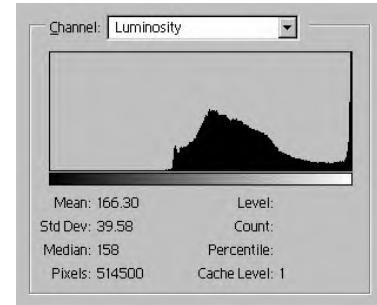
Luminance Histogram

Mean ++



$$\mathbf{J}(r, c, b) = \begin{cases} \mathbf{I}(r, c, b) + g, & \text{if } \mathbf{I}(r, c, b) + g < 256 \\ 255, & \text{if } \mathbf{I}(r, c, b) + g > 255 \end{cases}$$

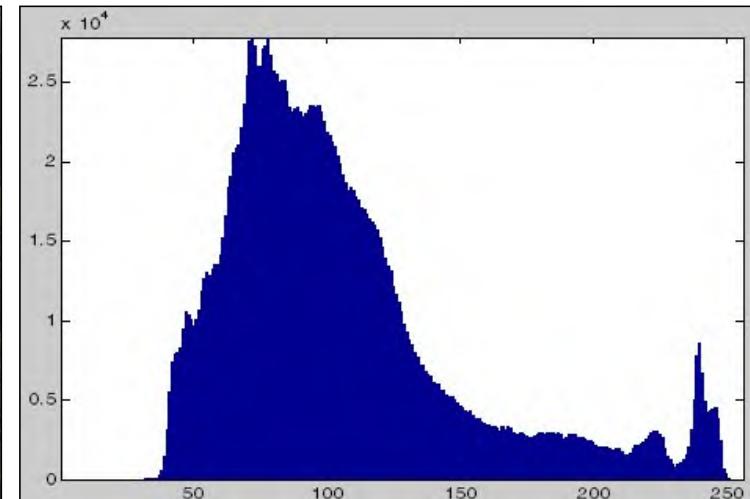
$g \geq 0$ and $b \in \{1, 2, 3\}$ is the band index.



Transformations in Histogram



Kinkaku-ji (Temple of the Golden Pavilion), also known as Rokuon-ji (Deer Garden Temple), is a Zen Buddhist temple in Kyoto, Japan.



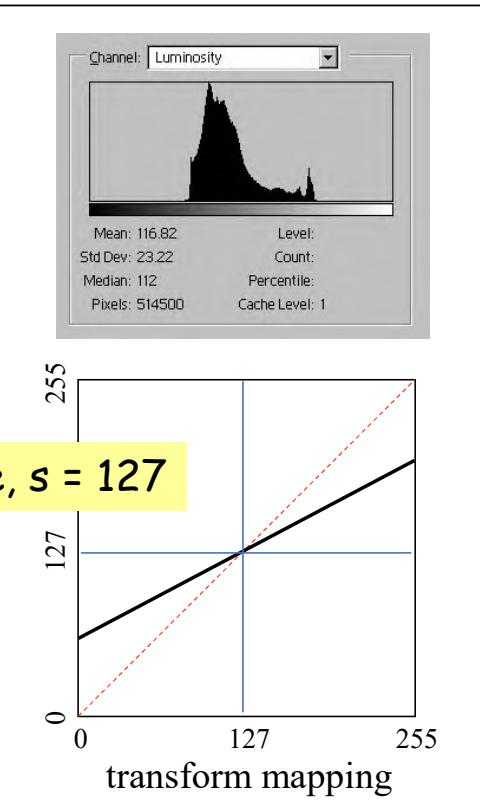
Luminance Histogram

Decrease Contrast

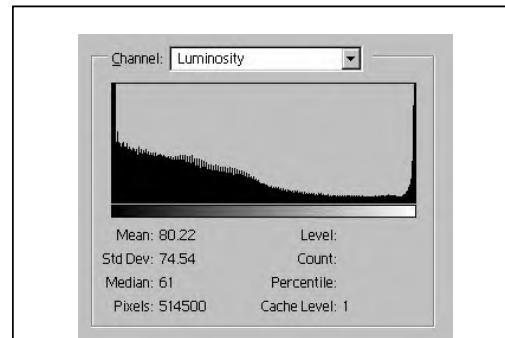


$\mathbf{T}(r,c,b) = a[\mathbf{I}(r,c,b) - s] + s$,
where $0 \leq a < 1.0$,
 $s \in \{0, 1, 2, \dots, 255\}$, and
 $b \in \{1, 2, 3\}$.

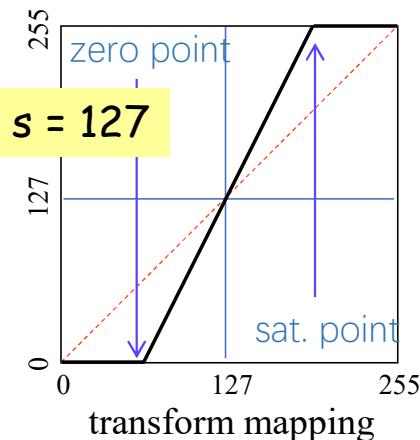
s is the center of the contrast function.



Increase Contrast



Here, $s = 127$

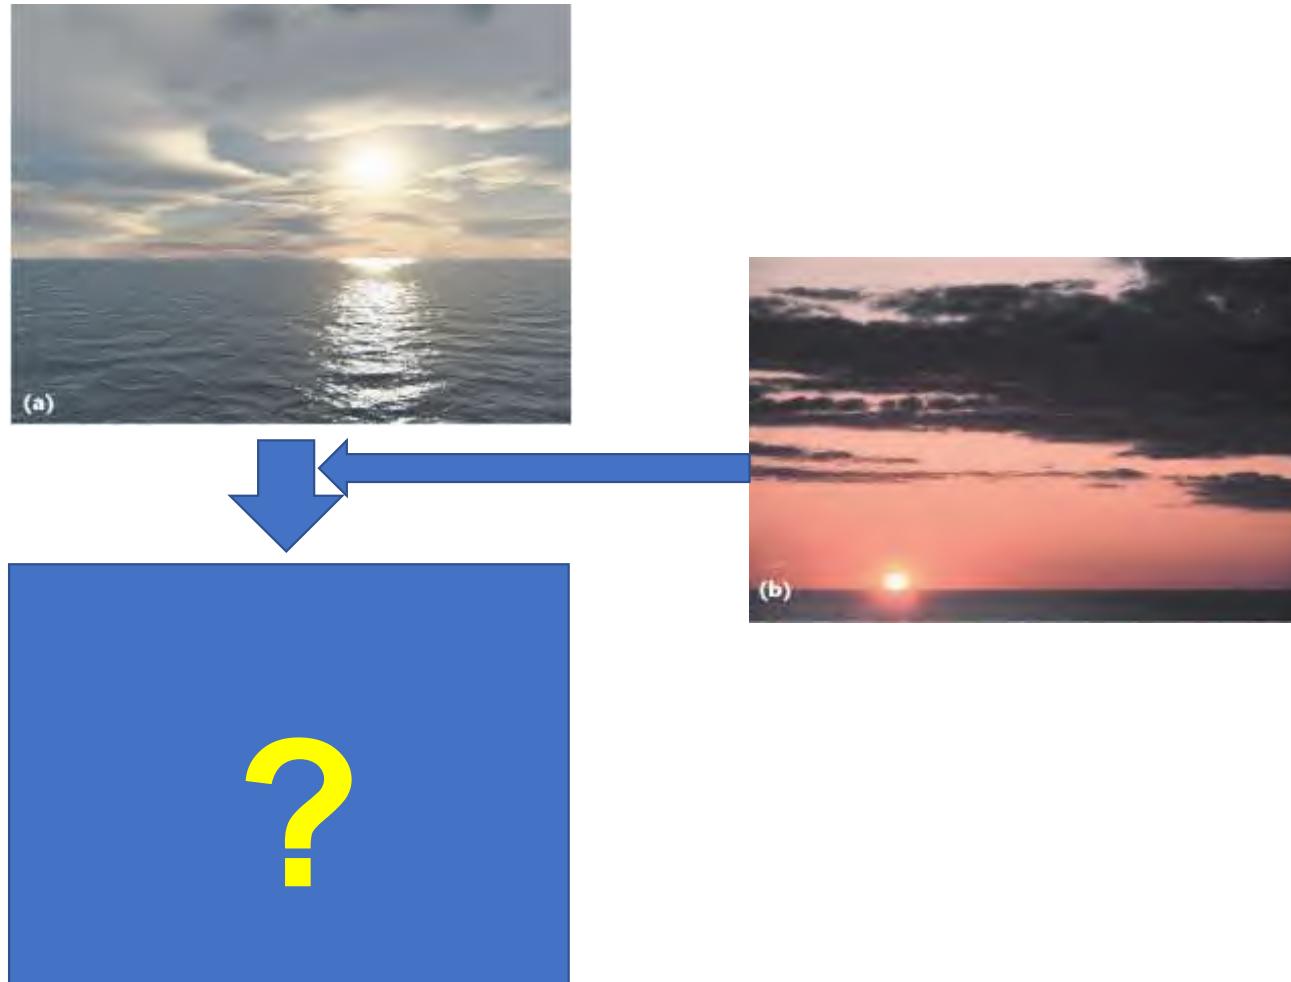


$$\mathbf{T}(r, c, b) = a[\mathbf{I}(r, c, b) - s] + s$$
$$\mathbf{J}(r, c, b) = \begin{cases} 0, & \text{if } \mathbf{T}(r, c, b) < 0, \\ \mathbf{T}(r, c, b), & \text{if } 0 \leq \mathbf{T}(r, c, b) \leq 255, \\ 255, & \text{if } \mathbf{T}(r, c, b) > 255. \end{cases}$$

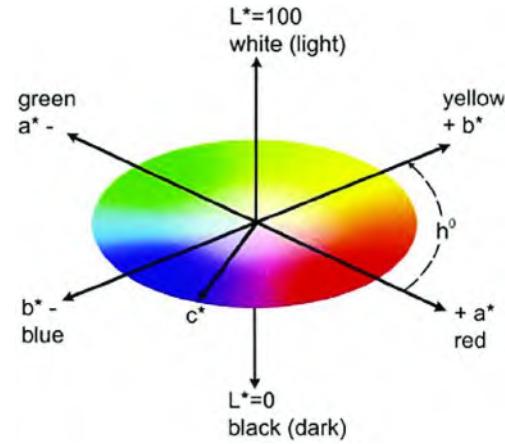
$a > 1, \quad s \in \{0, 127, 255\}, \quad b \in \{1, 2, 3\}$

More Applications

Example1: Color Transfer



Color Transfer Algorithm



$$l^* = l - \langle l \rangle$$

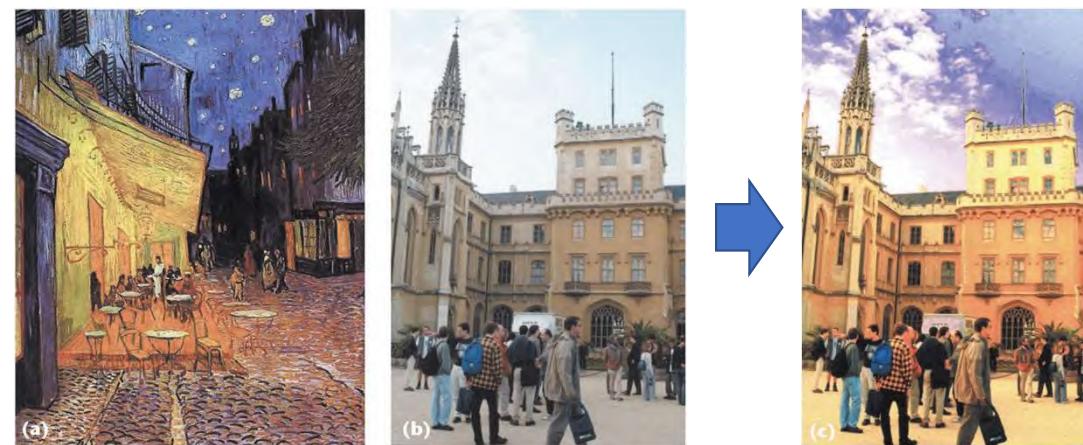
$$\alpha^* = \alpha - \langle \alpha \rangle$$

$$\beta^* = \beta - \langle \beta \rangle$$

$$l' = \frac{\sigma_t^l}{\sigma_s^l} l^*$$

$$\alpha' = \frac{\sigma_t^\alpha}{\sigma_s^\alpha} \alpha^*$$

$$\beta' = \frac{\sigma_t^\beta}{\sigma_s^\beta} \beta^*$$



[Color Transfer between Images](#). IEEE CG&A, 2001.

Example2: Color2Gray



$$0.299 * R + 0.587 * G + 0.114 * B$$



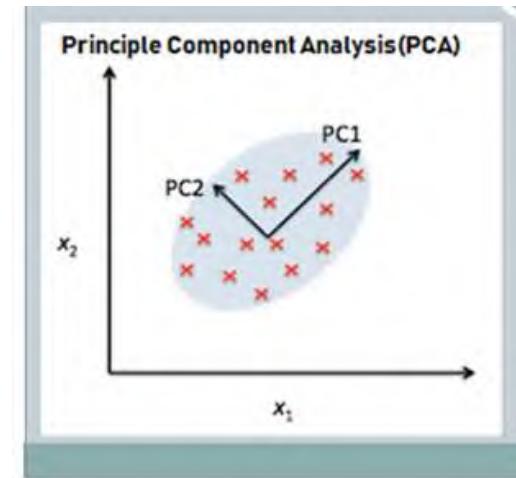
Example2: Color2Gray



Color2Gray Algorithm



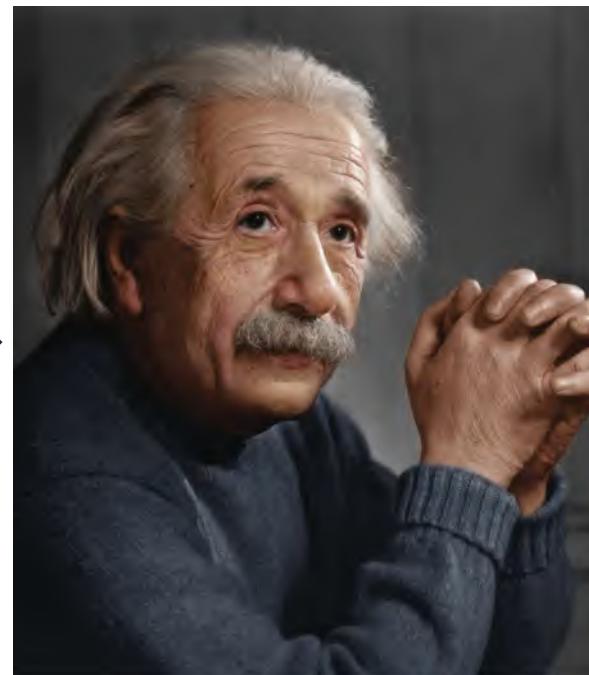
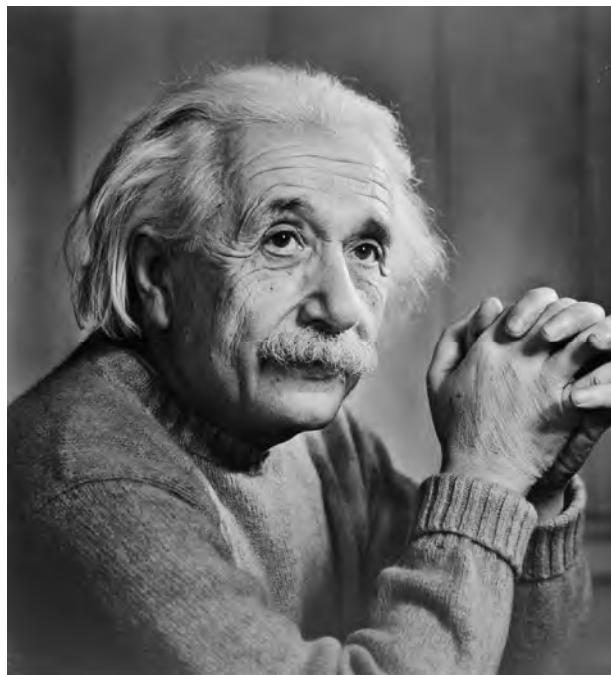
Dim Reduction



$$f(g) = \sum_{(i,j) \in \mathcal{K}} ((g_i - g_j) - \delta_{ij})^2$$

[Color2Gray: Salience-Preserving Color Removal](#). SIGGRAPH 2005.

Example3: Gray2Color (or Colorization)



Inherent Ambiguity



1D->3D, 欠定问题

Inherent Ambiguity



1D->3D, 欠定问题

Strand based Colorization

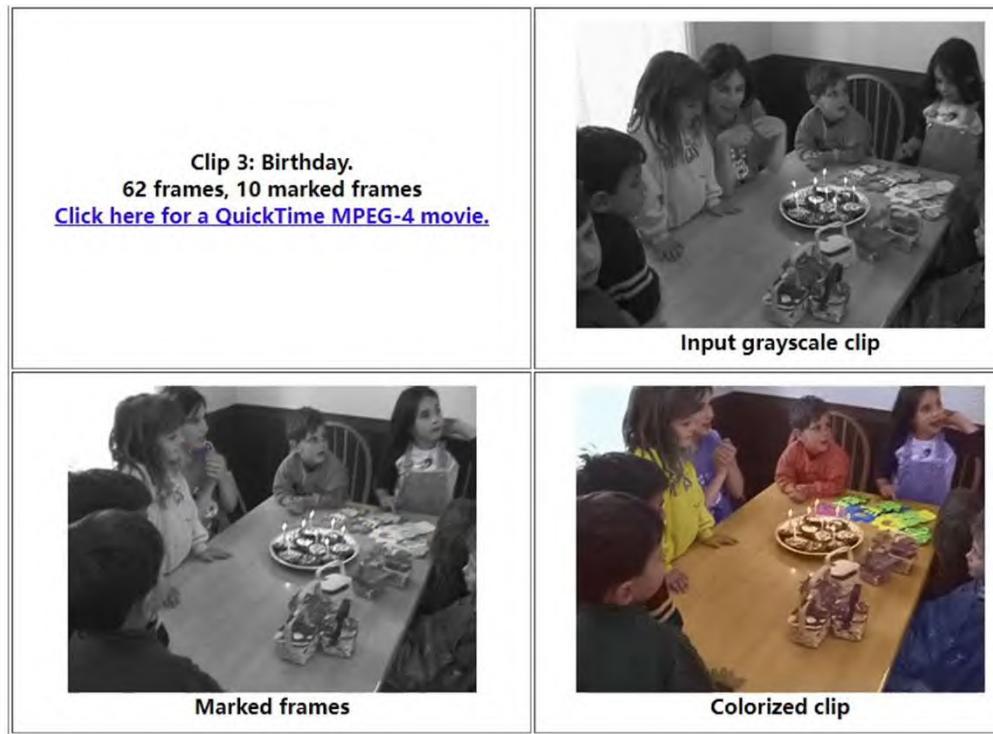


$$J(U) = \sum_{\mathbf{r}} \left(U(\mathbf{r}) - \sum_{\mathbf{s} \in N(\mathbf{r})} w_{\mathbf{rs}} U(\mathbf{s}) \right)^2$$

$$w_{\mathbf{rs}} \propto 1 + \frac{1}{\sigma_{\mathbf{r}}^2} (Y(\mathbf{r}) - \mu_{\mathbf{r}})(Y(\mathbf{s}) - \mu_{\mathbf{r}})$$

[Colorization using Optimization](#). SIGGRAPH 2004.

Strand based Colorization



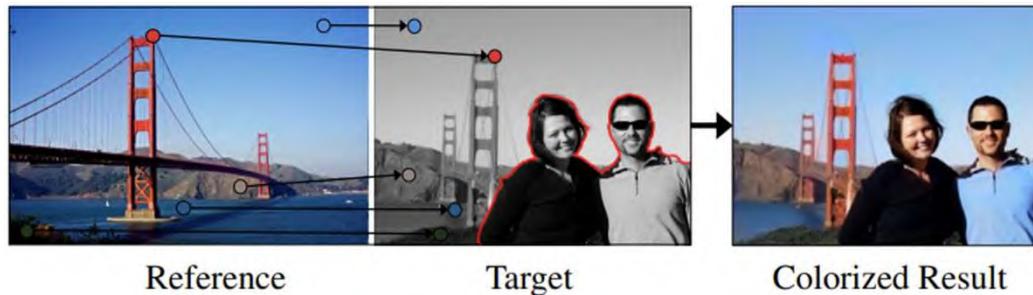
[Colorization using Optimization](#). SIGGRAPH 2004.

Reference based Colorization



[Deep Exemplar-based Colorization](#). SIGGRAPH 2018.

Example based Colorization



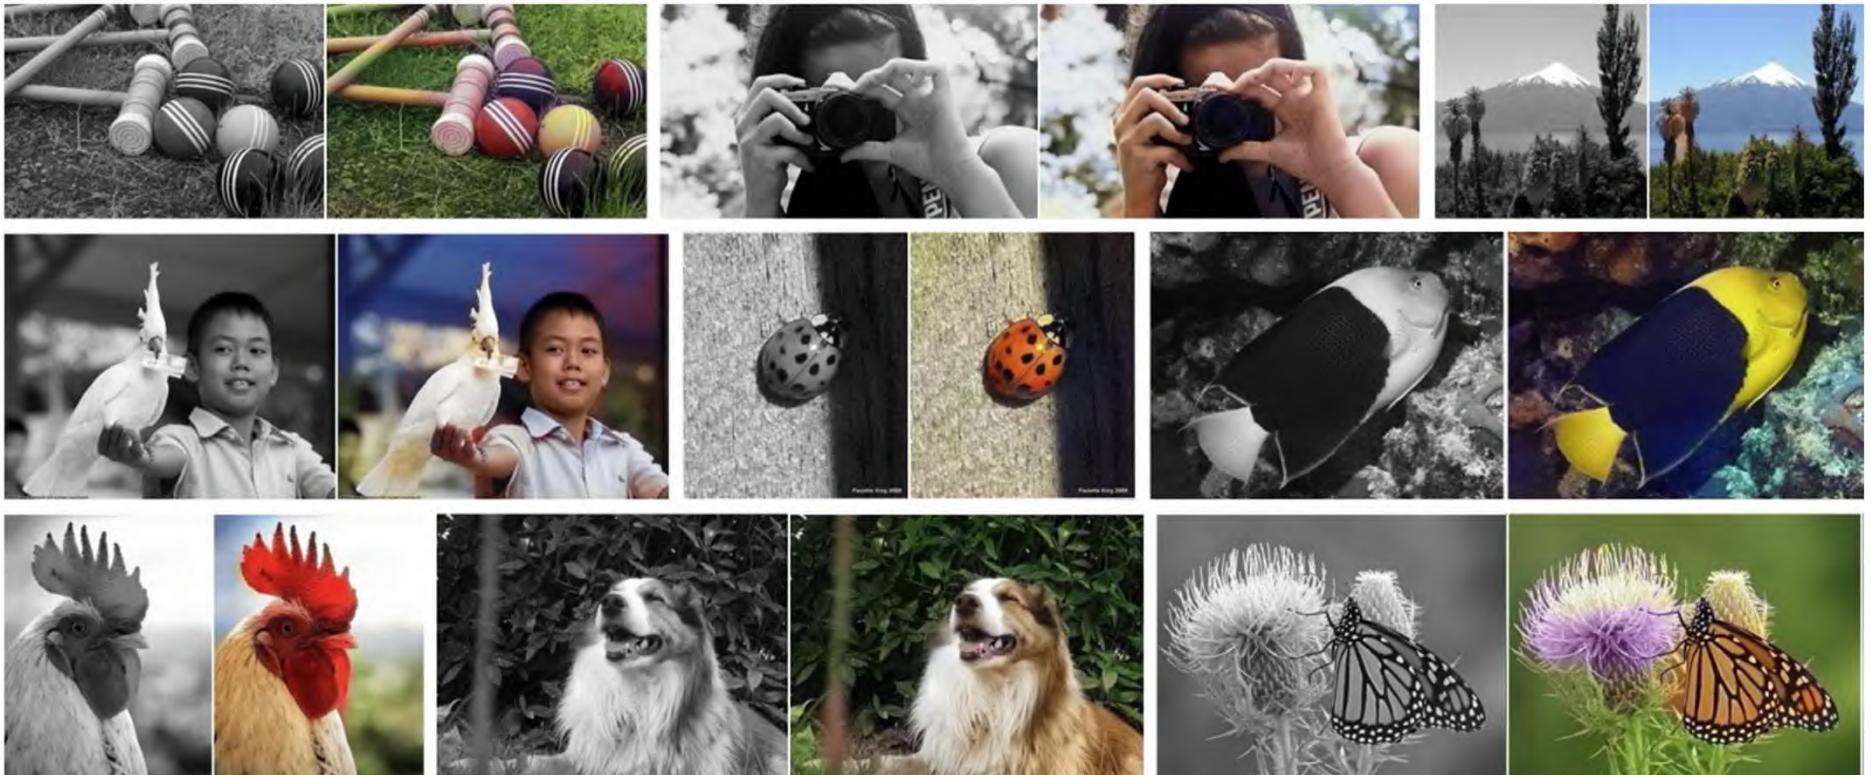
$$\begin{aligned} sim_{T \rightarrow R}^i(p) &= d(F_T^i(p), F_R^i(\phi_{T \rightarrow R}(p))), \\ sim_{R \rightarrow T}^i(p) &= d(F_T^i(\phi_{R \rightarrow T}(\phi_{T \rightarrow R}(p))), F_R^i(\phi_{T \rightarrow R}(p))) . \end{aligned} \quad (1)$$

As mentioned in Liao et al. [2017], cosine distance performs better in measuring feature similarity since it is more robust to appearance variances between image pairs. Thus, our similarity metric $d(x, y)$ between two deep features is defined as their cosine similarity:

$$d(x, y) = \frac{x^T y}{|x||y|} . \quad (2)$$

[Deep Exemplar-based Colorization](#). SIGGRAPH 2018.

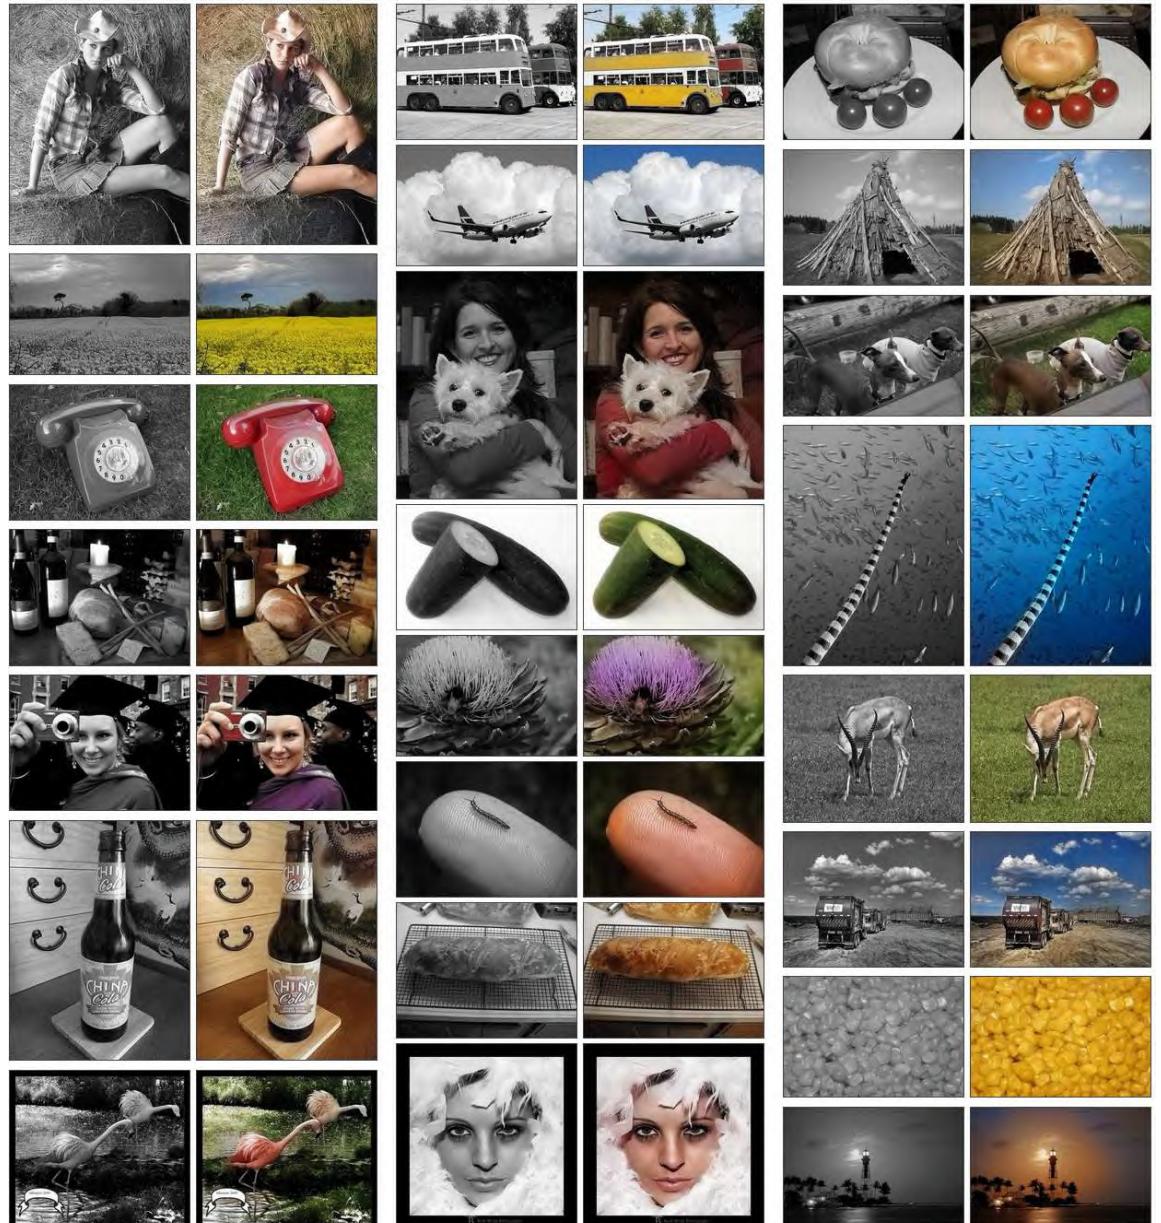
Reference free Colorization



[Colorful Image Colorization](#). ECCV 2016.

Data Driven

Color2gray is easy



Reference free Colorization

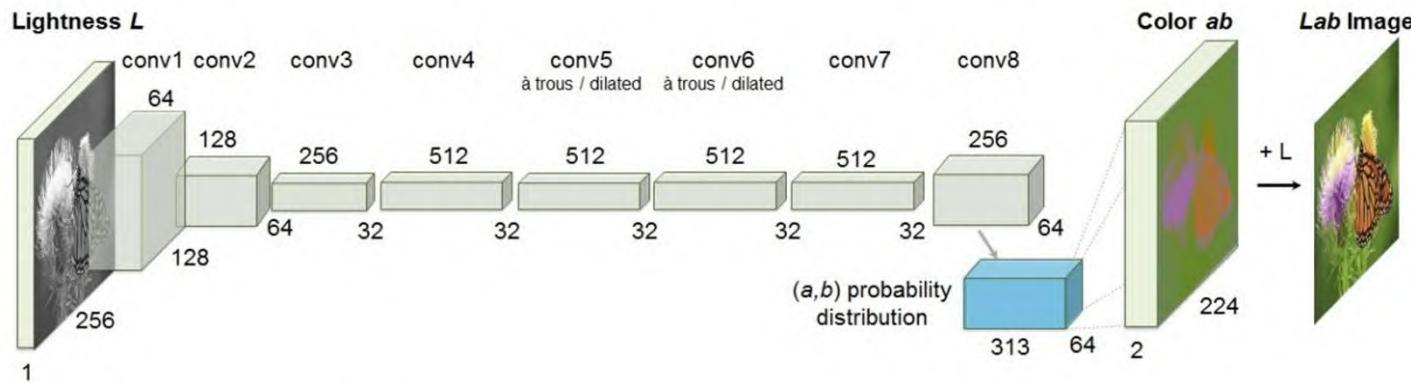
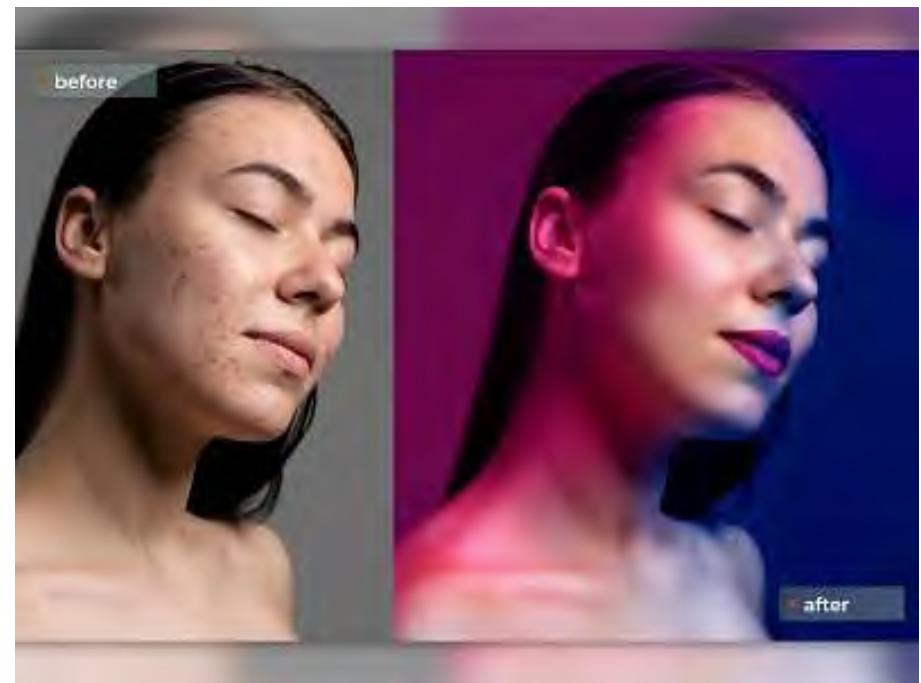
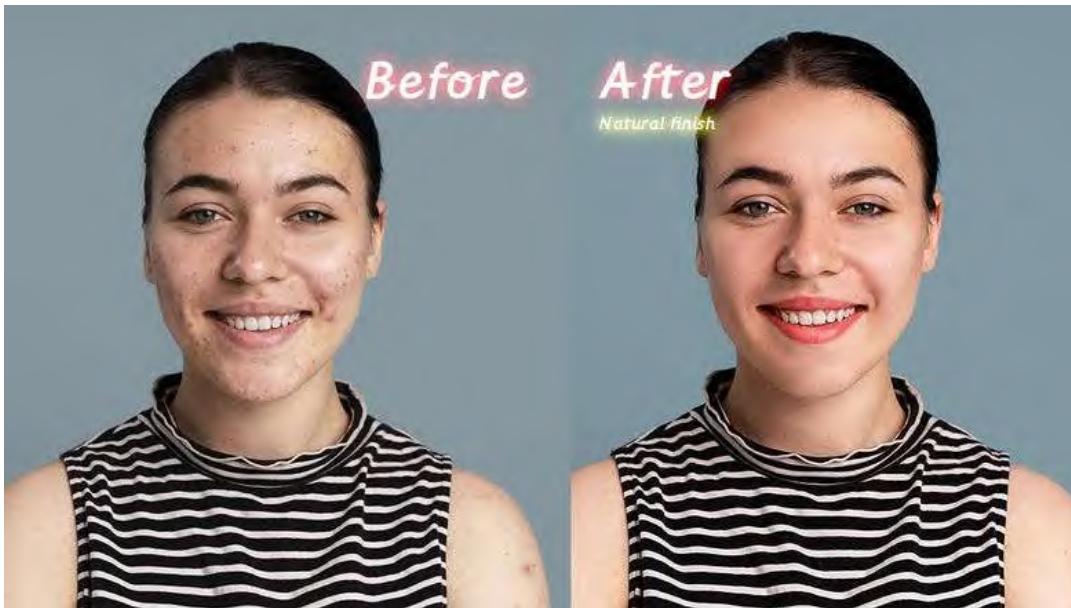


Fig. 2. Our network architecture. Each **conv** layer refers to a block of 2 or 3 repeated **conv** and ReLU layers, followed by a BatchNorm [30] layer. The net has no pool layers. All changes in resolution are achieved through spatial downsampling or upsampling between **conv** blocks.

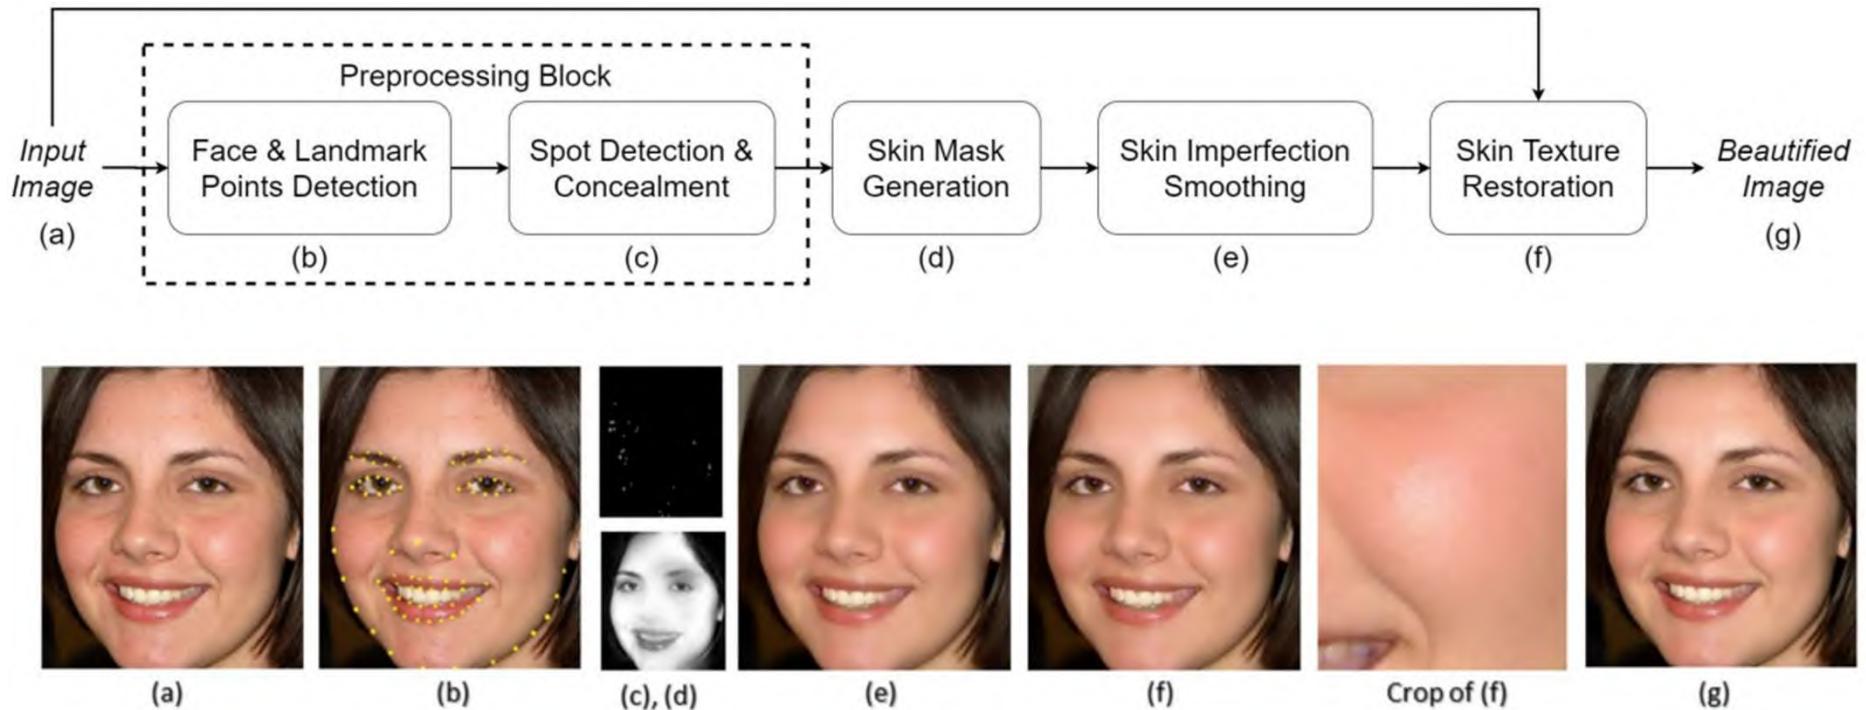
$$L_2(\hat{\mathbf{Y}}, \mathbf{Y}) = \frac{1}{2} \sum_{h,w} \|\mathbf{Y}_{h,w} - \hat{\mathbf{Y}}_{h,w}\|_2^2$$

[Colorful Image Colorization](#). ECCV 2016.

Example4: Face Retouch



Face Retouch Algorithm



[Face Beautification via Dynamic Skin Smoothing, Guided
Feathering, and Texture Restoration](#). CVPRW 2020.

Face Retouch Algorithm

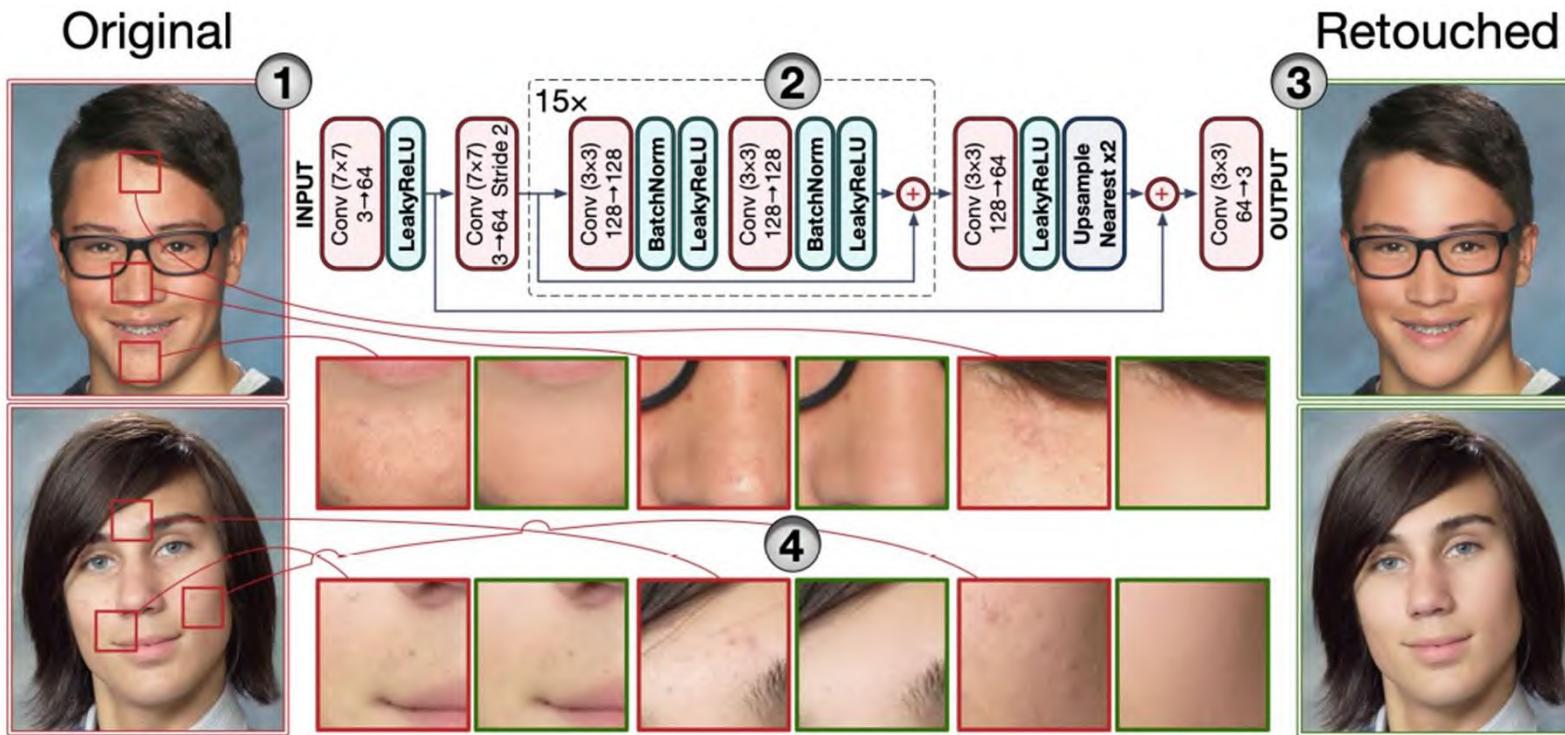
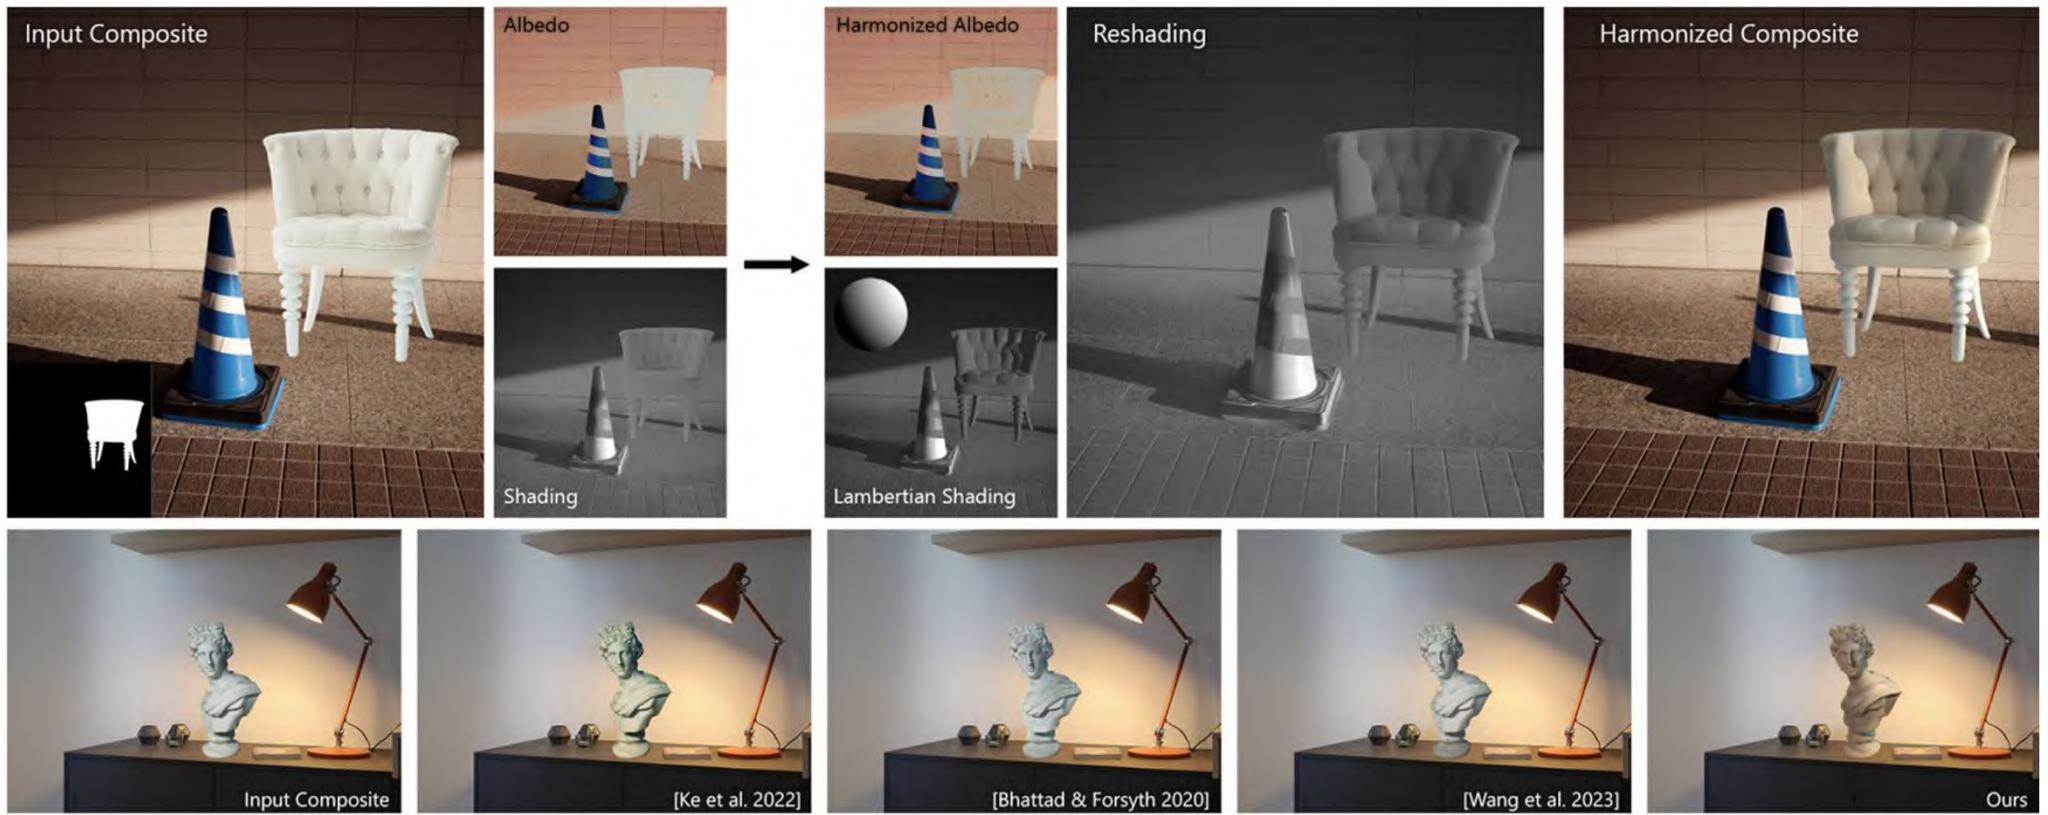


Figure 1: Overview of AutoRetouch: (1) input portraits cropped to the head, (2) the images pass through the fully convolutional architecture of AutoRetouch, (3) the output portraits are retouched while preserving fine textures. (4) Sample before-after patches scaled up. We empirically show that our final models have desirable generalization properties.

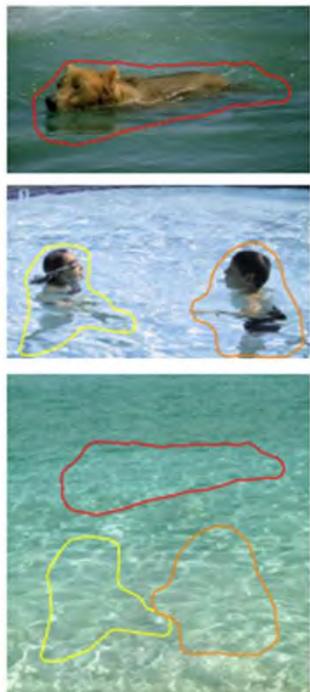
[AutoRetouch: Automatic Professional Face Retouching](#). WACV 2021.

Example5: Image Composition



[Intrinsic Harmonization for Illumination-Aware Compositing. SIGGRAPH Asia 2023.](#)

A Simple Image Composition



sources/destinations

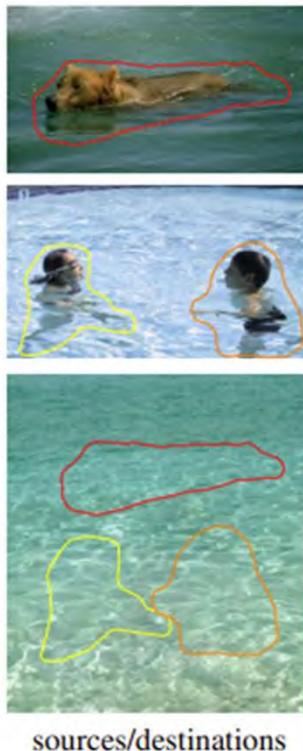


cloning



seamless cloning

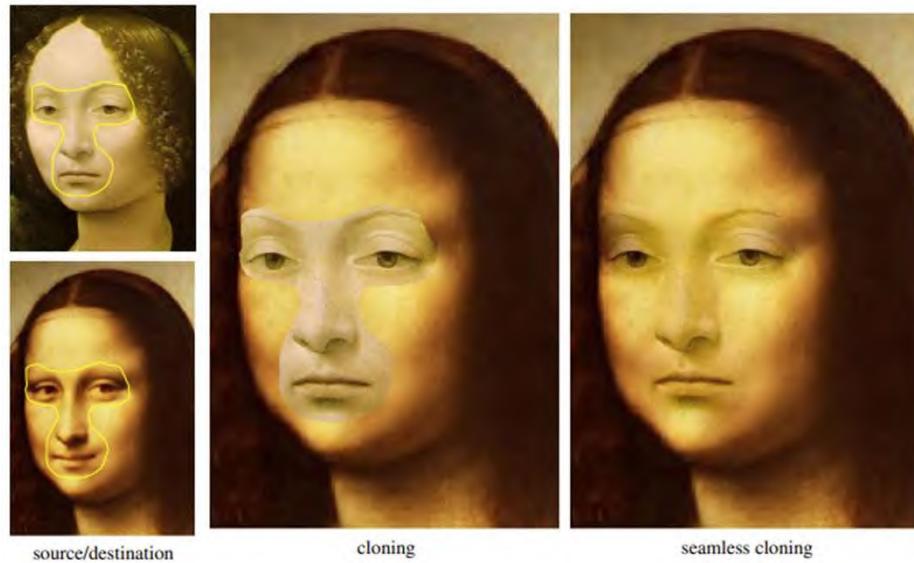
A Simple Image Composition



Same Boundary

Same Color Delta

A Simple Image Composition



A *guidance field* is a vector field \mathbf{v} used in an extended version of the minimization problem (1) above:

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (3)$$

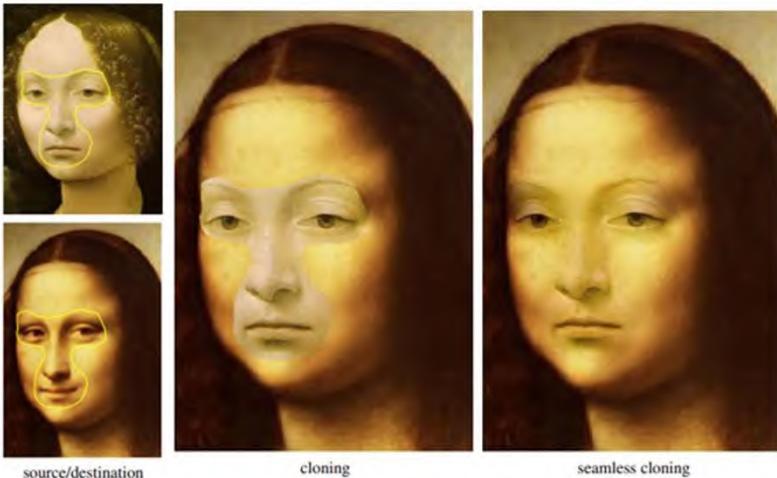
whose solution is the unique solution of the following Poisson equation with Dirichlet boundary conditions:

$$\Delta f = \operatorname{div} \mathbf{v} \text{ over } \Omega, \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (4)$$

where $\operatorname{div} \mathbf{v} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y}$ is the divergence of $\mathbf{v} = (u, v)$. This is the

[Poisson Image Editing](#). TOG 2003.

Example Review



$$J(U) = \sum_{\mathbf{r}} \left(U(\mathbf{r}) - \sum_{\mathbf{s} \in N(\mathbf{r})} w_{\mathbf{rs}} U(\mathbf{s}) \right)^2$$

$$w_{\mathbf{rs}} \propto 1 + \frac{1}{\sigma_{\mathbf{r}}^2} (Y(\mathbf{r}) - \mu_{\mathbf{r}})(Y(\mathbf{s}) - \mu_{\mathbf{r}})$$

A *guidance field* is a vector field \mathbf{v} used in an extended version of the minimization problem (1) above:

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (3)$$

whose solution is the unique solution of the following Poisson equation with Dirichlet boundary conditions:

$$\Delta f = \operatorname{div} \mathbf{v} \text{ over } \Omega, \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (4)$$

where $\operatorname{div} \mathbf{v} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y}$ is the divergence of $\mathbf{v} = (u, v)$. This is the

Shared transformation for all pixels

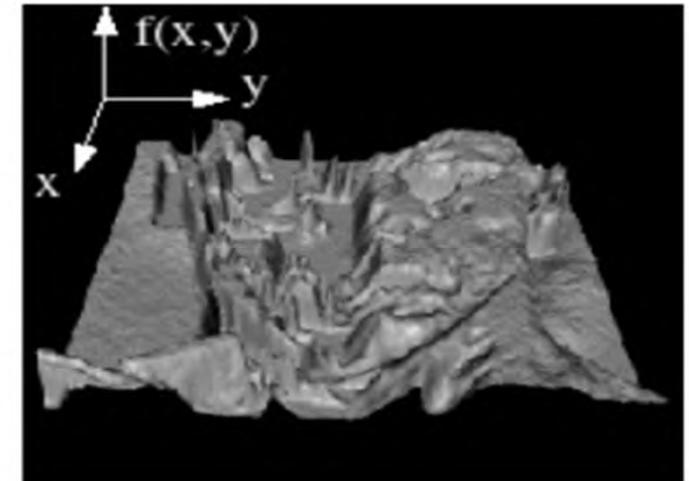
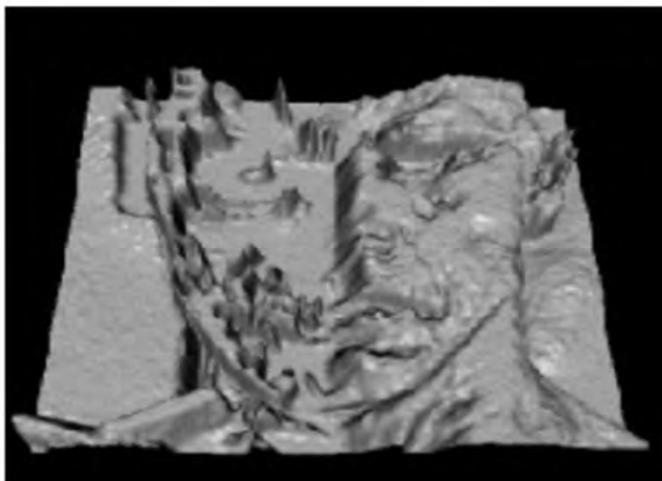
- In a digital image, point = pixel.
- Point processing transforms a pixel's value as function of its value alone;
- it does not depend on the values of the pixel's neighbors.



Neighboring / Coordinate
Dependent Transform

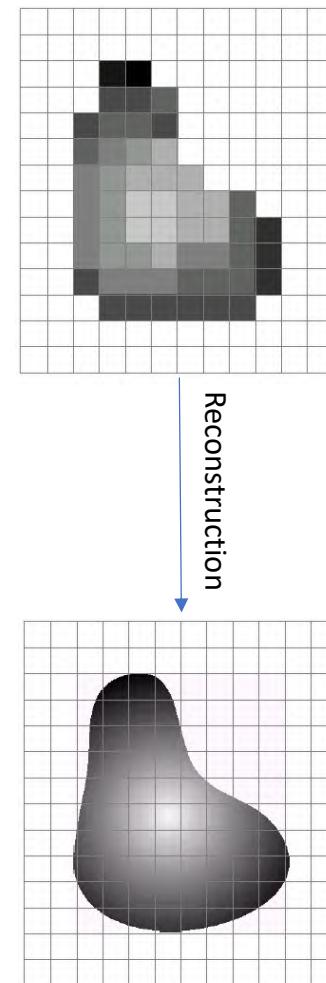
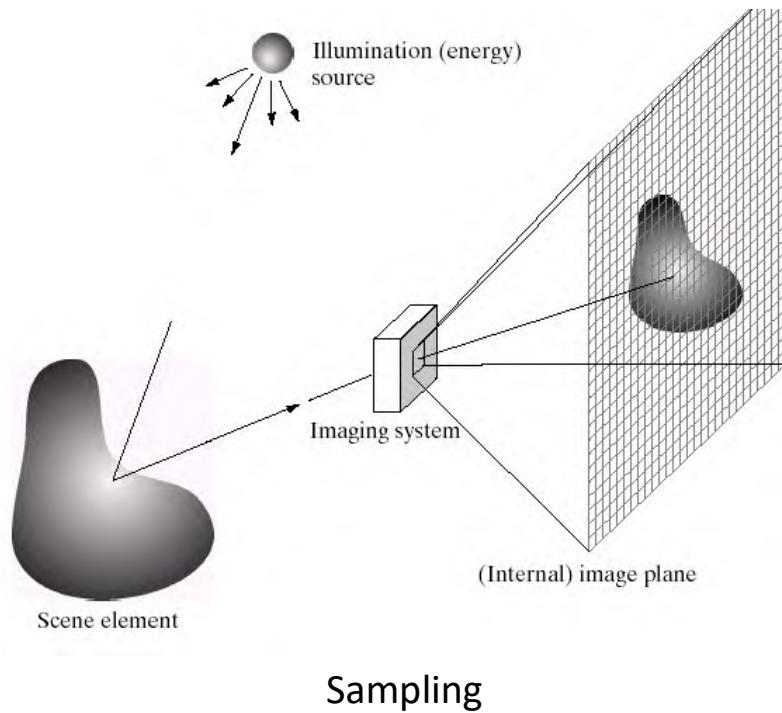
Point Processing of Images
类比：三维几何点云的刚性变换

Image as **Continuous** function of Coordinates

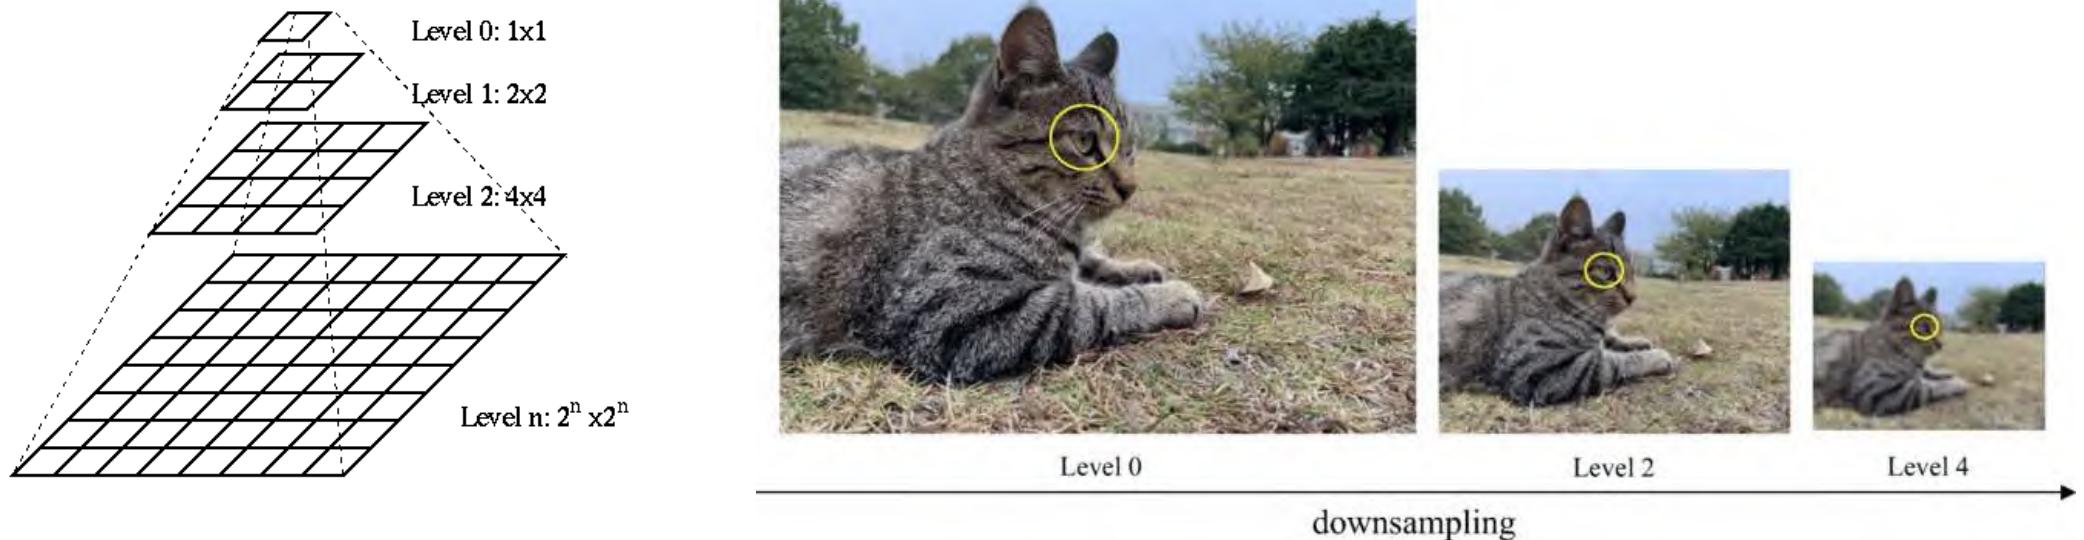


Visualization of single channel grayscale image $f(x,y)$

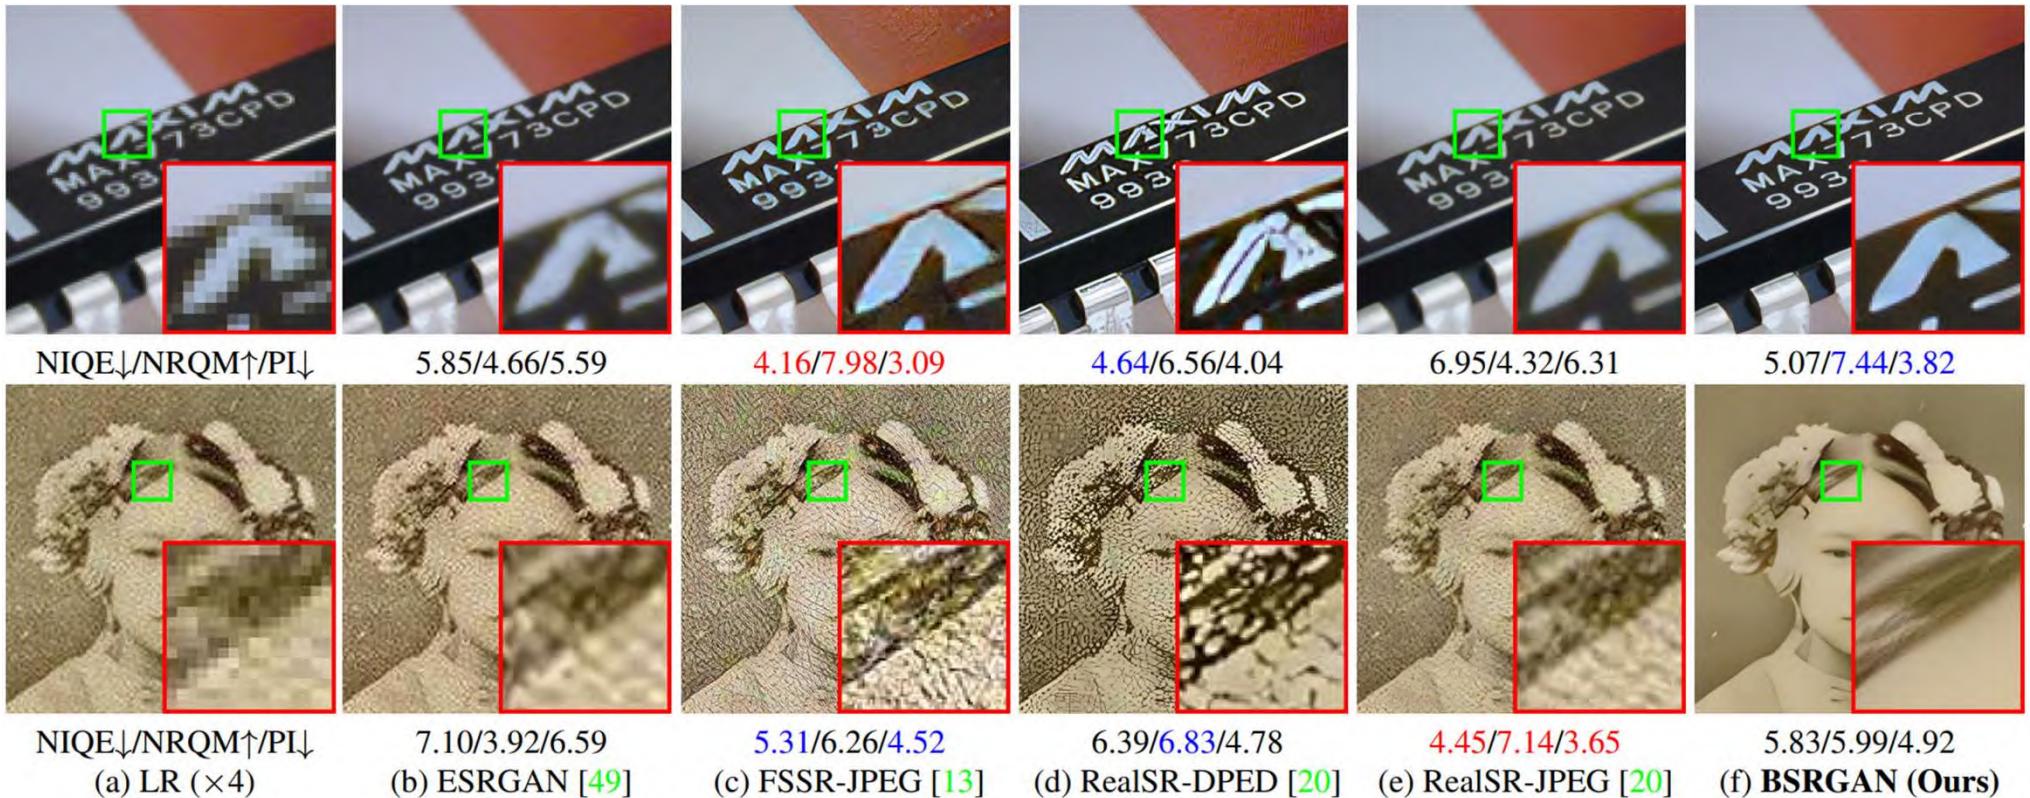
Sampling (Non-continuous) and Reconstruction



Downsampling



Upsampling (Super-resolution)

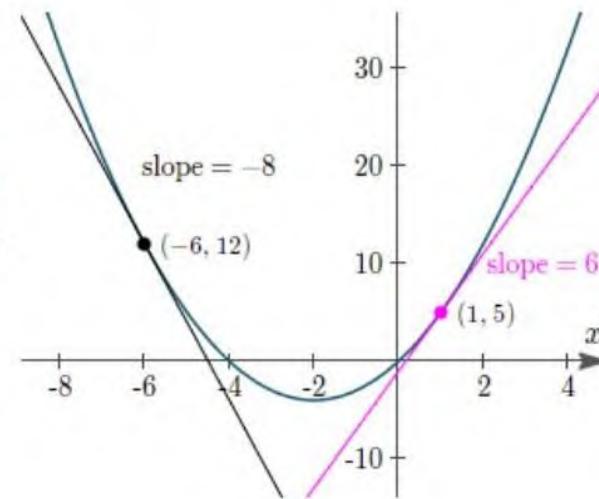


Discrete Operations

DIFFERENTIAL CALCULUS

Gradient of any point on CURVE

$$\frac{dy}{dx} = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$



T. Prashant

Image Gradient

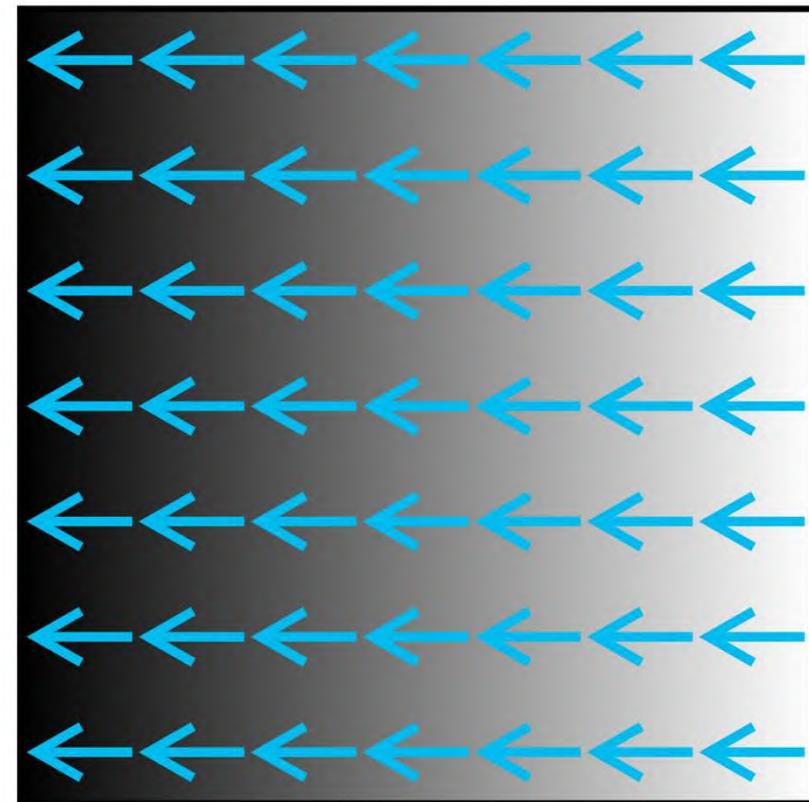
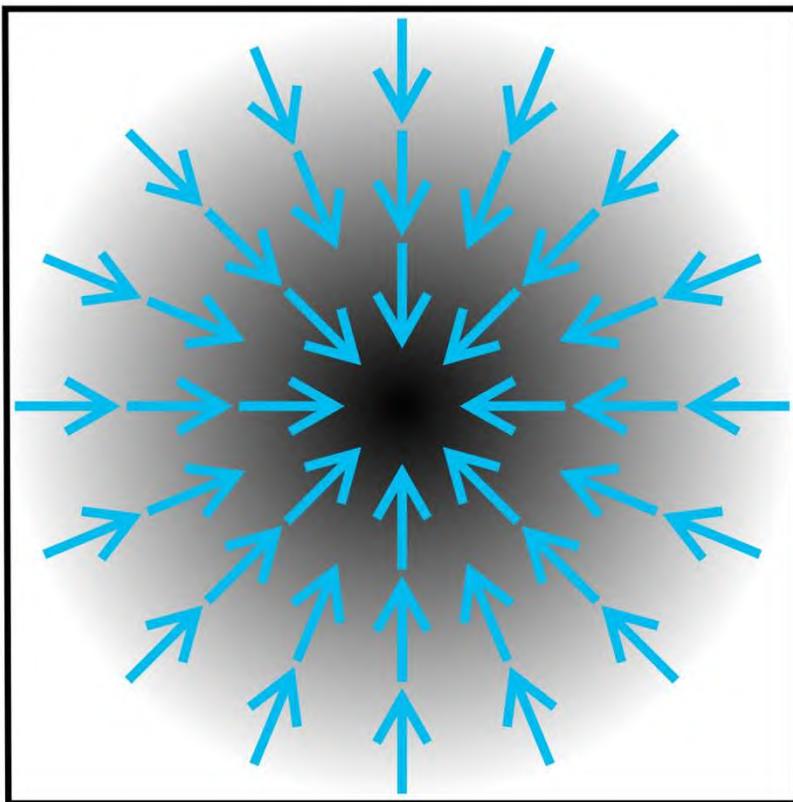
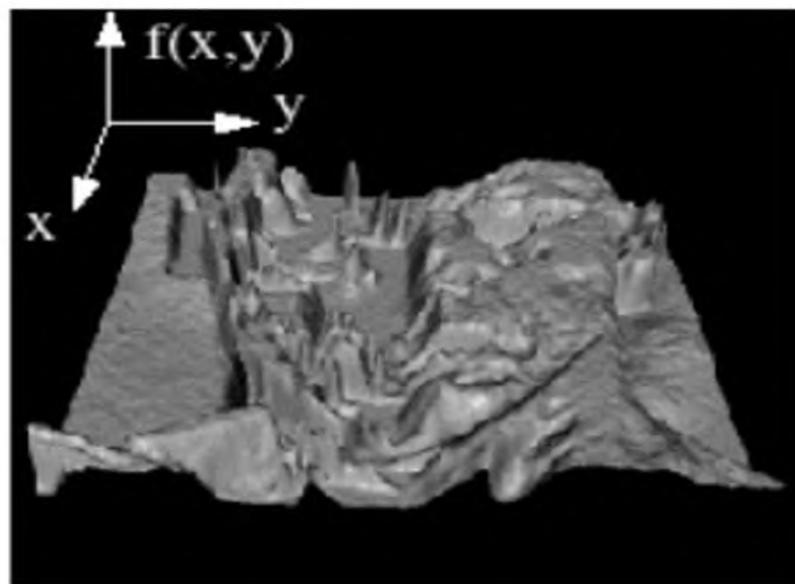


Image Gradient

$$\nabla f = \begin{bmatrix} g_x \\ g_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$



$$g_x = f(x + 1, y) - f(x, y)$$

$$g_y = f(x, y + 1) - f(x, y)$$

or

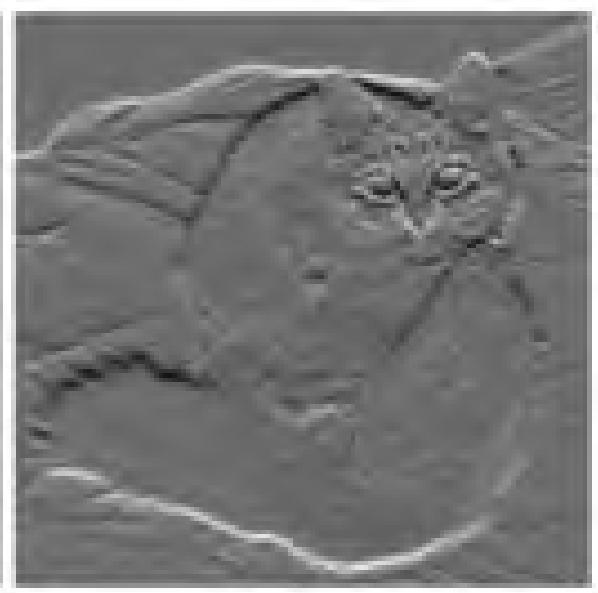
$$g_x = f(x + 1, y) - f(x - 1, y)$$

$$g_y = f(x, y + 1) - f(x, y - 1)$$

Image Gradient Visualization

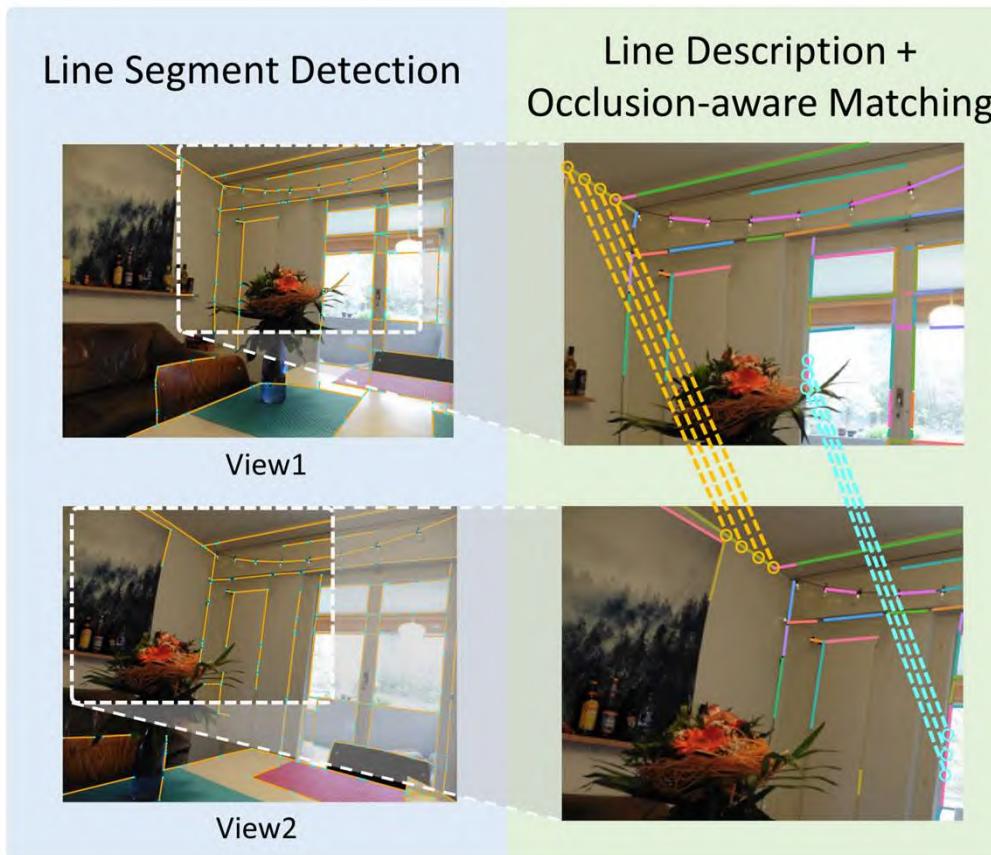


X-gradient



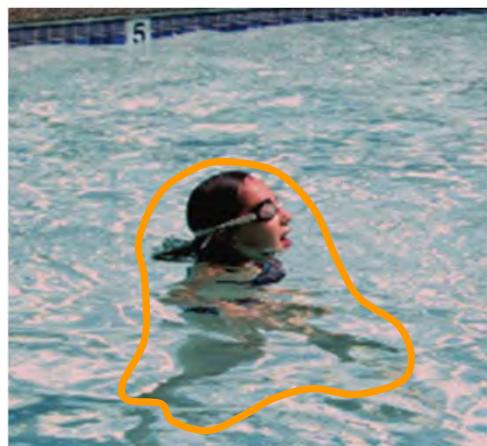
Y-gradient

Application: Image -> Edge (Sketch)



[SOLD2: Self-supervised Occlusion-aware Line Description and Detection. CVPR 2021.](#)

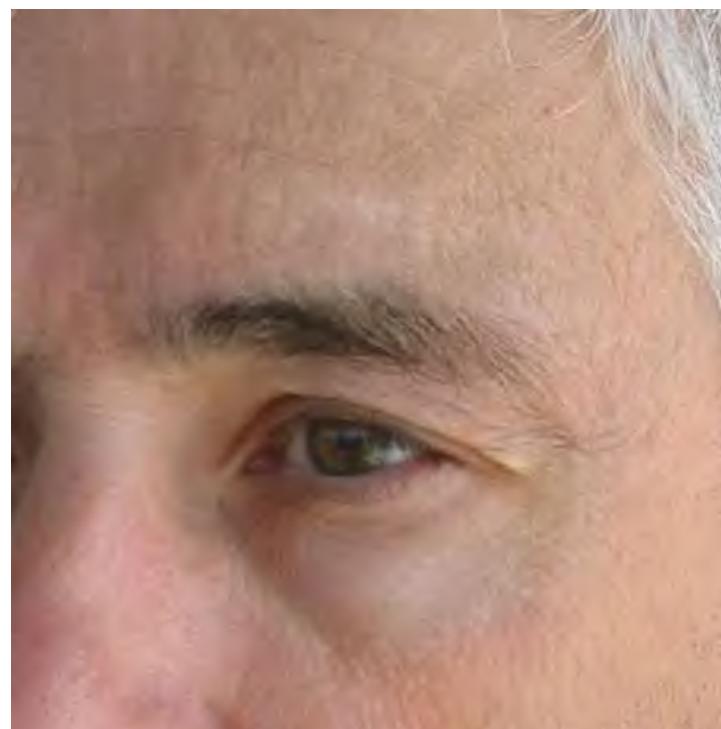
Application: Image Composition



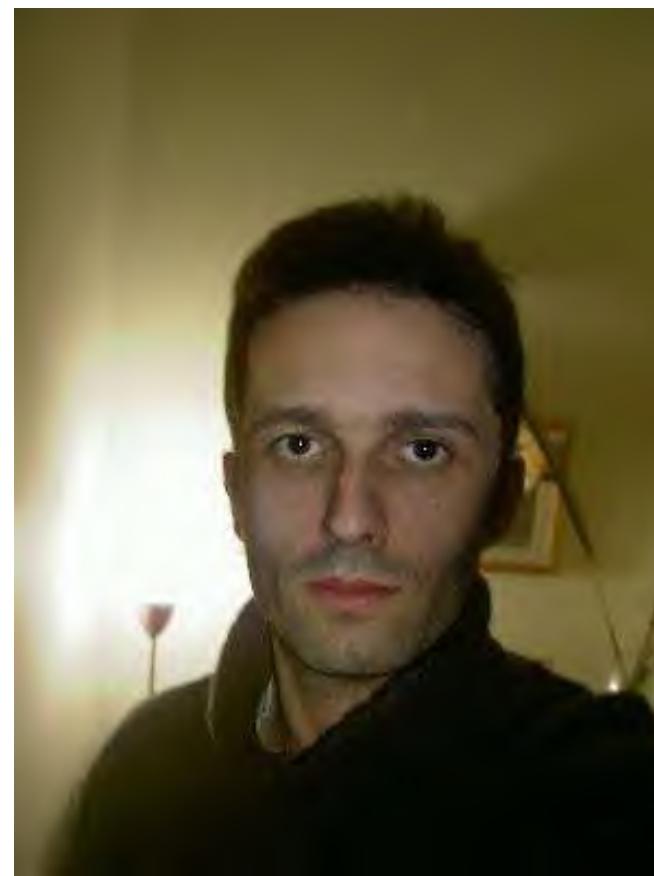
Application: Texture Exchange



Application: Wrinkle Removal



Application: Lighting Interpolation



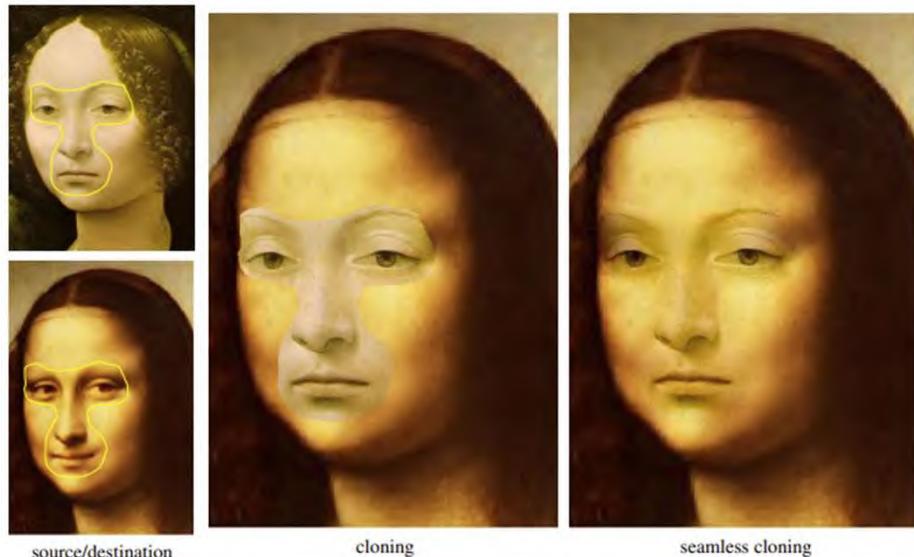
Application: Color Transfer



Application: Color Transfer



How to efficiently implement



A *guidance field* is a vector field \mathbf{v} used in an extended version of the minimization problem (1) above:

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (3)$$

whose solution is the unique solution of the following Poisson equation with Dirichlet boundary conditions:

$$\Delta f = \operatorname{div} \mathbf{v} \text{ over } \Omega, \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (4)$$

where $\operatorname{div} \mathbf{v} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y}$ is the divergence of $\mathbf{v} = (u, v)$. This is the

$O(N^2)$

Filtering / Convolution

$$(f * h)(x, y) = \sum_{i=-k}^k \sum_{j=-k}^k f(x - i, y - j) \times h(i, j)$$

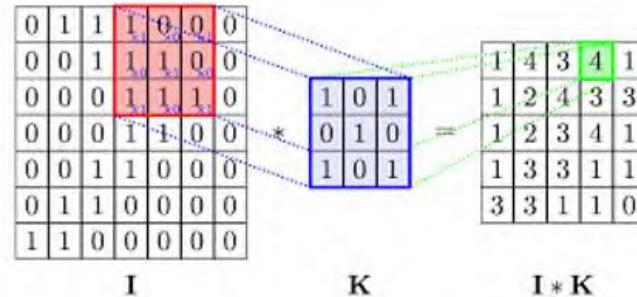


Figure 1: Convolution Algorithm.

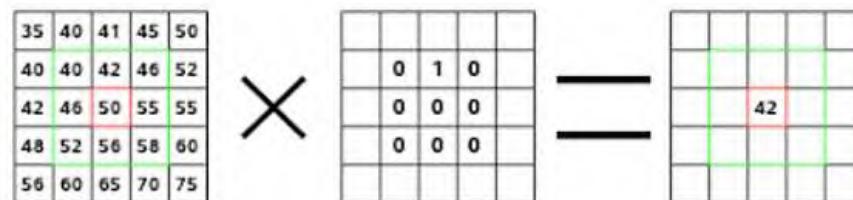
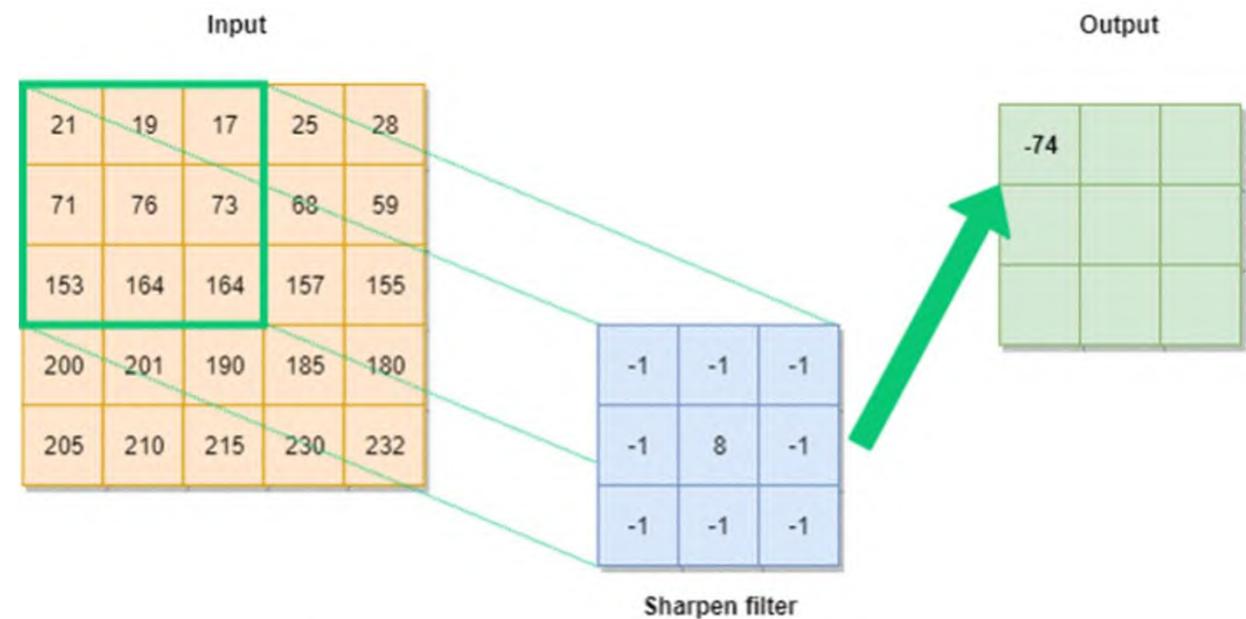
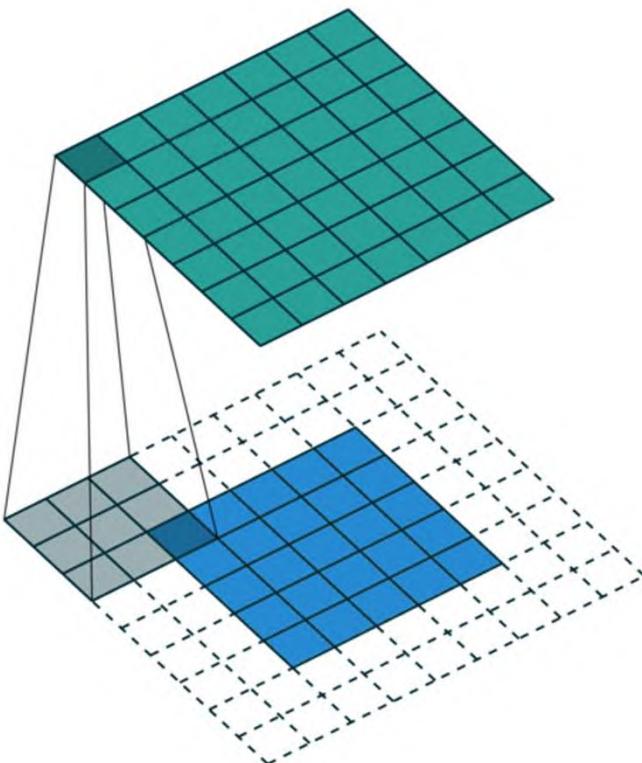


Figure 2: Example of convolution

Example of adjusted convolution mask:

Box blur: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \Rightarrow \text{at the edge } \frac{1}{4} \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$

Image Filtering / Convolution



2D Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$h[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

2D Box Filter

$$h[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Gaussian Blur

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

0.075	0.124	0.075
0.124	0.204	0.124
0.075	0.124	0.075



Sharpen

$$\begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$



XY-Gradient

X – Direction Kernel

-1	0	1
-2	0	2
-1	0	1

Y – Direction Kernel

-1	-2	-1
0	0	0
1	2	1



Laplacian

0	-1	0
-1	4	-1
0	-1	0

-1	-1	-1
-1	8	-1
-1	-1	-1



$$\begin{aligned} & \left(f(x+1, y) - f(x, y) \right) - \left(f(x, y) - f(x-1, y) \right) \\ + & \left(f(x, y+1) - f(x, y) \right) - \left(f(x, y) - f(x, y-1) \right) \end{aligned}$$

Mixture of Filters

Canny Edge Detection is a popular edge detection algorithm. It was developed by John F. Canny in

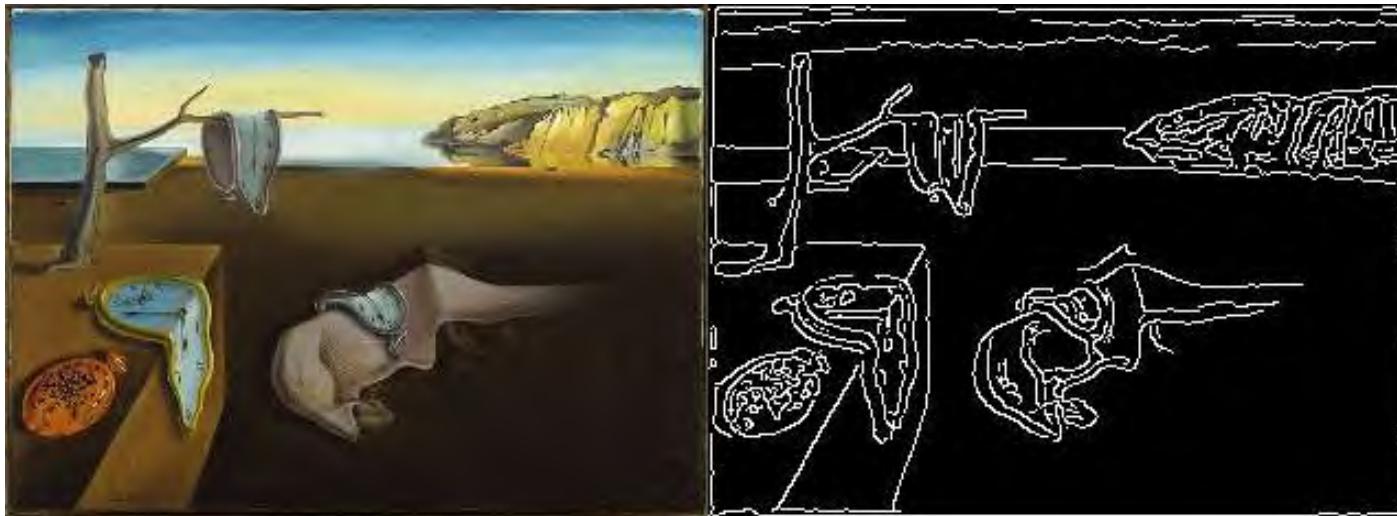
1. It is a multi-stage algorithm and we will go through each stages.
2. **Noise Reduction**

Since edge detection is susceptible to noise in the image, first step is to remove the noise in the image with a 5x5 Gaussian filter. We have already seen this in previous chapters.

3. **Finding Intensity Gradient of the Image**

Smoothed image is then filtered with a Sobel kernel in both horizontal and vertical direction to get first derivative in horizontal direction (G_x) and vertical direction (G_y). From these two images, we can find edge gradient and direction for each pixel as follows:

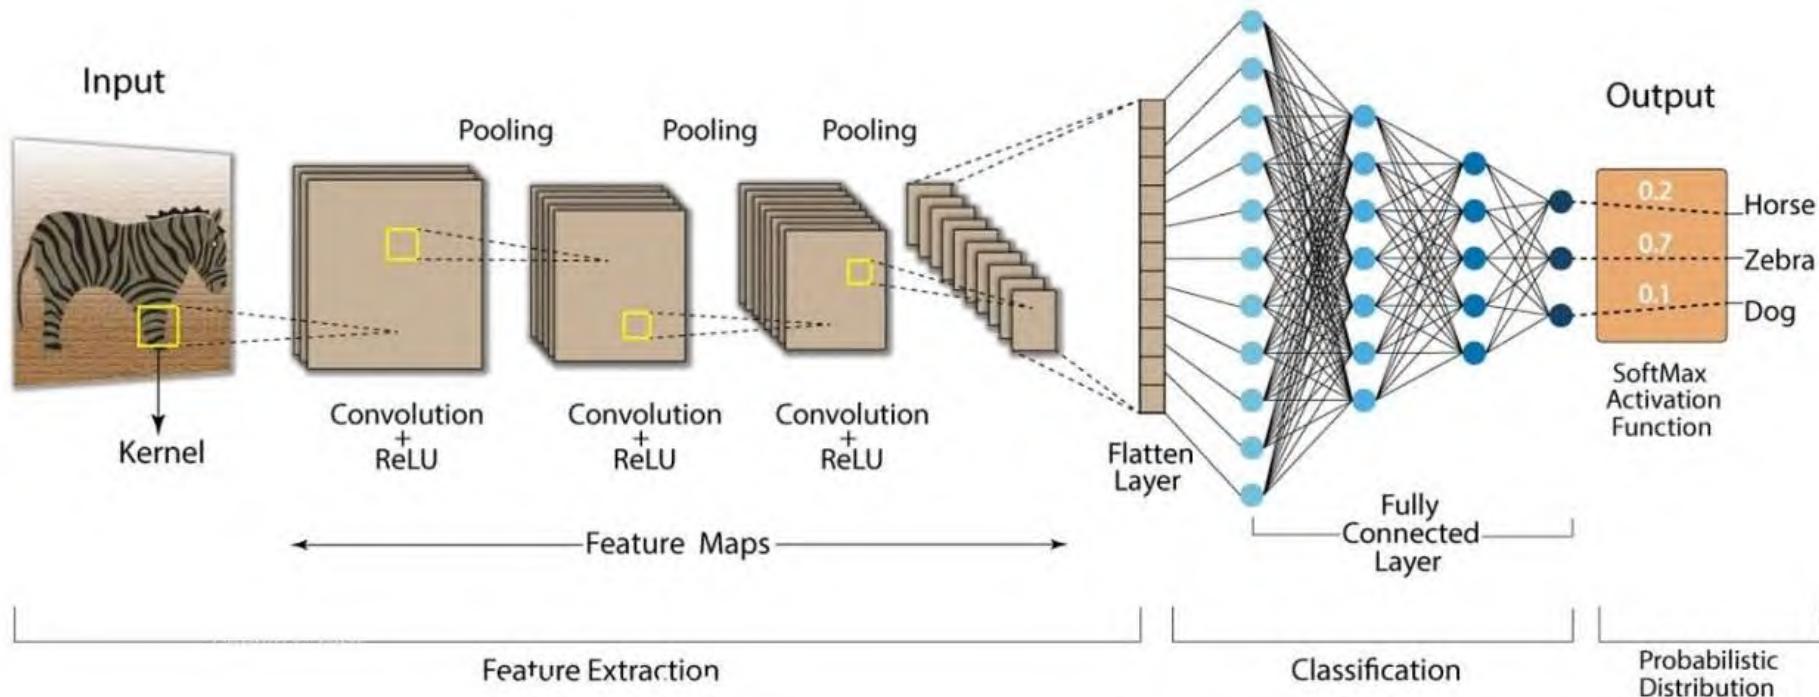
$$\text{Edge_Gradient } (G) = \sqrt{G_x^2 + G_y^2} \text{Angle } (\theta) = \tan^{-1} \left(\frac{G_y}{G_x} \right)$$



Mixture of Learnable Filters

h_{00}	h_{01}	h_{02}
h_{10}	h_{11}	h_{12}
h_{20}	h_{21}	h_{22}

Convolution Neural Network (CNN)



Efficiently Conv—With Fourier Transform

$$F(k, l) = \sum_{i=0}^{N-1} \sum_{j=0}^{N-1} f(i, j) e^{-2\pi i \left(\frac{ki}{N} + \frac{lj}{N} \right)}$$

Fast Fourier transform

Article Talk

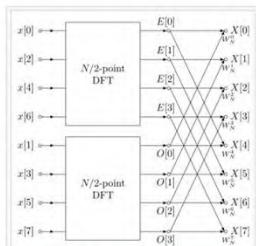
29 languages ▾

Read Edit View history Tools ▾

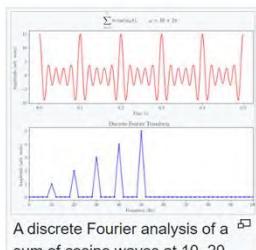
From Wikipedia, the free encyclopedia

"FFT" redirects here. For other uses, see [FFT \(disambiguation\)](#).

A **fast Fourier transform (FFT)** is an algorithm that computes the [Discrete Fourier Transform \(DFT\)](#) of a sequence, or its inverse ([IDFT](#)). Fourier analysis converts a signal from its original domain (often time or space) to a representation in the [frequency domain](#) and vice versa. The DFT is obtained by decomposing a [sequence](#) of values into components of different frequencies.^[1] This operation is useful in many fields, but computing it directly from the definition is often too slow to be practical. An FFT rapidly computes such transformations by [factorizing](#) the [DFT matrix](#) into a product of [sparse](#) (mostly zero) factors.^[2] As a result, it manages to reduce the [complexity](#) of computing the DFT from $O(n^2)$, which arises if one simply applies the definition of DFT, to $O(n \log n)$, where n is the data size. The difference in speed can be enormous, especially for long data sets where n may be in the thousands or millions. In the presence of [round-off error](#), many FFT algorithms are much more accurate than evaluating the DFT definition directly or indirectly. There are many different FFT algorithms based on a wide range of published theories, from simple [complex-number arithmetic](#) to [group theory](#) and [number theory](#).



An example FFT algorithm structure, using a decomposition into half-size FFTs



A discrete Fourier analysis of a sum of cosine waves at 10-20



Frequency Magnitude
Center at (0,0)

Convolution Theorem

In mathematics, the **convolution theorem** states that under suitable conditions the Fourier transform of a convolution of two functions (or signals) is the product of their Fourier transforms. More generally, convolution in one domain (e.g., [time domain](#)) equals point-wise multiplication in the other domain (e.g., [frequency domain](#)). Other versions of the convolution theorem are applicable to various [Fourier-related transforms](#).

Convolution theorem

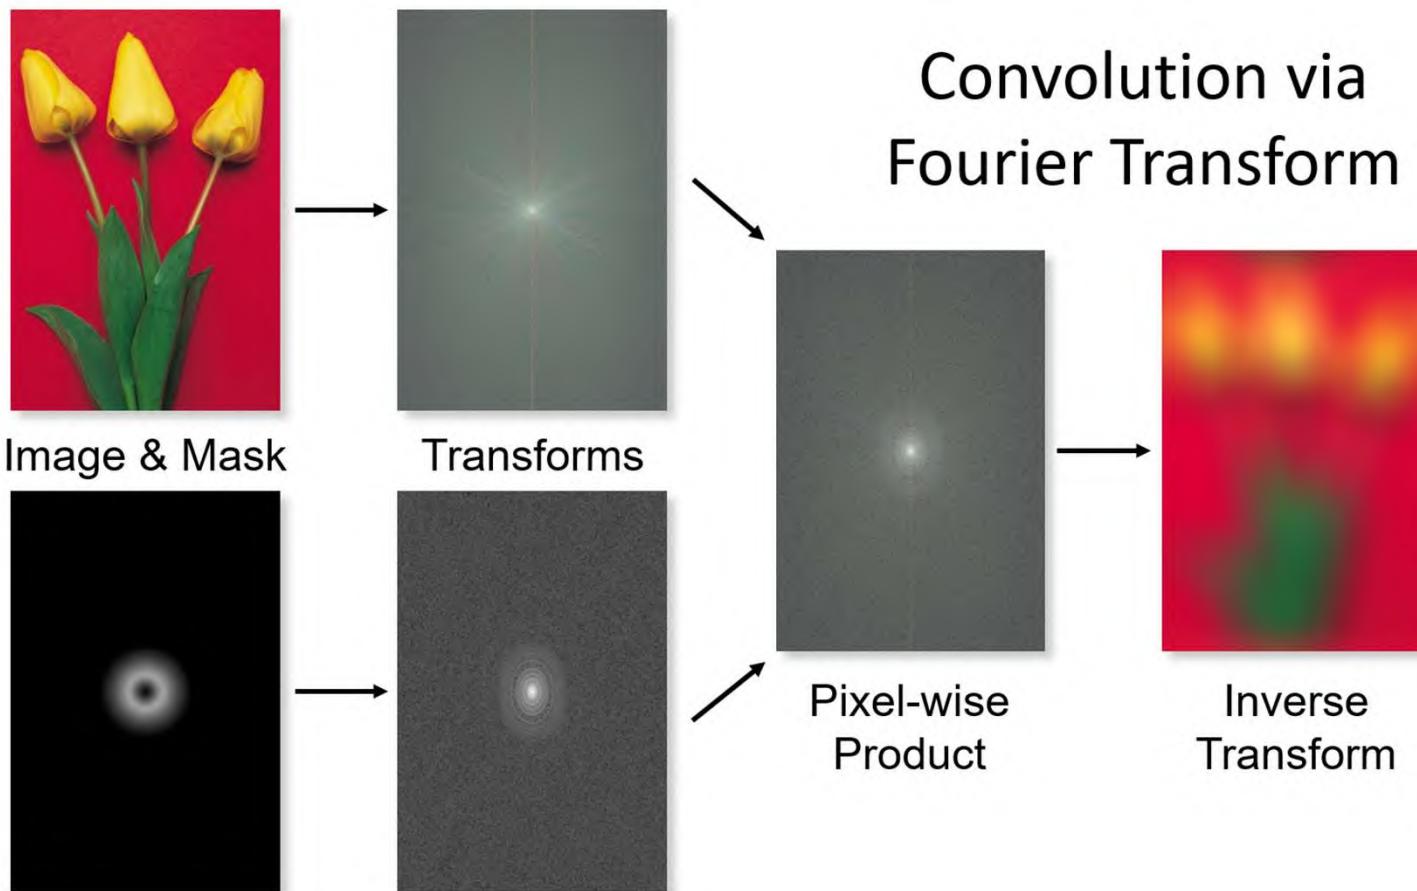
$$r(x) = \{u * v\}(x) = \mathcal{F}^{-1}\{U \cdot V\}. \quad (\text{Eq.1b})$$

$$F(k, l) = \sum_{i=0}^{N-1} \sum_{j=0}^{N-1} f(i, j) e^{-2\pi i \left(\frac{ki}{N} + \frac{lj}{N} \right)}$$

$$f(i, j) = \frac{1}{N^2} \sum_{k=0}^{N-1} \sum_{l=0}^{N-1} F(k, l) e^{2\pi i \left(\frac{ki}{N} + \frac{lj}{N} \right)}$$

$$f(i, j) = \sum_{k=0}^{N-1} \sum_{l=0}^{N-1} \left(\sum_{m=0}^{N-1} \sum_{n=0}^{N-1} f(m, n) e^{-2\pi i \left(\frac{km}{N} + \frac{ln}{N} \right)} \right) e^{2\pi i \left(\frac{ki}{N} + \frac{lj}{N} \right)}$$

Fast Conv Computation



So called filtering

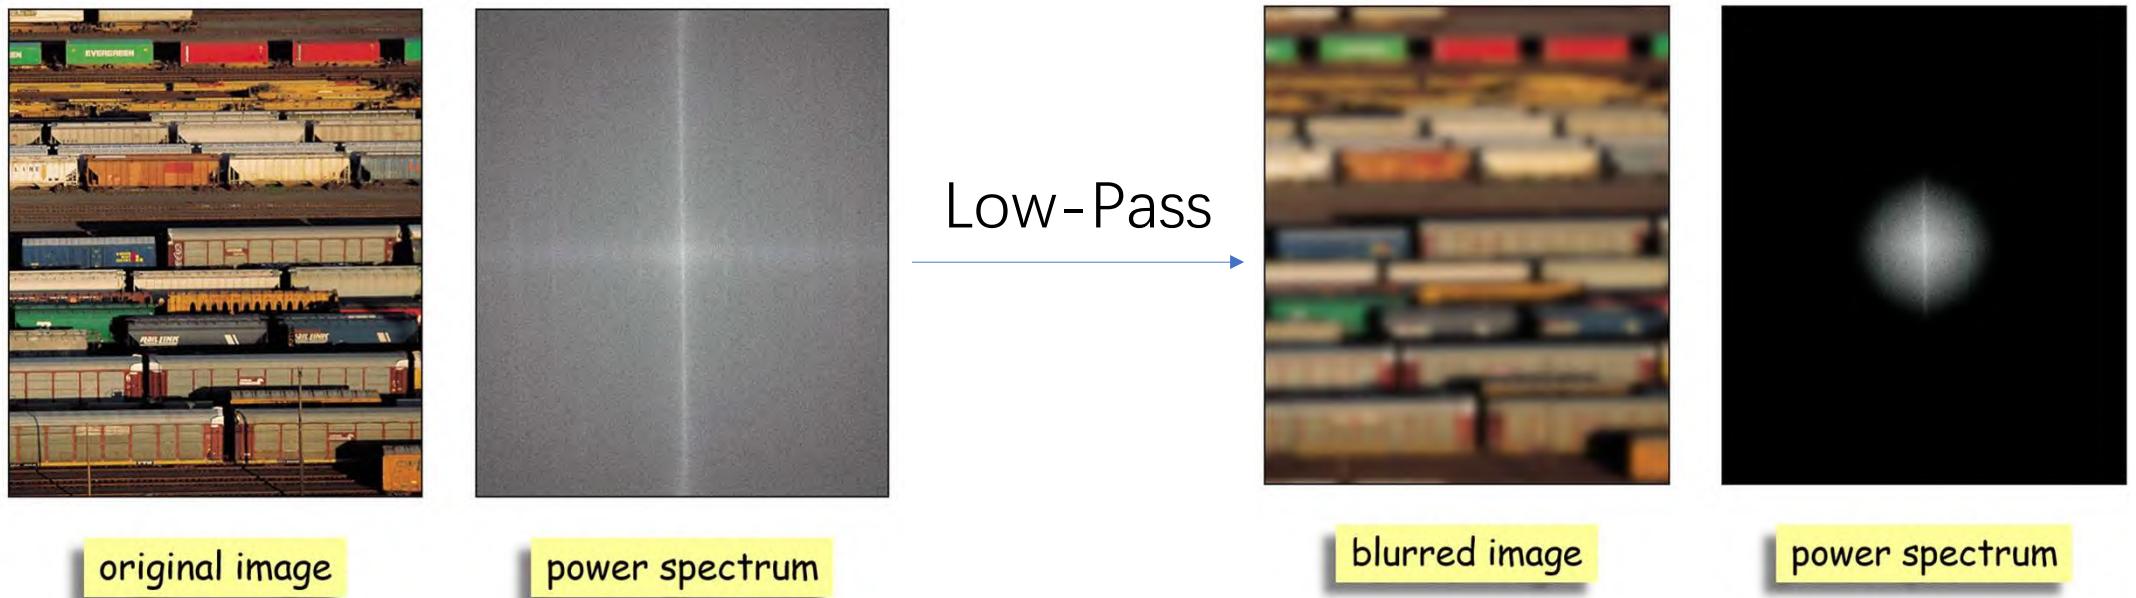


original image

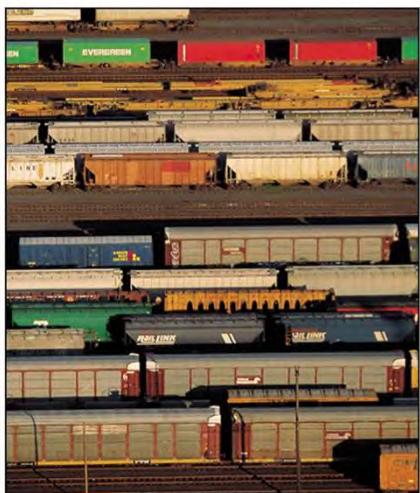


power spectrum

So called filtering



So called filtering

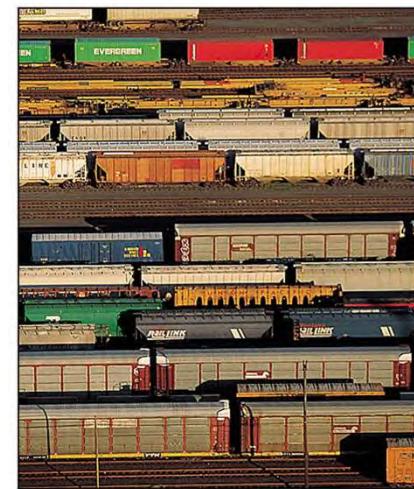


original image

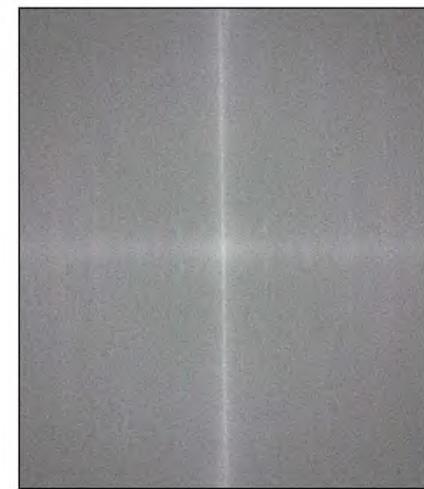


power spectrum

More High

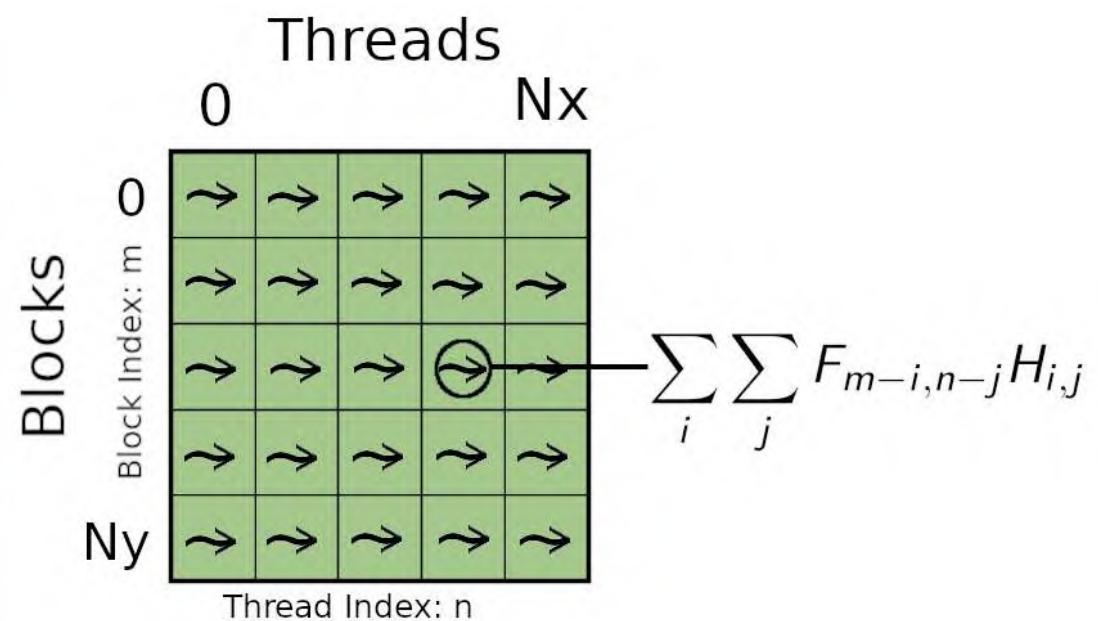
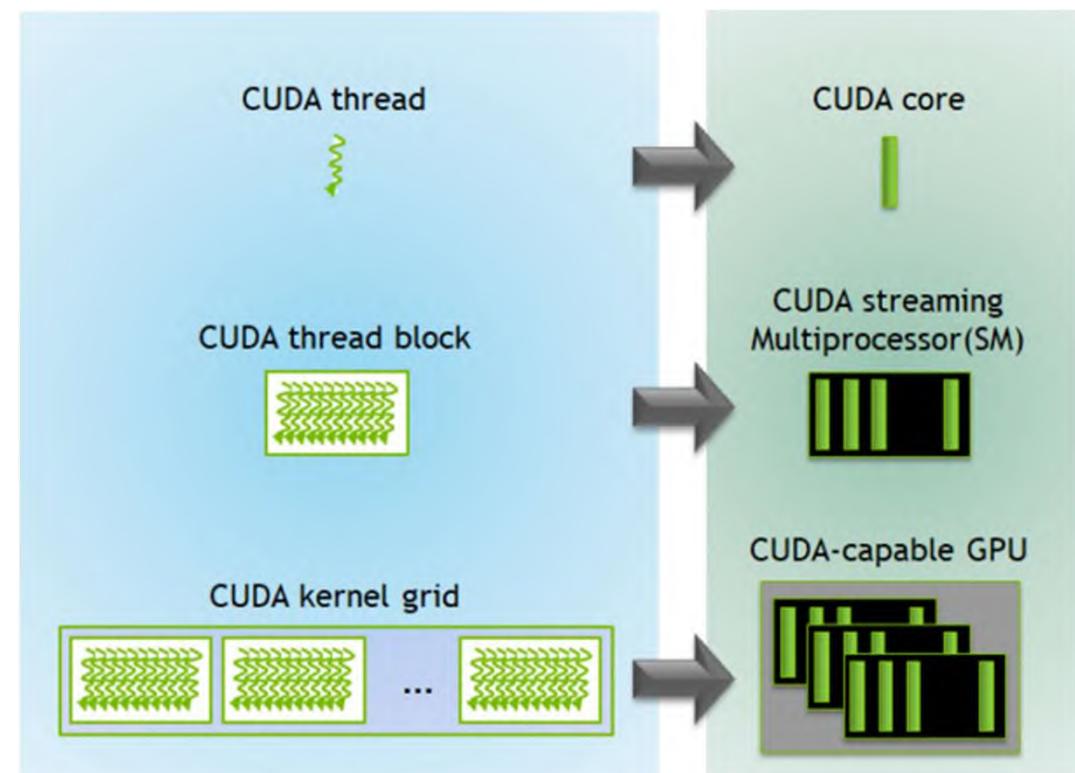


sharpened image



power spectrum

But not used in Modern GPU



Easy to Use

```
1 import cv2
2 import numpy as np
3
4 image = cv2.imread('test.jpg')
5
6 # Print error message if image is null
7 if image is None:
8     print('Could not read image')
9
10 # Apply identity kernel
11 kernel1 = np.array([[0, 0, 0],
12                     [0, 1, 0],
13                     [0, 0, 0]])
14
15 identity = cv2.filter2D(src=image, ddepth=-1, kernel=kernel1)
16
17 cv2.imshow('Original', image)
18 cv2.imshow('Identity', identity)
19
20 cv2.waitKey()
21 cv2.imwrite('identity.jpg', identity)
22 cv2.destroyAllWindows()
23
24 # Apply blurring kernel
25 kernel2 = np.ones((5, 5), np.float32) / 25
26 img = cv2.filter2D(src=image, ddepth=-1, kernel=kernel2)
27
28 cv2.imshow('Original', image)
29 cv2.imshow('Kernel Blur', img)
30
31 cv2.waitKey()
32 cv2.imwrite('blur_kernel.jpg', img)
33 cv2.destroyAllWindows()
```

torch.nn.functional.conv2d

```
torch.nn.functional.conv2d(input, weight, bias=None, stride=1, padding=0,
dilation=1, groups=1) → Tensor
```

Applies a 2D convolution over an input image composed of several input planes.

This operator supports `TensorFloat32`.

See `Conv2d` for details and output shape.

* NOTE

In some circumstances when given tensors on a CUDA device and using CuDNN, this operator may select a nondeterministic algorithm to increase performance. If this is undesirable, you can try to make the operation deterministic (potentially at a performance cost) by setting `torch.backends.cudnn.deterministic = True`. See [Reproducibility](#) for more information.

* NOTE

This operator supports complex data types i.e. `complex32`, `complex64`, `complex128`.

Parameters

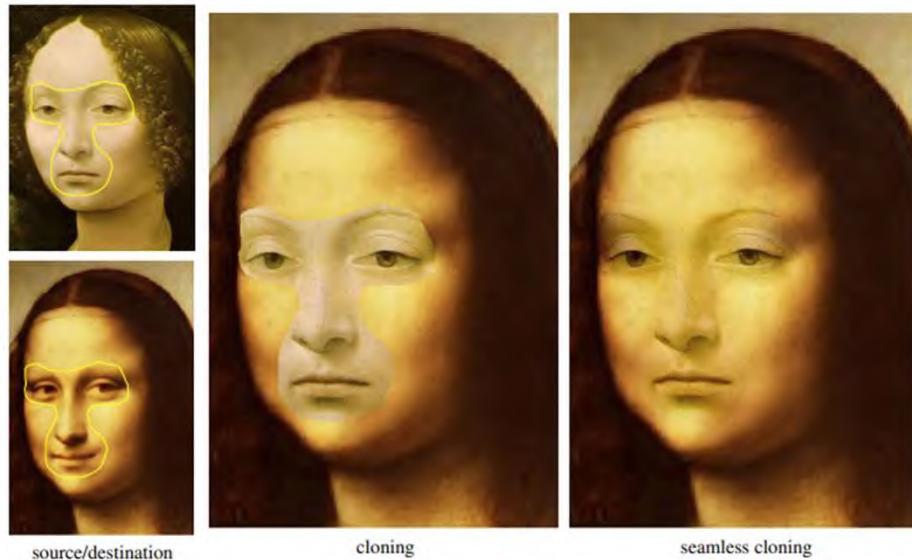
- `input` – input tensor of shape (`minibatch, in_channels, iH, iW`)
- `weight` – filters of shape (`out_channels, $\frac{\text{in_channels}}{\text{groups}}$, kH, kW`)
- `bias` – optional bias tensor of shape (`out_channels`). Default: `None`
- `stride` – the stride of the convolving kernel. Can be a single number or a tuple (`sH, sW`). Default: 1
- `padding` – implicit paddings on both sides of the input. Can be a string {'valid', 'same'}, single number or a tuple (`padH, padW`). Default: 0 `padding='valid'` is the same as no padding. `padding='same'` pads the input so the output has the same shape as the input. However, this mode doesn't support any stride values other than 1.

* WARNING

For `padding='same'`, if the `weight` is even-length and `dilation` is odd in any dimension, a full `pad()` operation may be needed internally. Lowering performance.

- `dilation` – the spacing between kernel elements. Can be a single number or a tuple (`(dH, dW)`). Default: 1
- `groups` – split input into groups, both `in_channels` and `out_channels` should be divisible by the number of groups. Default: 1

How to implement PIE



0	-1	0
-1	4	-1
0	-1	0

A *guidance field* is a vector field \mathbf{v} used in an extended version of the minimization problem (1) above:

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (3)$$

whose solution is the unique solution of the following Poisson equation with Dirichlet boundary conditions:

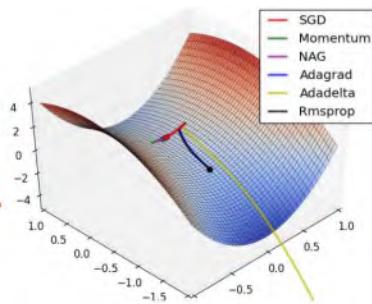
$$\Delta f = \operatorname{div} \mathbf{v} \text{ over } \Omega, \text{ with } f|_{\partial\Omega} = f^*|_{\partial\Omega}, \quad (4)$$

where $\operatorname{div} \mathbf{v} = \frac{\partial u}{\partial x} + \frac{\partial v}{\partial y}$ is the divergence of $\mathbf{v} = (u, v)$. This is the

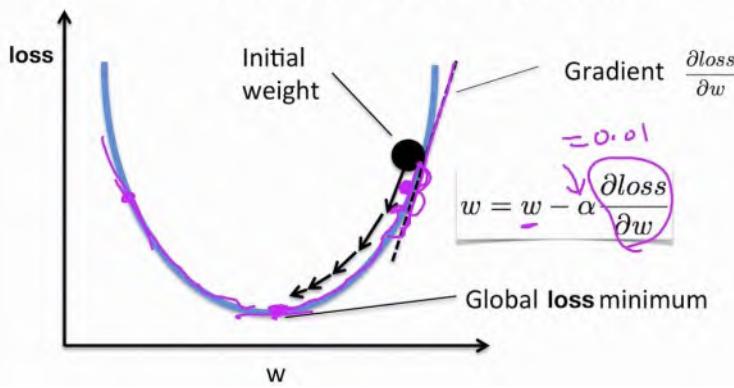
`torch.nn.functional.conv2d(I, K)` or

`I[1:-1, 1:-1]*4 - I[1:-1, :-2] - I[1:-1, 2:] - I[:-2, 1:-1] - I[2:, 1:-1]`

Optimize using Gradient Descent



Gradient descent algorithm



```
import torch
import torch.optim as optim

# 构造数据，假设我们的目标函数是  $y = 2x^2 + 3x + 1$ 
x = torch.tensor([1.0, 2.0, 3.0, 4.0, 5.0])
y_true = 2 * x * x + 3 * x + 1

# 初始化需要优化的参数 a, b, c，设置为需要梯度的变量
a = torch.tensor(0.0, requires_grad=True)
b = torch.tensor(0.0, requires_grad=True)
c = torch.tensor(0.0, requires_grad=True)

# 定义优化器，使用随机梯度下降
optimizer = optim.SGD([a, b, c], lr=0.01)

# 开始训练
for epoch in range(1000):
    # 预测  $y = ax^2 + bx + c$ 
    y_pred = a * x * x + b * x + c

    # 计算均方误差损失
    loss = torch.mean((y_pred - y_true) ** 2)

    # 反向传播计算梯度
    optimizer.zero_grad()
    loss.backward()

    # 更新参数
    optimizer.step()

    # 每 100 个 epoch 打印一次损失和参数
    if (epoch + 1) % 100 == 0:
        print(f'Epoch [{epoch+1}/1000], Loss: {loss.item():.4f}, a: {a.item():.4f}, b: {b.item():.4f}, c: {c.item():.4f}'
```



中国科学技术大学

University of Science and Technology of China

Q & A