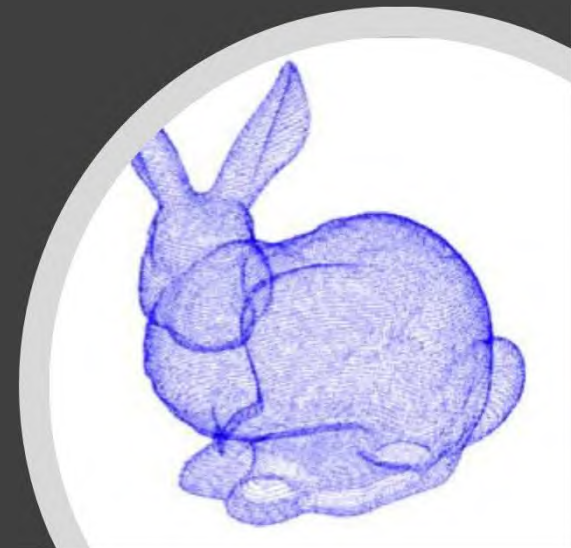
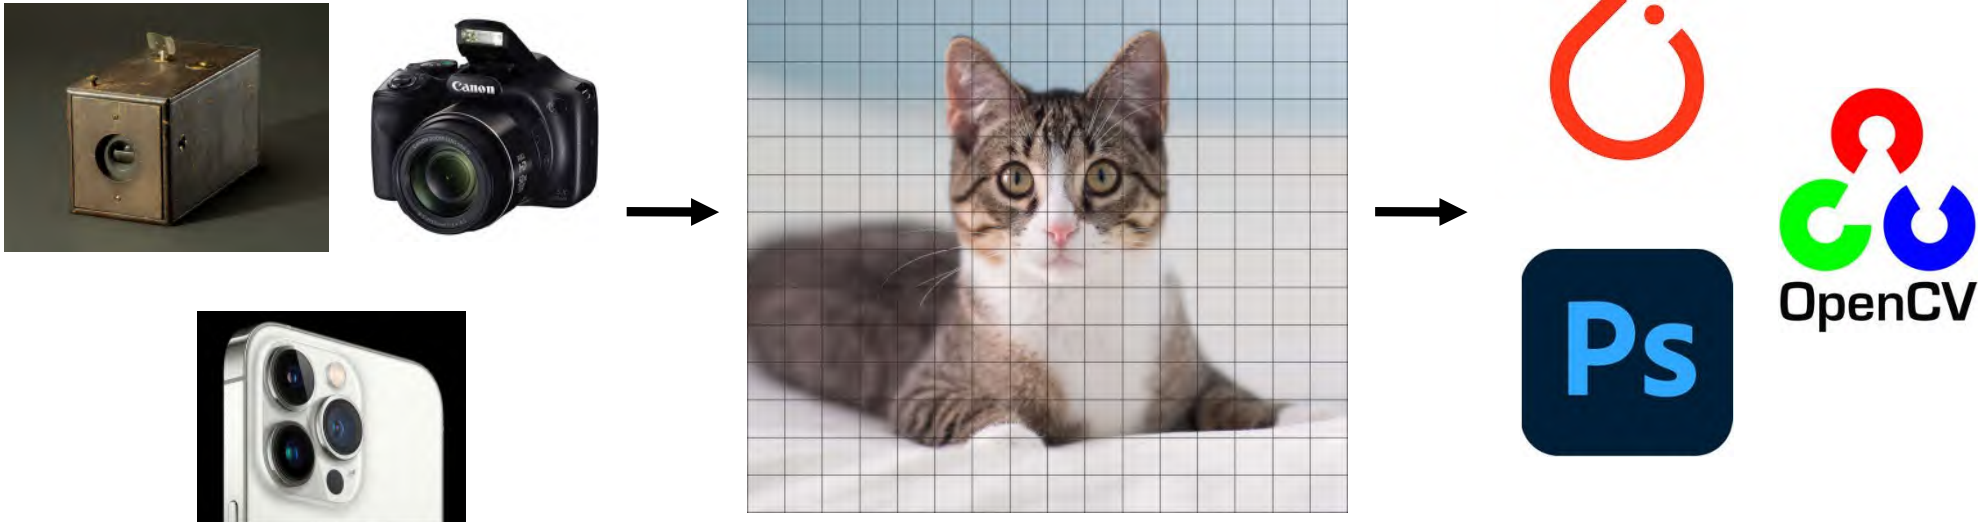


# 三维表示

## 3D Representations

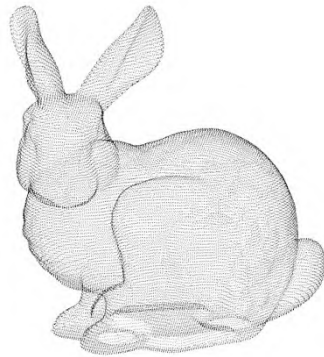
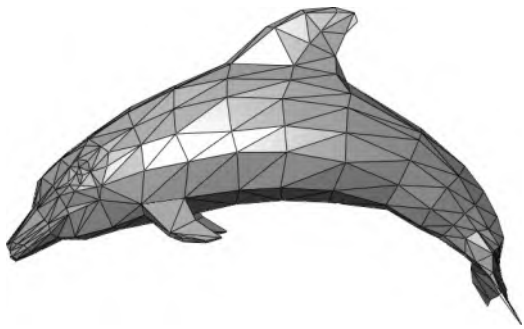


# Representation is simple for 2D-land...



Unified representation across  
applications, tools, and sensors

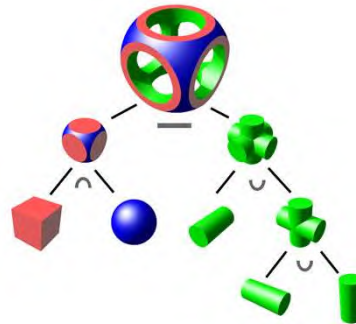
# Meanwhile in 3D-land



$$\{\mathbf{p} \mid f(\mathbf{p}) = 0\}$$



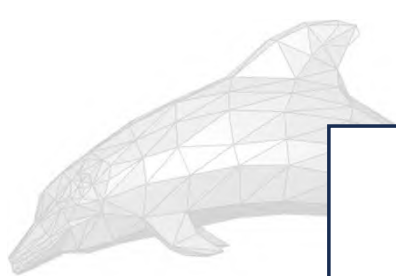
$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3$$



and many more ...



# Meanwhile in 3D-land



What is the **best** representation?  
It **depends** on applications, input, requirements

Reconstruction / Editing / Animation

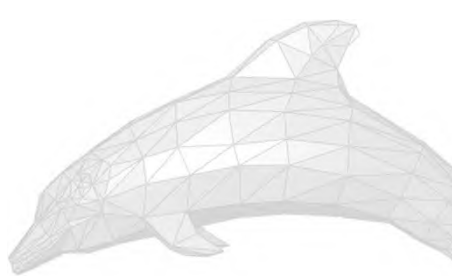


... more ...

$$\{\mathbf{p} \mid f(\mathbf{p}) = 0\}$$

$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3$$

# Meanwhile in 3D-land



Goals:

Understand different representations  
(and basic implementation)

Limitations and Benefits

many more ...



$$\{\mathbf{p} \mid f(\mathbf{p}) = 0\}$$



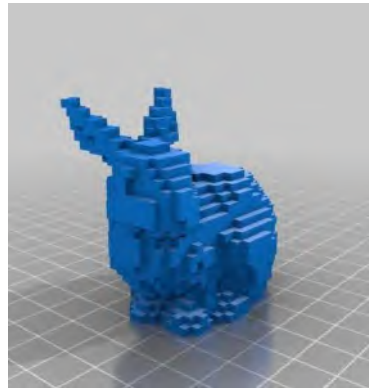
$$f(\mathbf{u}) = \mathbf{p} \in \mathbb{R}^3$$



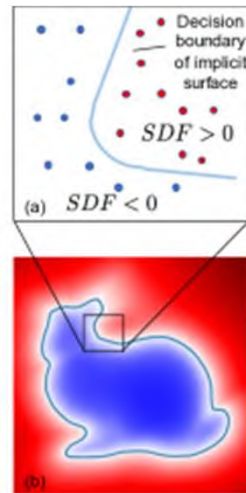
# Basic 3D representations



2.5D / Image Based Rendering

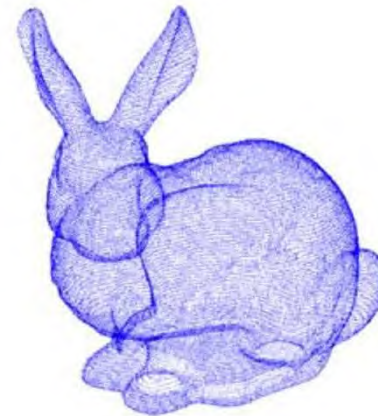


Explicit

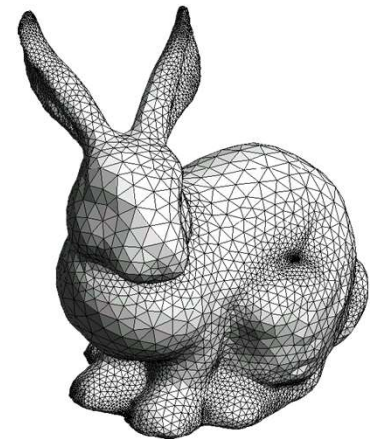


Volumetrics

Implicit



Point clouds

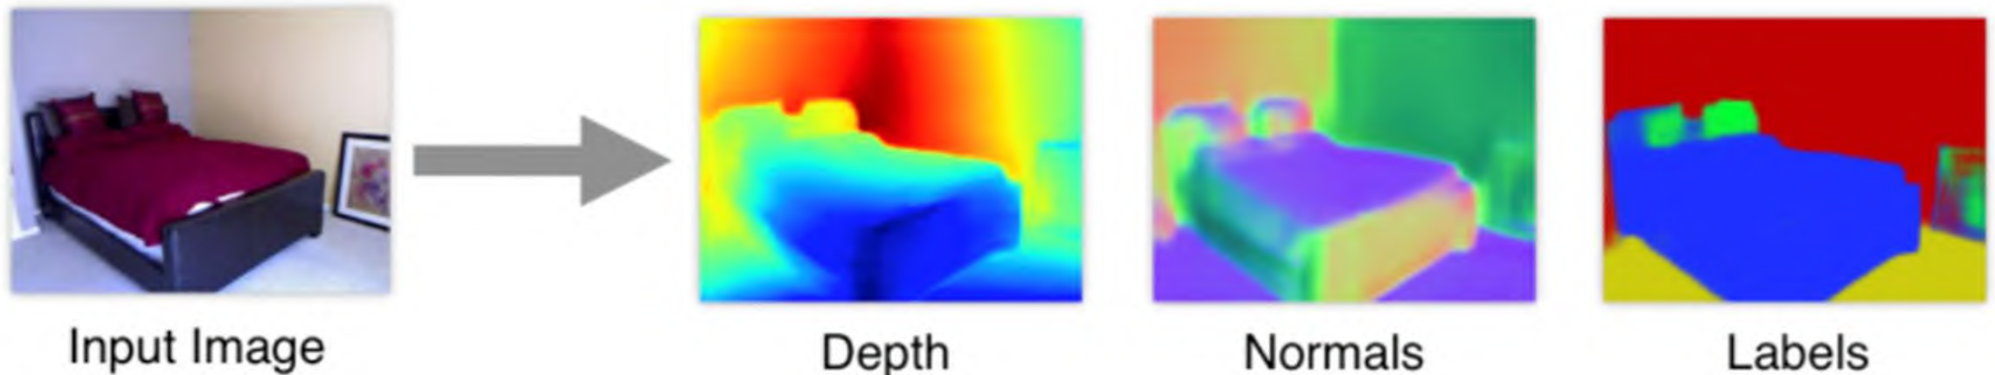


Meshes

## 2.5D — ex. Depth Maps



## 2.5D Representations

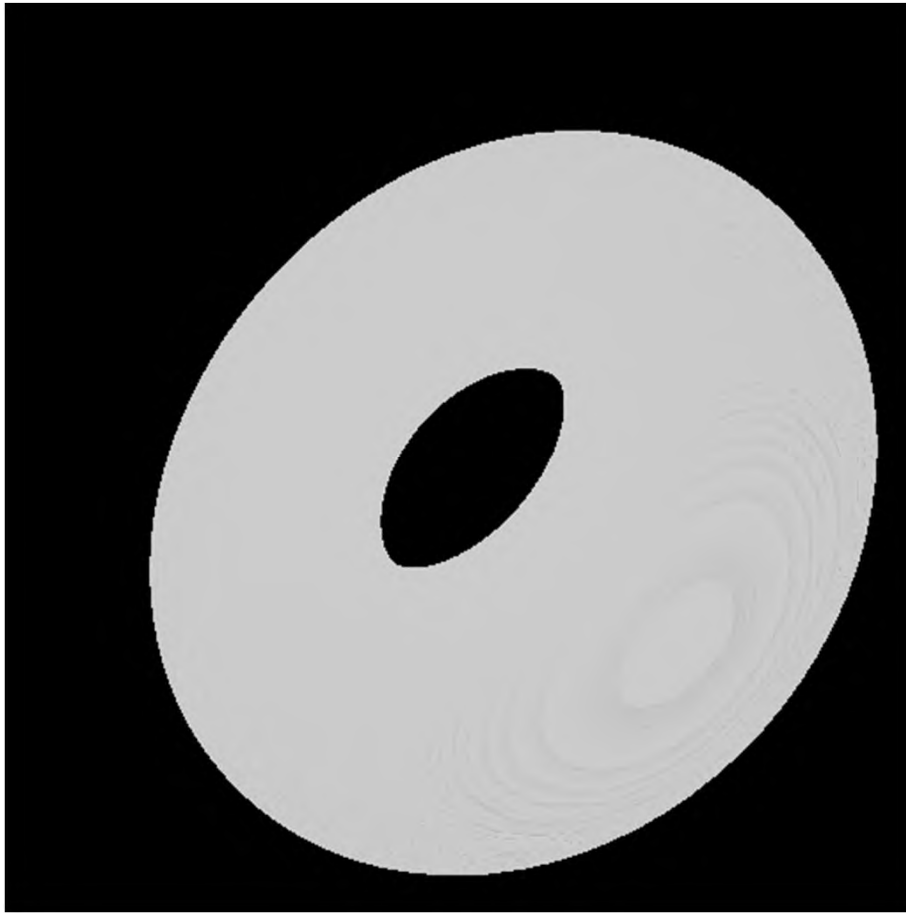


Eigen and Fergus ICCV '15

- Depth-map = **Per-pixel** depth or disparity value ( $H \times W \times 1$ )
- Normal map = per-pixel normal values ( $H \times W \times 3$ )
- Output is aligned with the input
- 2.5D = Not suitable to look around (no large baseline)
- Can think of each pixel as a colored 3D point



## 2.5D — Representing the Visible



Does not capture the 'full' 3D structure

Properties associated to the **visible** image pixels in the **input view**

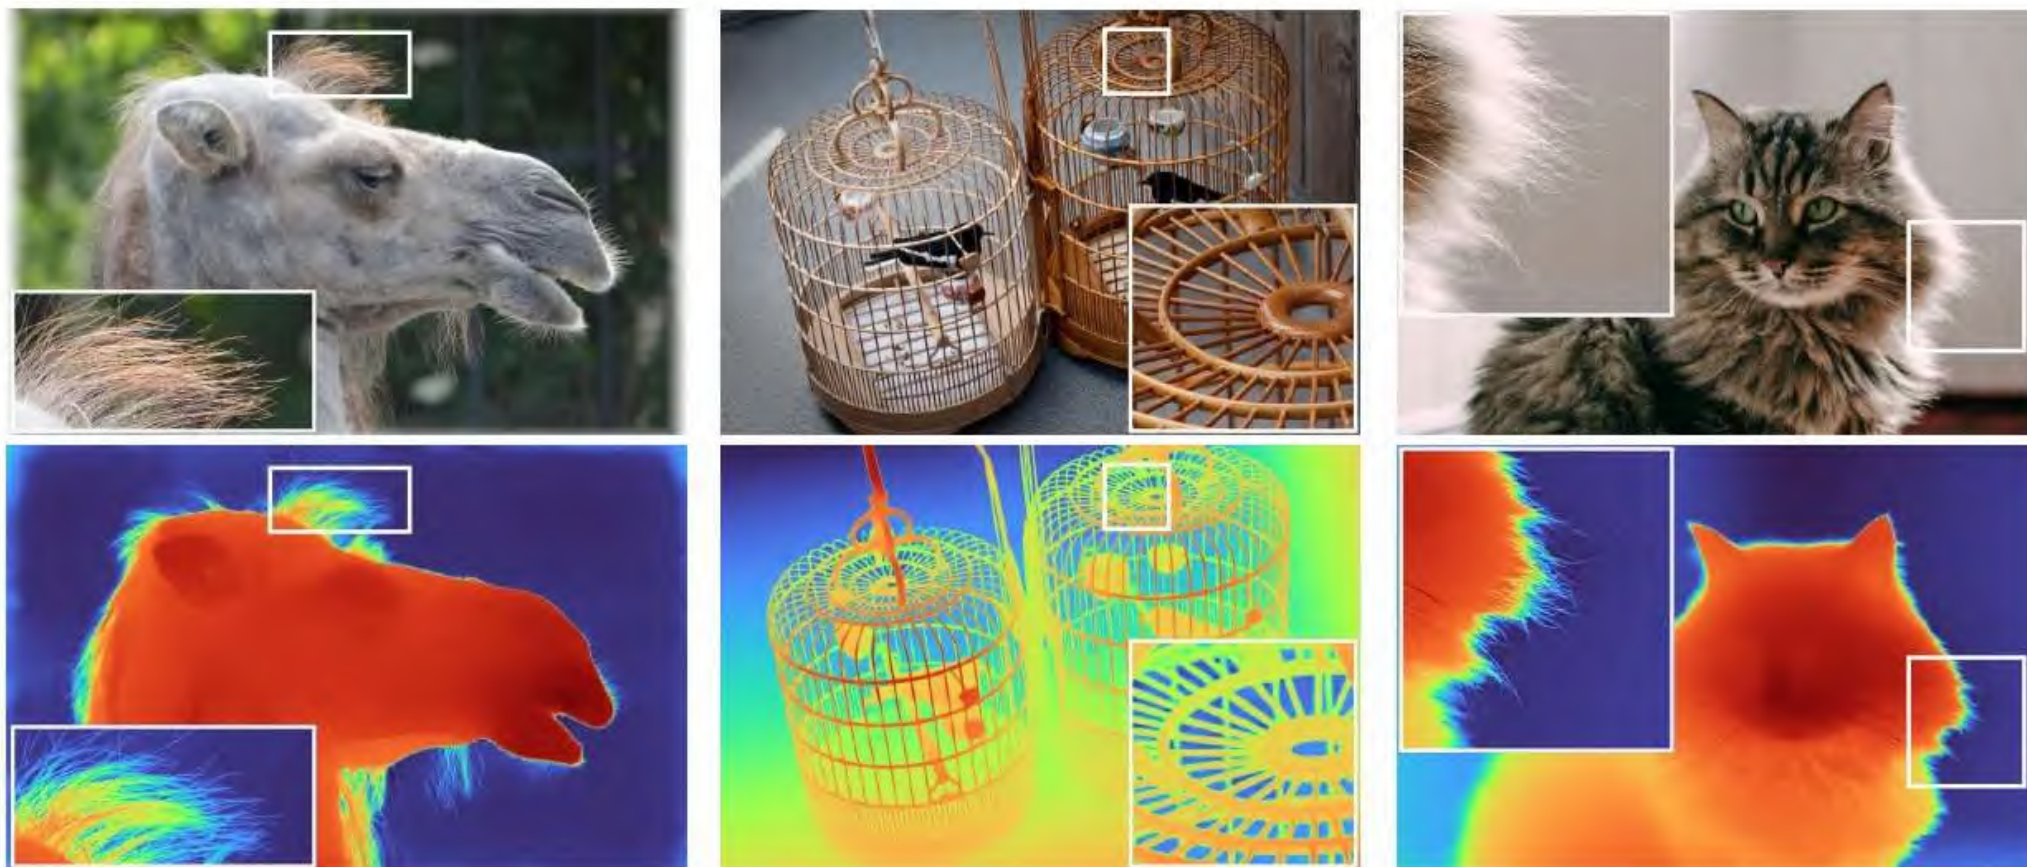
# Easy way to obtain — Monocular Depth

- Predicted from a single image!
- Usually learn from many many image-depth pairs
- Predicts relative depth

$$(s, t) = \arg \min_{s, t} \sum_{i=1}^M (s \mathbf{d}_i + t - \mathbf{d}_i^*)^2$$



# Even for sharp structures



Depth Pro: Sharp Monocular Metric Depth in Less Than a Second. ArXiv 2024.

# How to render?



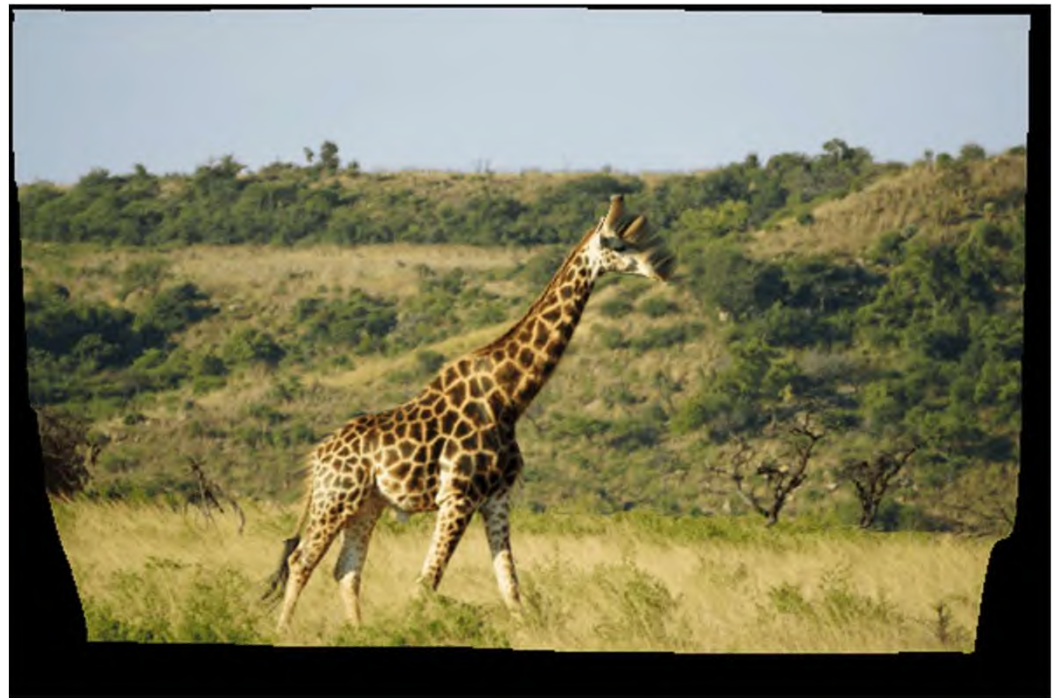


# Treat it as Colored 3D Points

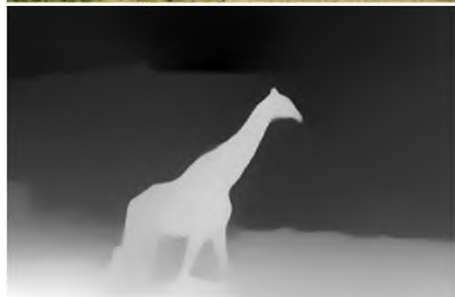


All monocular depth works can't produce perfect depth  
But still useful for some applications

# 3D Photo: View synthesis from single image



# 3D Photo: View synthesis from single image



Synthesize the missing regions!



# 3D Photo via inpainting missing regions



Niklaus et al. ToG 2019



Shih et al. CVPR 2020

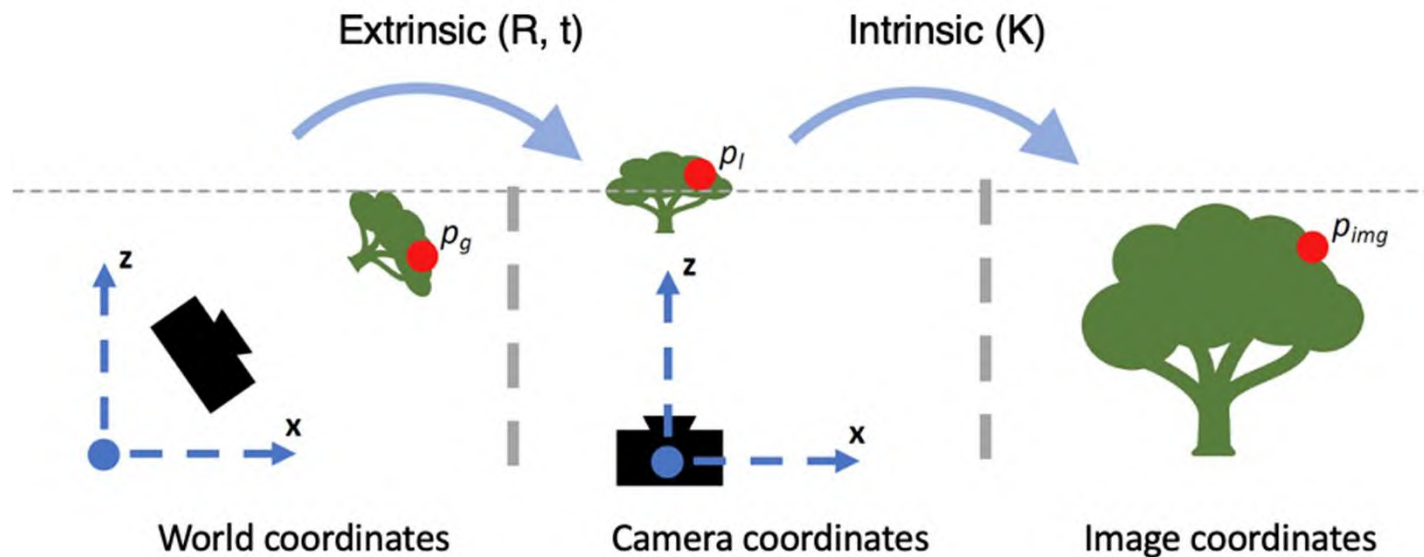


# Inpaint RGB and D & Repeat → Infinite Nature Perpetual View Generation



# Recall: Depth $\rightarrow$ 3D Point

- Lift the 2D plane to 3D
- Inverse the projection process (the last class)

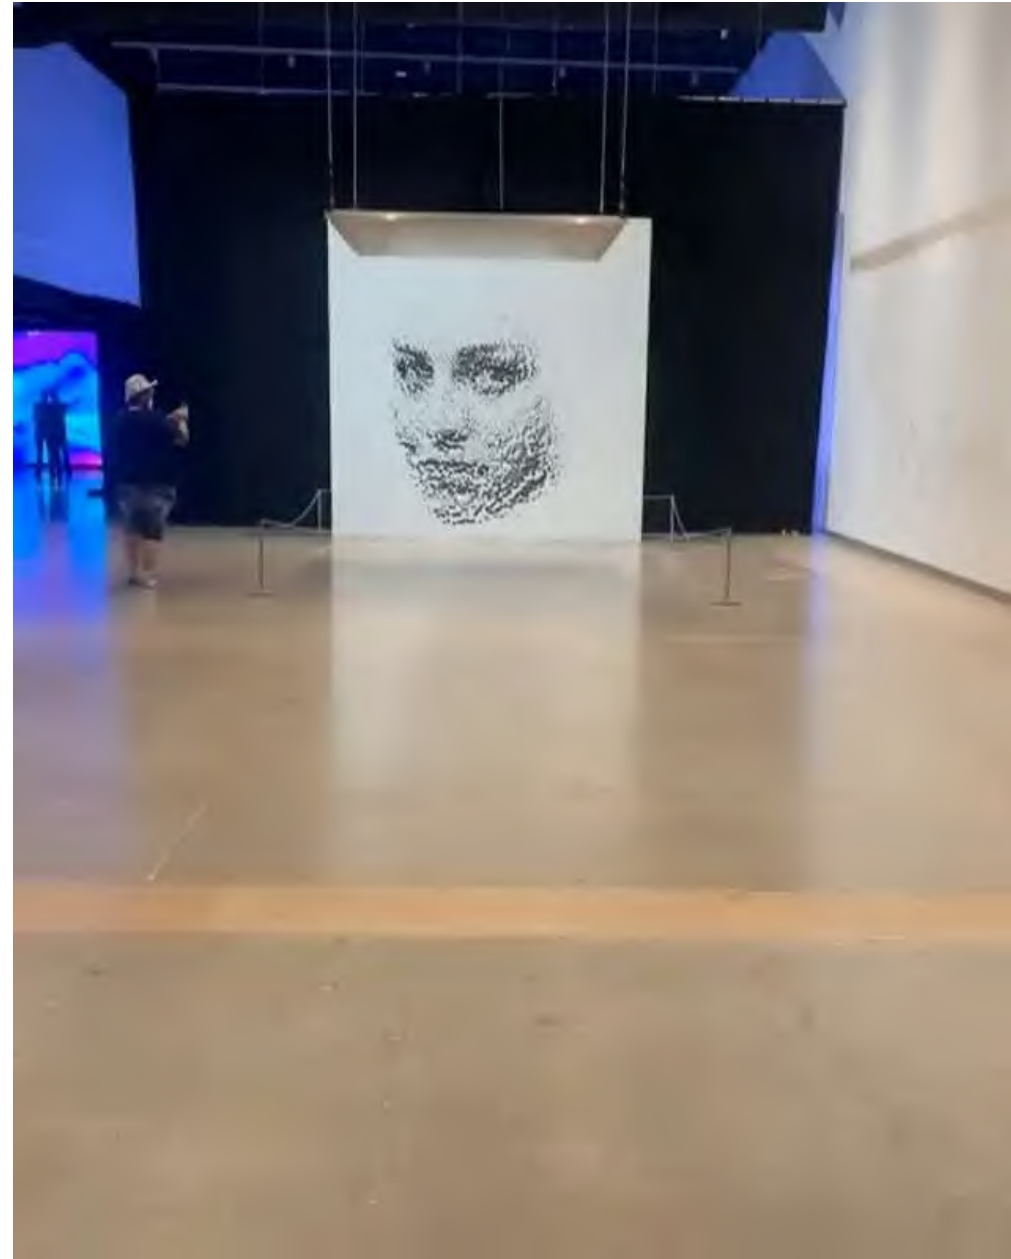


# Point Clouds (PCL)

A basic point cloud =  $\{(x_i, y_i, z_i), i \in [1, n]\}$   
Can have other attributes (color, normal, ... )

Obtained from

- Depth images / Lidar
- Single Image (with deep learning)
- Multiple RGB Images
- Scanner outputs
- Converted from a 3D mesh

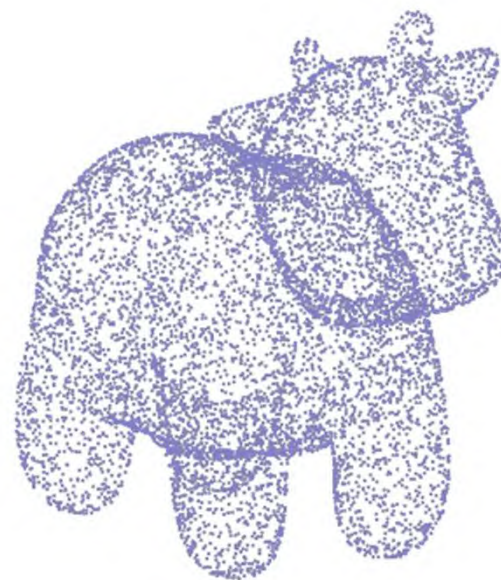


# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \cdots, \mathbf{p}_N\}$$

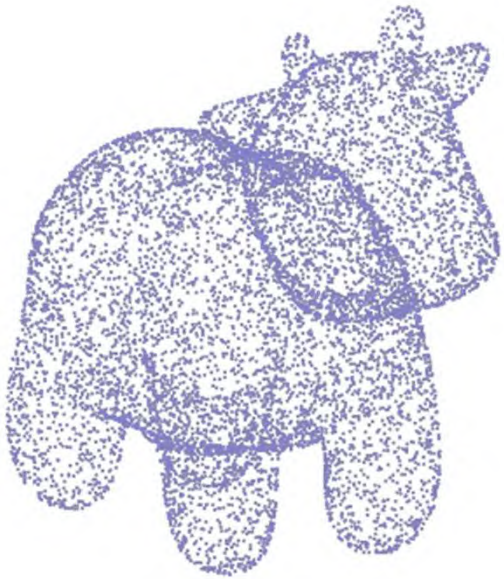
Permutation  $\sigma$



$$\{\mathbf{p}_{\sigma(1)}, \mathbf{p}_{\sigma(2)}, \cdots, \mathbf{p}_{\sigma(N)}\}$$



# Point Clouds



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

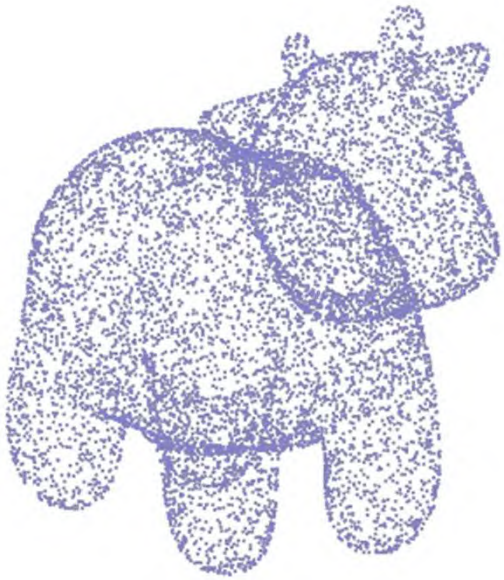
**Unordered set of points**

x1,y1,z1
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.
.

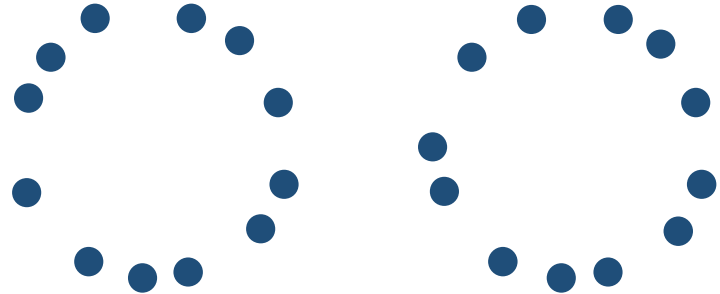
Often represented as a  $N \times 3$  array, but ordering does **not** matter (unlike images)

Need processing/generation methods that are permutation invariant (e.g. fully connected layer will not work)

# Point Clouds



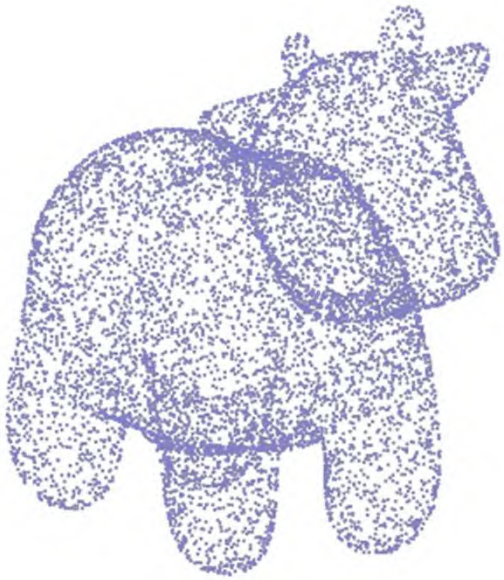
No explicit 'connectivity'  
information



$$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$$

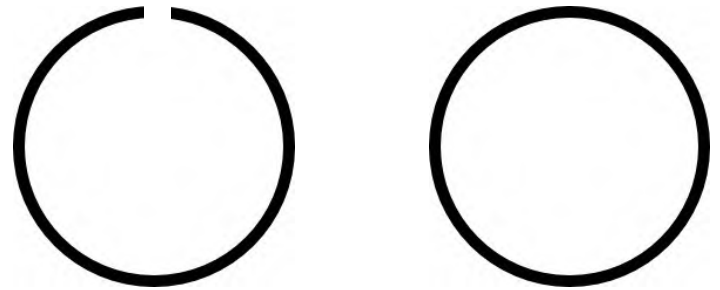
**Unordered set of points**

# Point Clouds



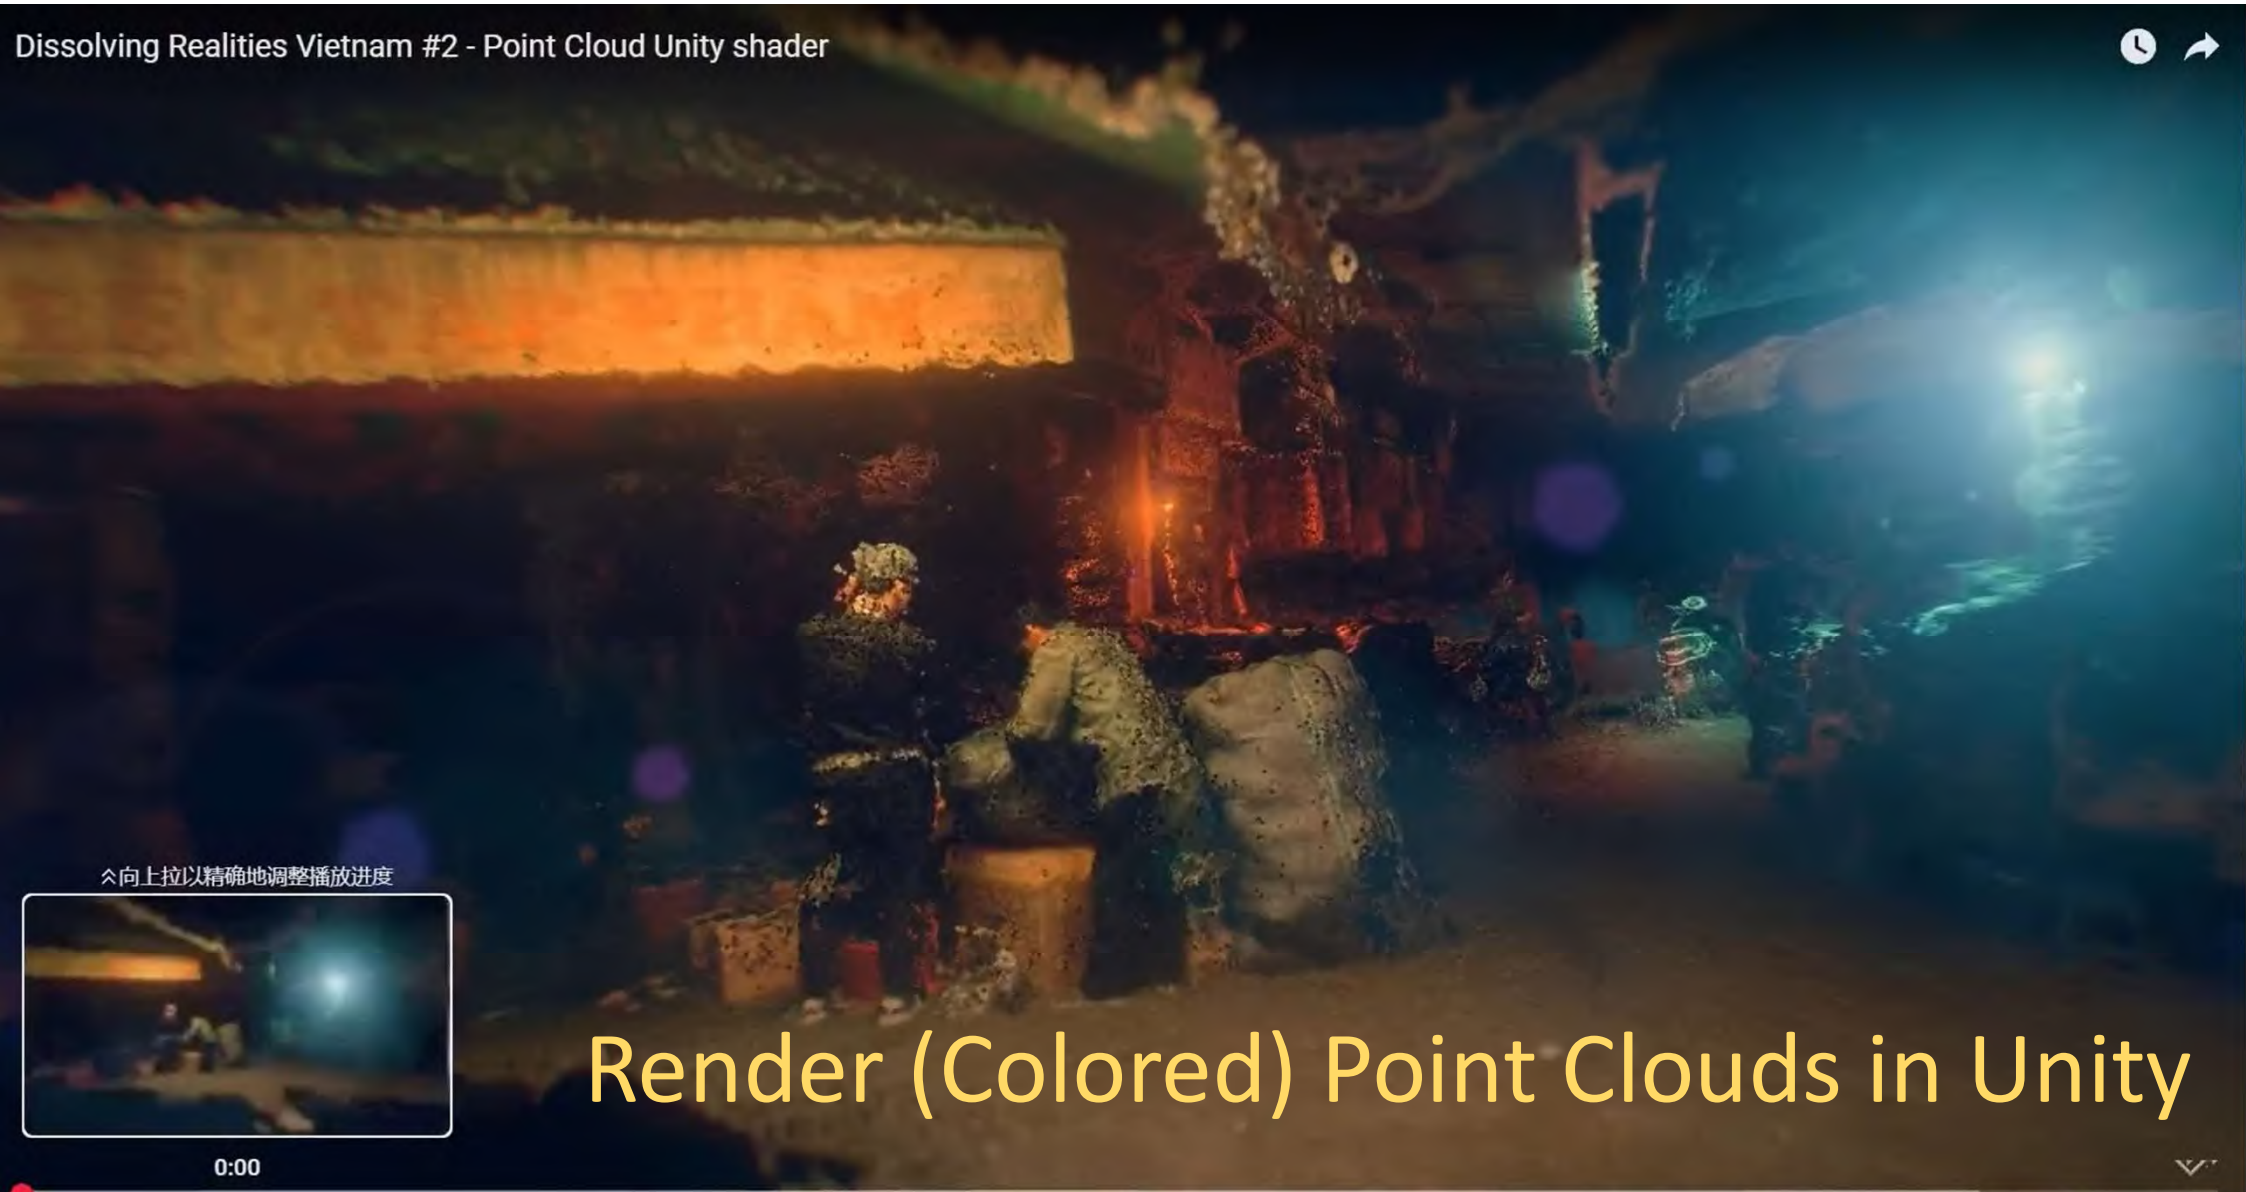
$\{\mathbf{p}_1, \mathbf{p}_2, \dots, \mathbf{p}_N\}$   
**Unordered set of points**

No explicit 'connectivity'  
information



(So it's more efficient to add  
edges (connectivity))

Dissolving Realities Vietnam #2 - Point Cloud Unity shader

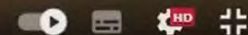


⤴ 向上拉以精确地调整播放进度

Render (Colored) Point Clouds in Unity

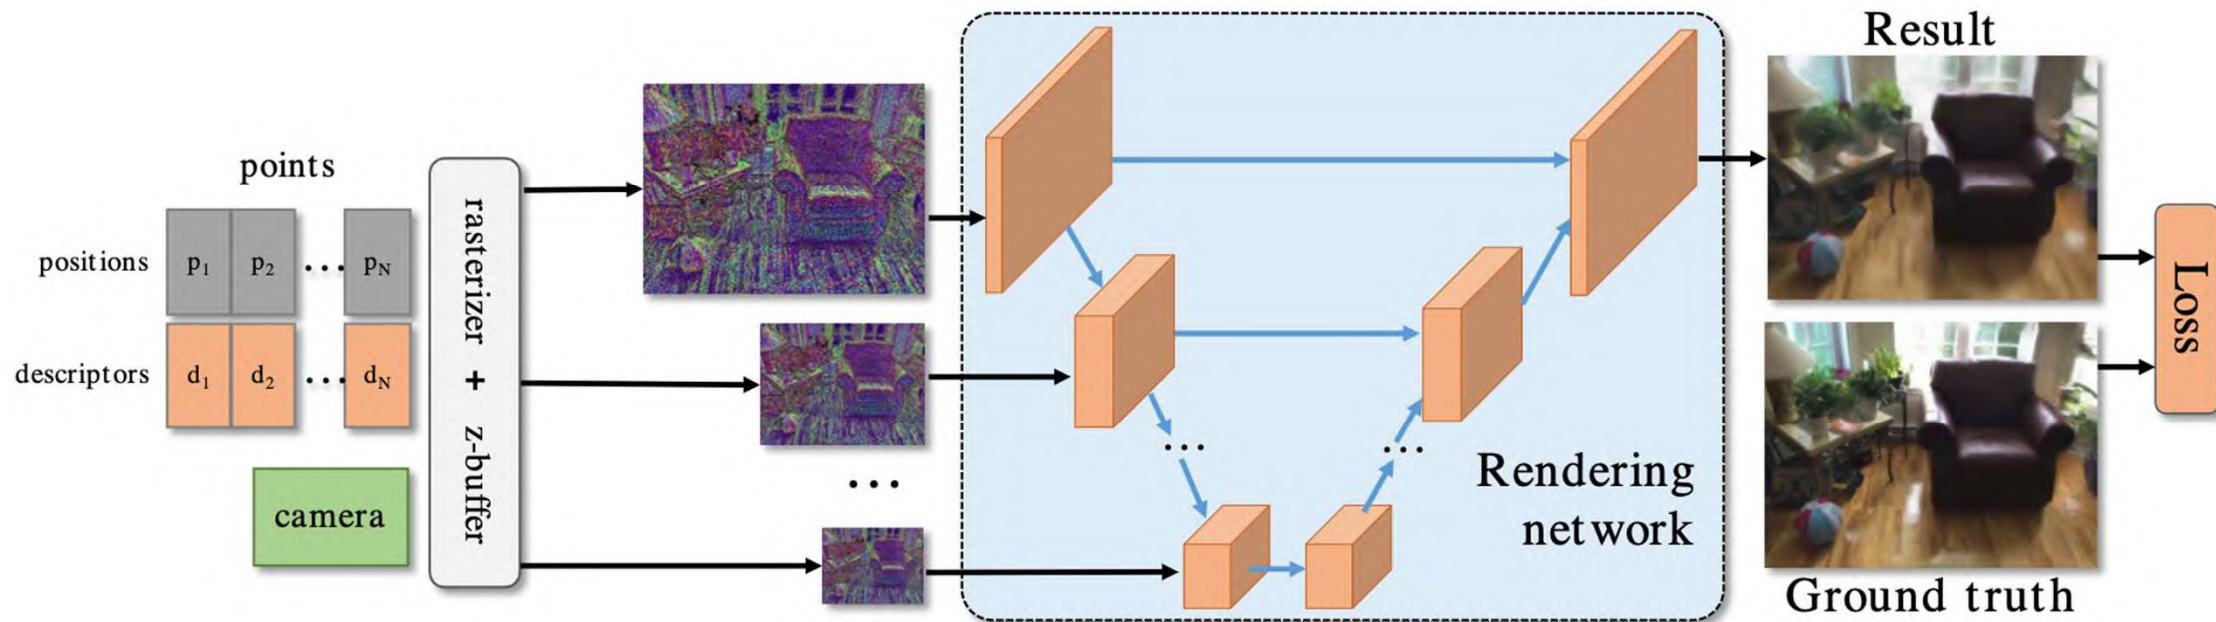
0:00

▶ ⏮ 🔊 0:00 / 0:59





# Neural Rendering Point Clouds



Setup: View-synthesis from available views  
Memorize one specific scene

Neural Point Based Graphics, Aliev et al. ECCV 2020

# Neural Rendering Point Clouds

Nearest  
available  
view



Pointcloud w/ color rendered

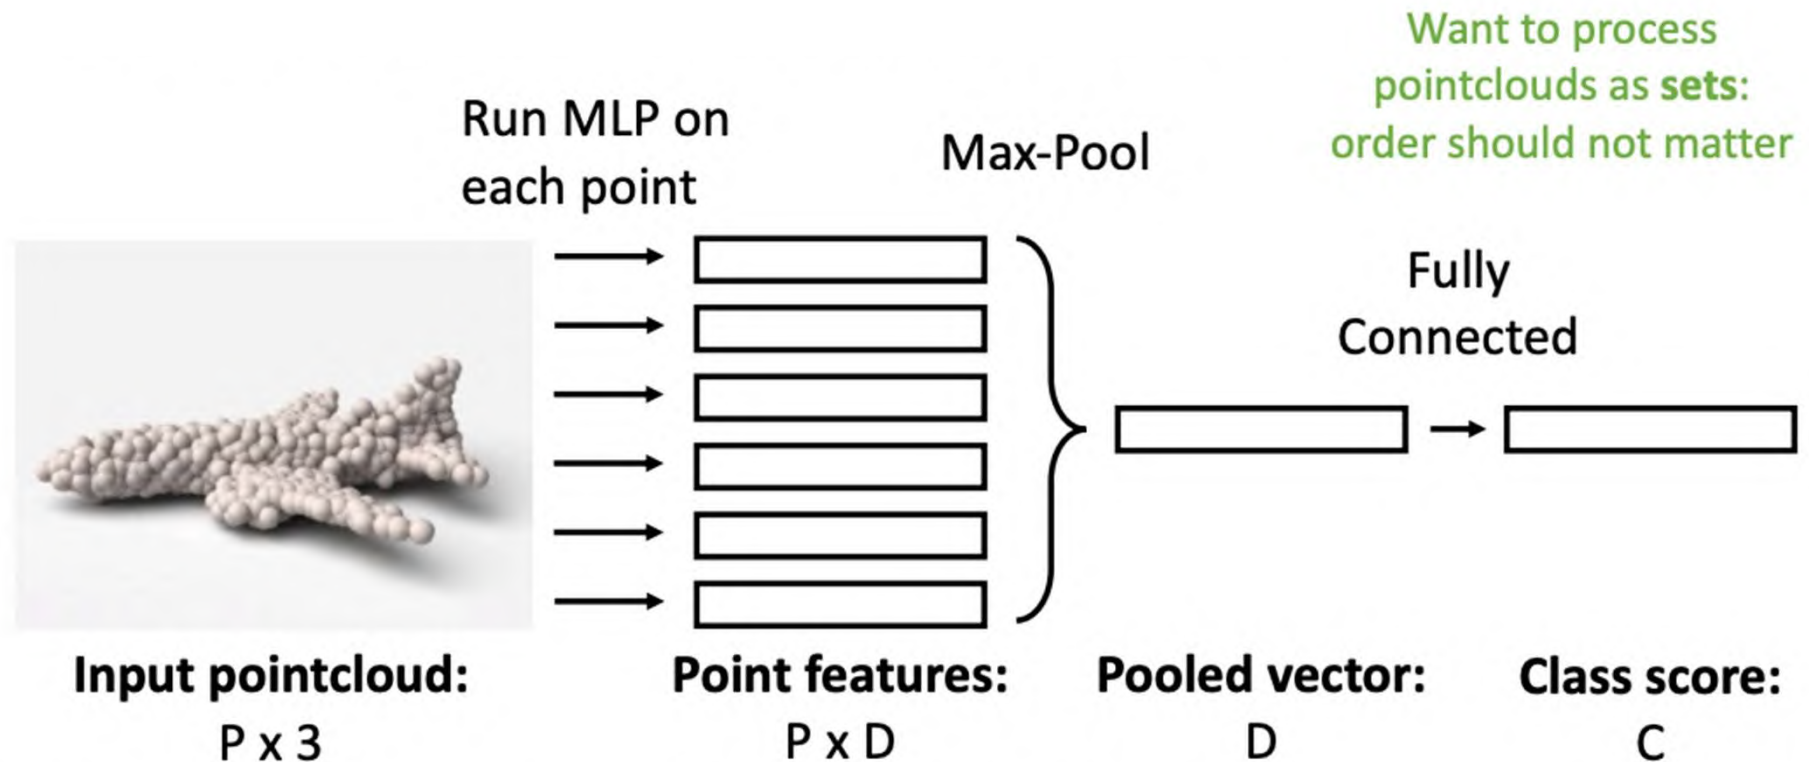


After neural rendering

Neural Point Based Graphics, Aliev et al. ECCV 2020



# Neural Processing Point Clouds



Qi et al, "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", CVPR 2017

Qi et al, "PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space", NeurIPS 2017

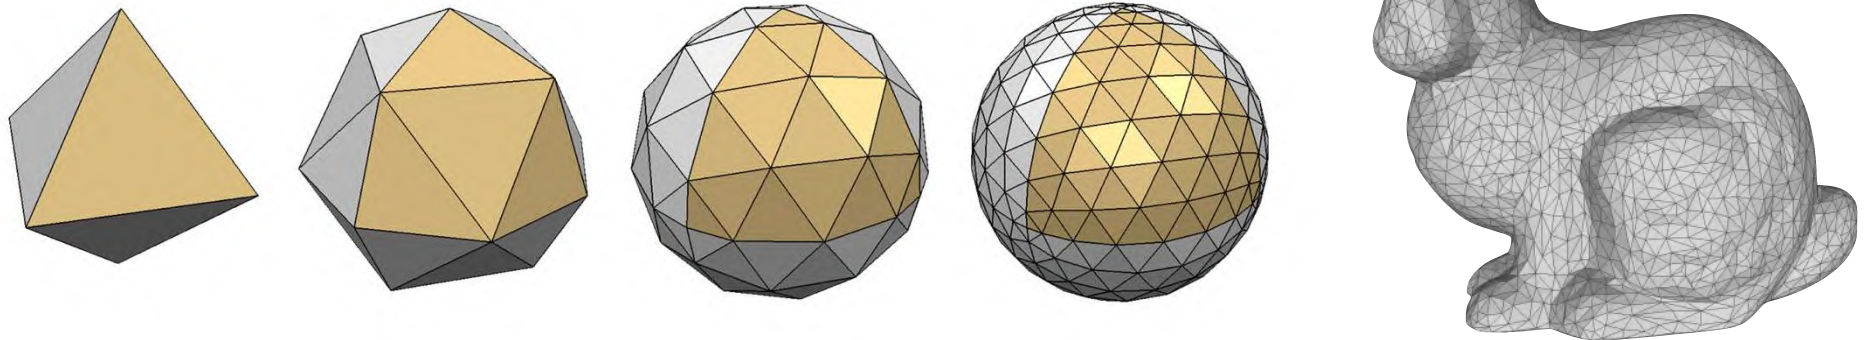
# Meshes: Connected Point Clouds

- Point clouds are order-less set of points
- Permutation invariant
- Meshes vertices are ordered and connected



# Polygon Meshes

- A mesh is a set of vertices with faces that defines the topology
- Mesh = {Vertices, Faces}
  - Vertices:  $N \times 3$
  - Faces:  $F \times \{3, 4, \dots\}$  specifying the edges of a polygon
  - Triangle faces most common but tetrahedrons (tets) are also.
- Surface is explicitly modeled by the faces
- Most common modeling representation



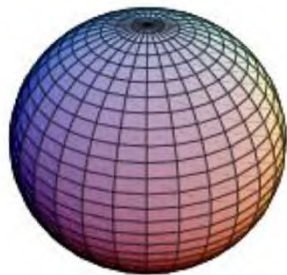
# Topology: Genus

Makes it hard to *directly* predict meshes of arbitrary objects from images

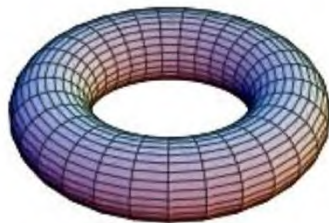
Informally, the number of **holes** or **handles**

Meshes can represent arbitrary topology.

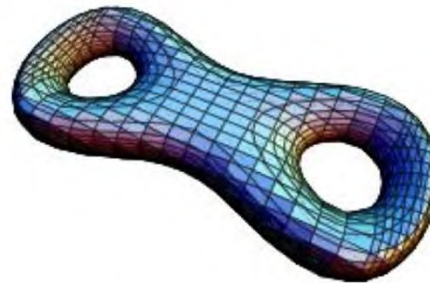
BUT two surfaces with different genres are not homeomorphic (can't be transformed without cutting / gluing)



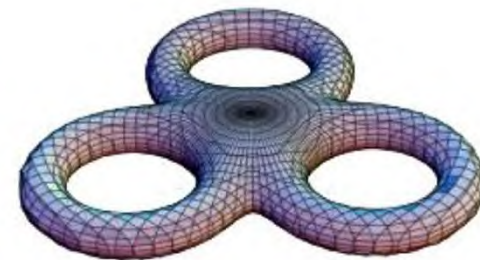
Genus 0



Genus 1



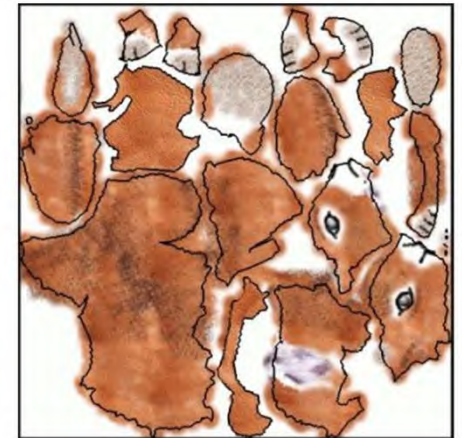
Genus 2



Genus 3

# Meshes are great for texturing

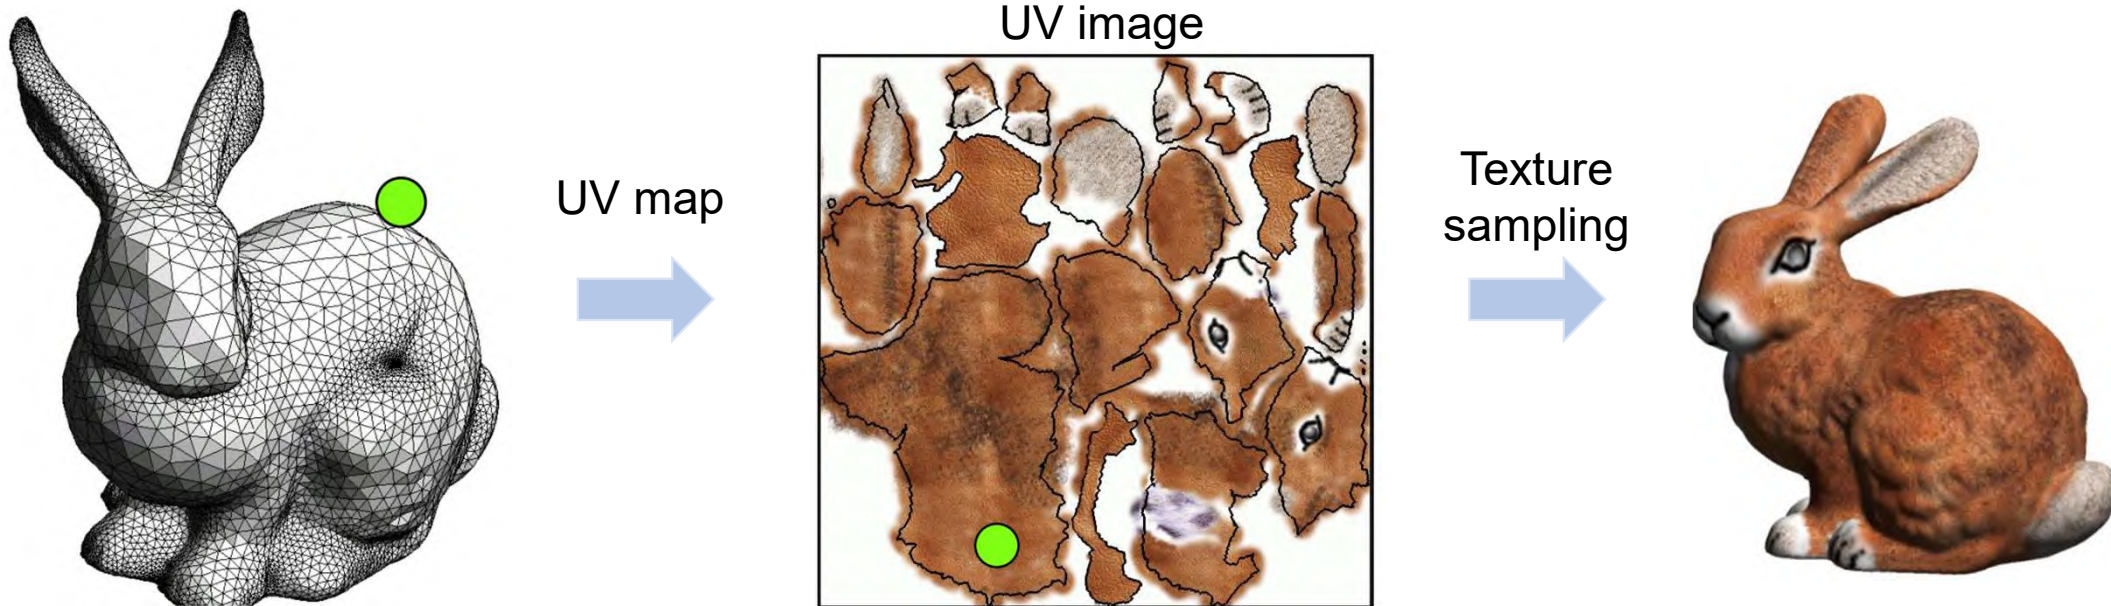
UV Image





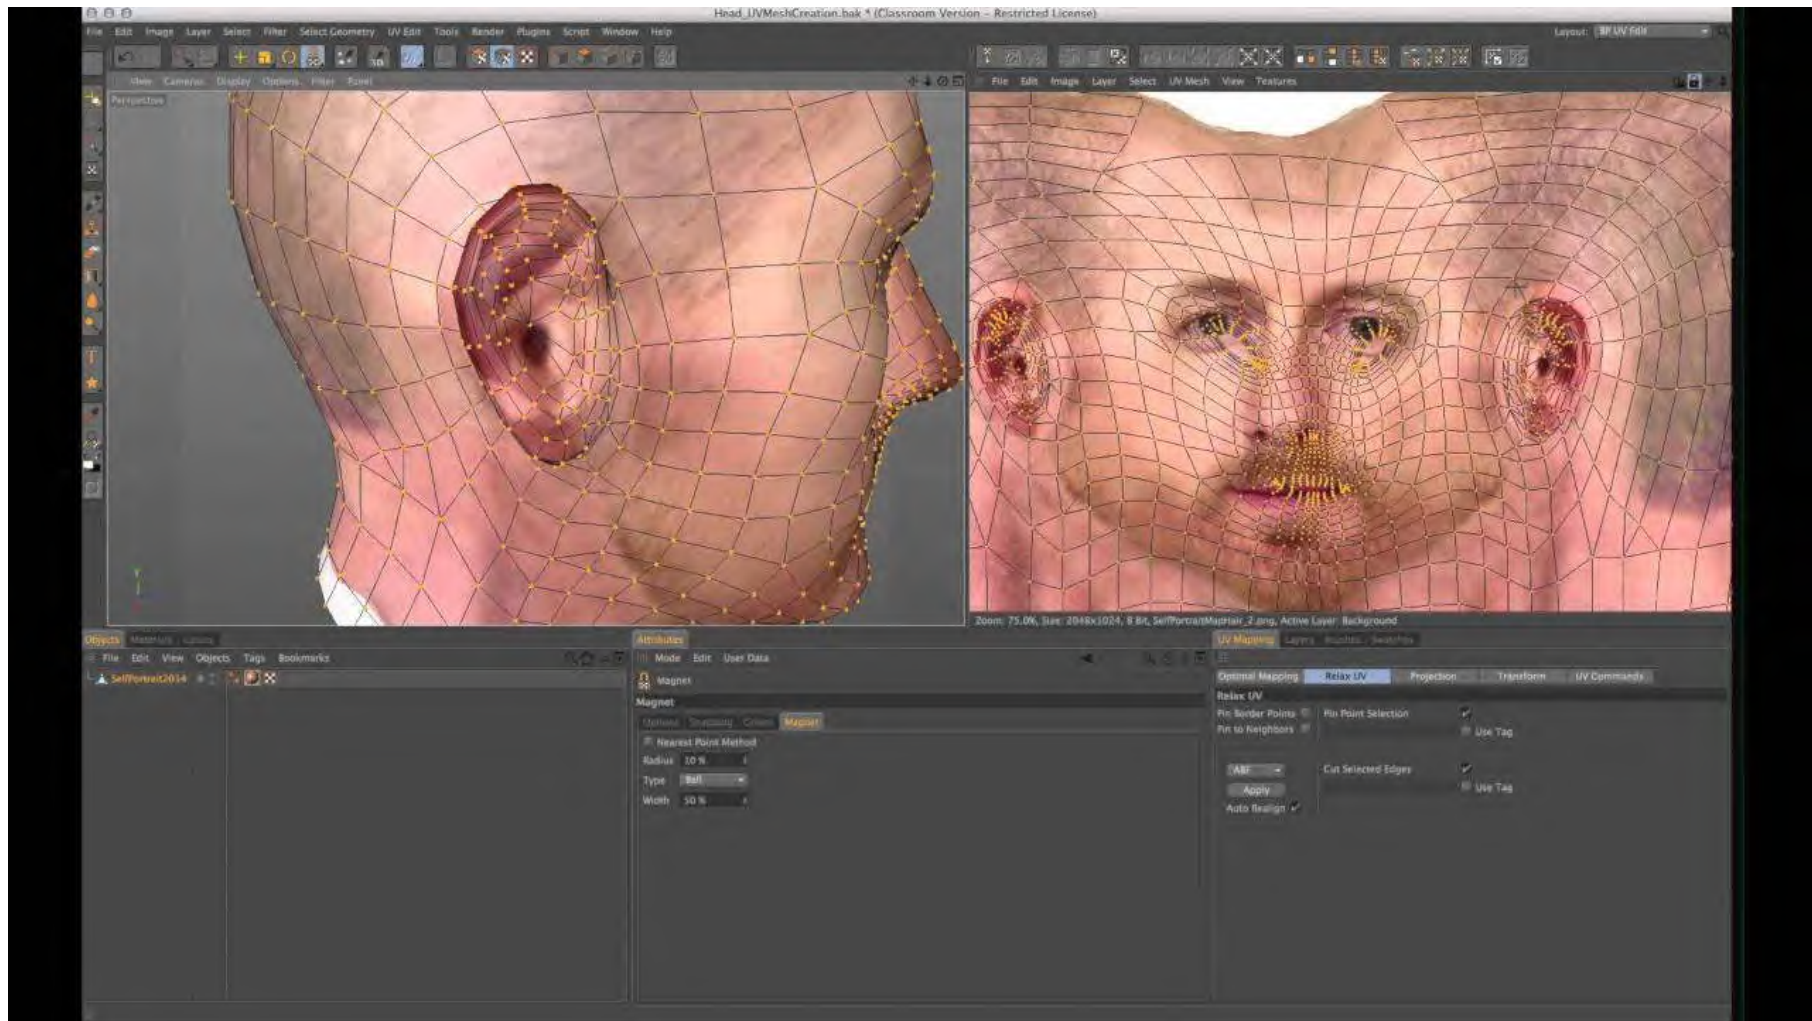
# Every single vertex has a UV coordinate

- Defined by UV mapping :  $(x,y,z) \rightarrow (u,v)$
- “texture coordinates”

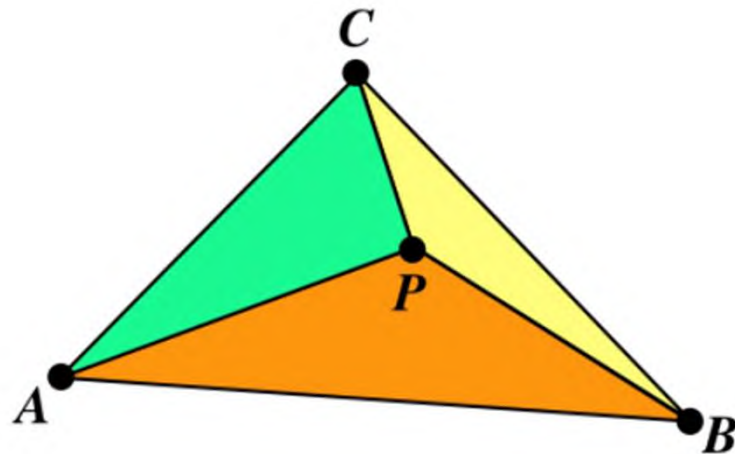




# Benefit of UV texturing: Continuous Color



# Barycentric coordinates to get UV Coordinates



$$P = w_A \times A + w_B \times B + w_C \times C$$

$$w_A = \frac{\Delta PBC}{\Delta ABC} = \frac{\text{yellow triangle}}{\text{whole triangle}}$$

$$w_B = \frac{\Delta PCA}{\Delta ABC} = \frac{\text{green triangle}}{\text{whole triangle}}$$

$$w_C = \frac{\Delta PAB}{\Delta ABC} = \frac{\text{orange triangle}}{\text{whole triangle}}$$

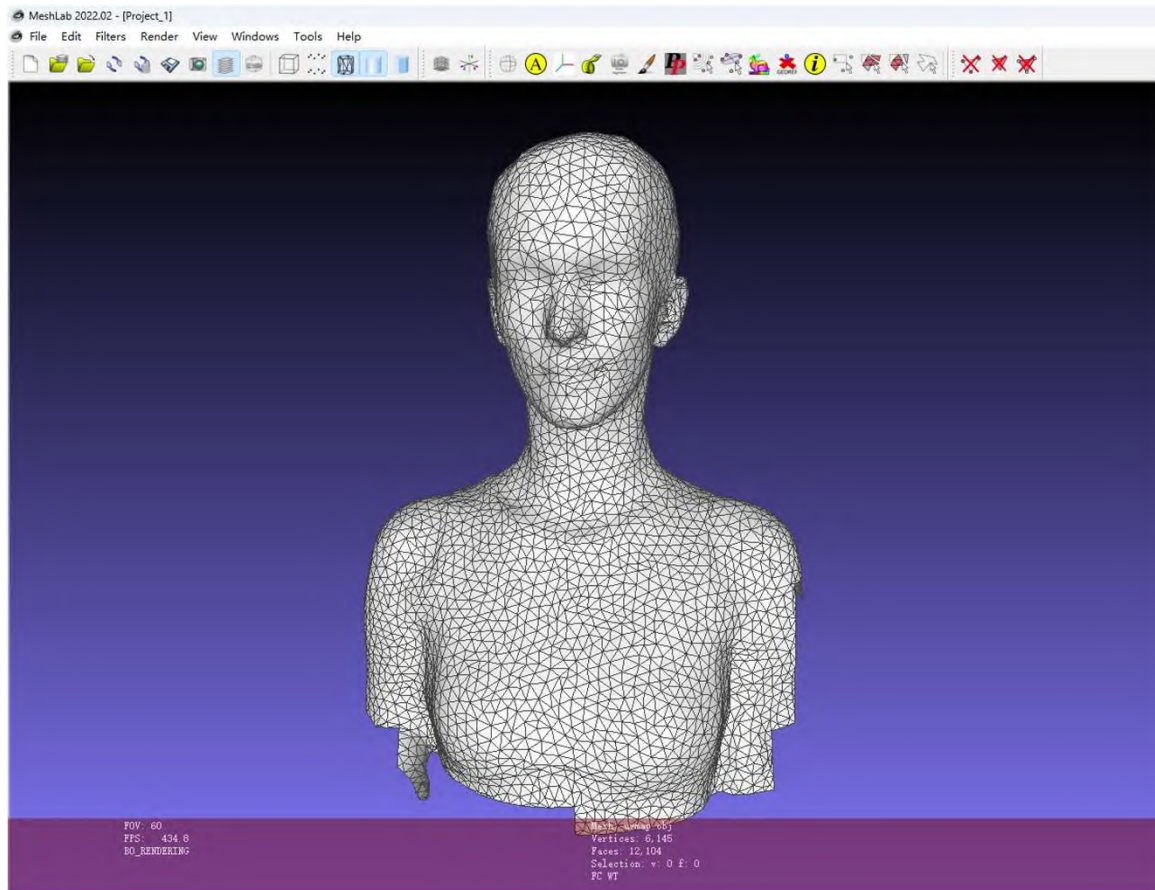
inside condition

$$0 \leq w_A, w_B, w_C \leq 1 \quad w_A + w_B + w_C = 1$$

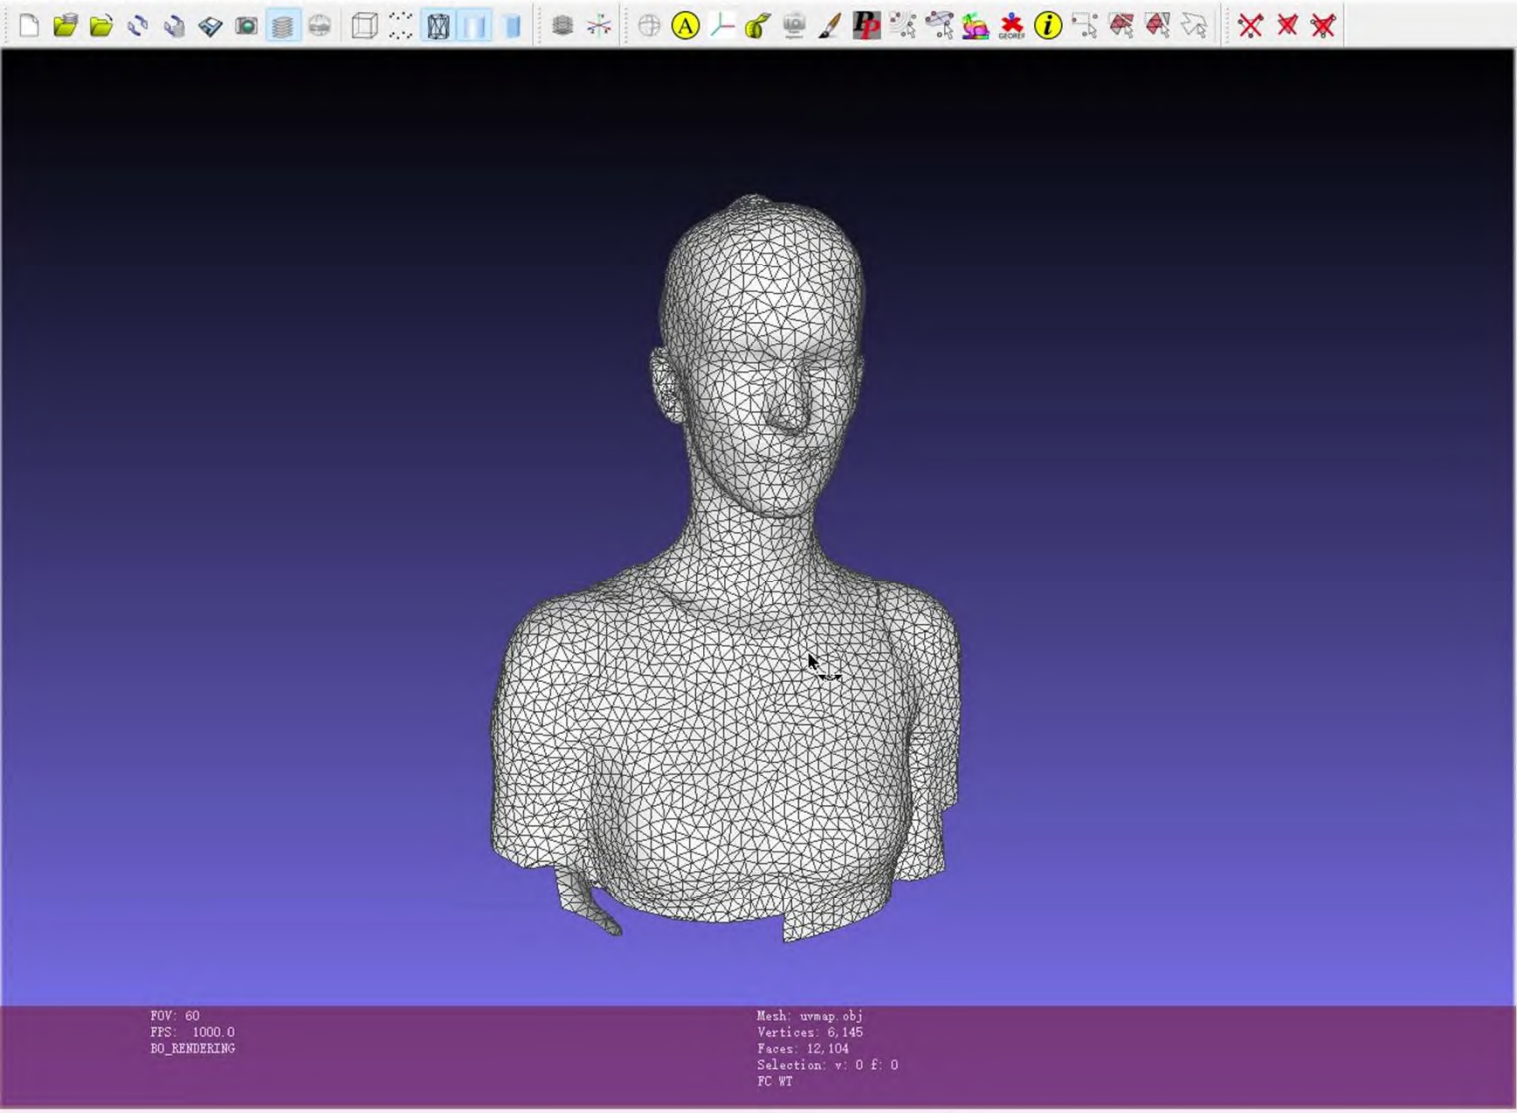
# Texturing Process

- Precompute the UV mapping
- Sample  $X$  many points on each triangle
- Figure out their UV coordinates (compute once)
- Get a UV image
- Sample the UV image

# Easy For Rendering







1 2 3 4 1 2 3 4

uvmap.obj

Shading **Vert** Face None

Color Face Mesh **User-Def**

Back-Face Single **Double** Fancy Cull

Texture Coord On **Off**

apply to all visible layers ☐

Opened mesh C:  
Users\Yudong\Downloads\textured\_head\uvmap.obj  
in 492 msec  
All files opened in 494 msec

FOV: 60  
FPS: 1000.0  
BO\_RENDERING

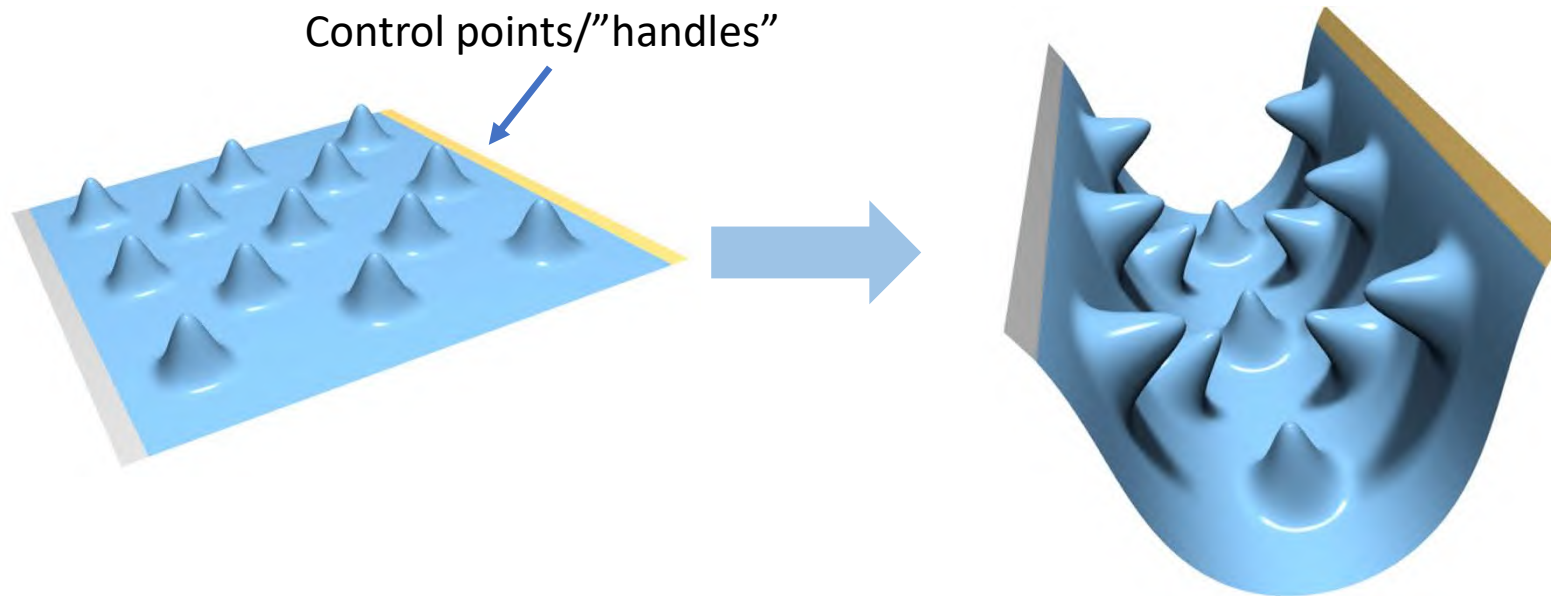
Mesh: uvmap.obj  
Vertices: 6,145  
Faces: 12,104  
Selection: v: 0 f: 0  
FC WT

ee: 8800MB vbo - 8800MB t



# Good For Deformation

- For: shape/character editing, sculpting, modelling
- Problem: Deform some 3D representation given some target



# Good For Animation



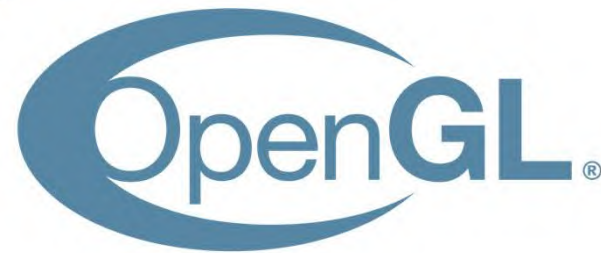


Very Easy For Using

GUI Software

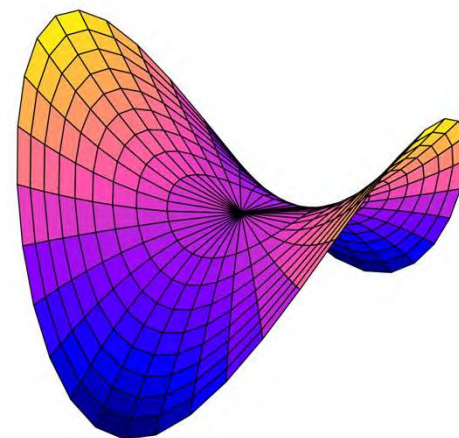
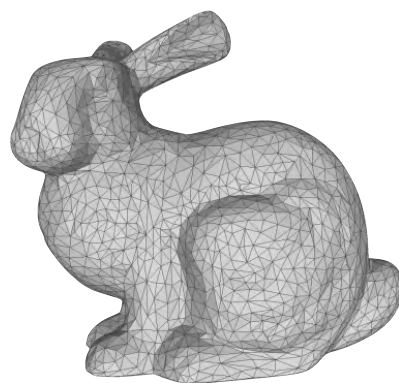
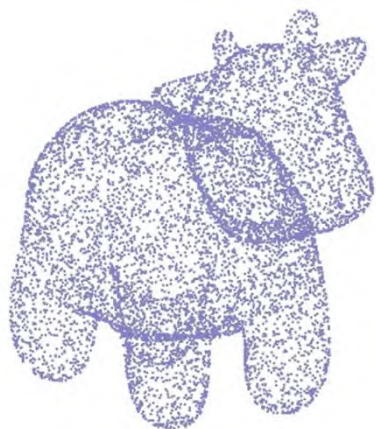


**MeshLab**

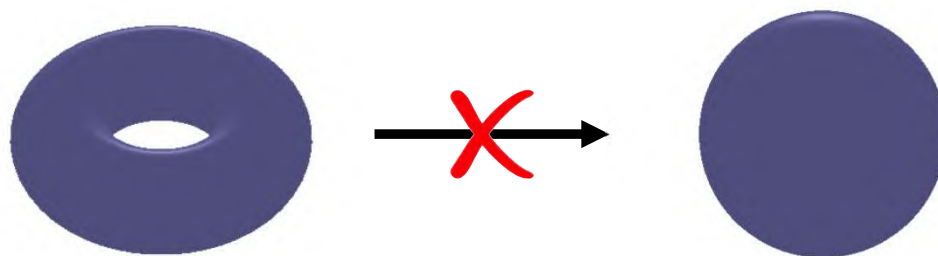


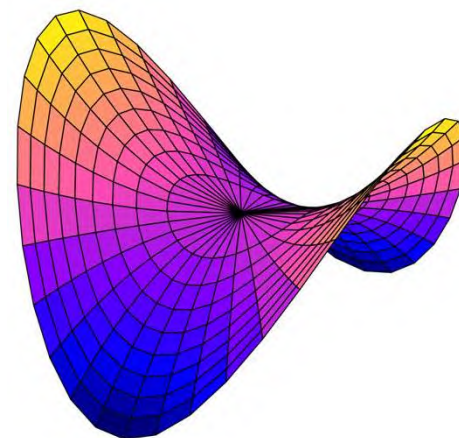
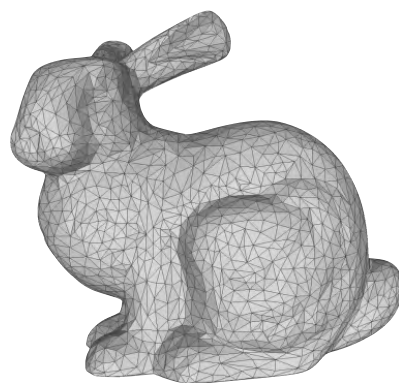
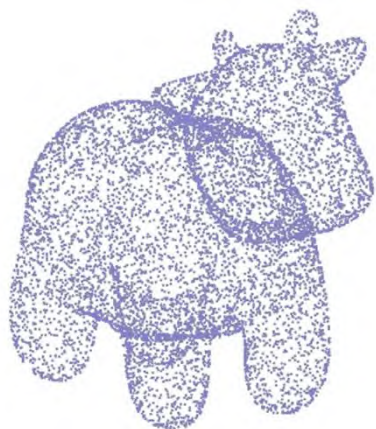
Coding Library



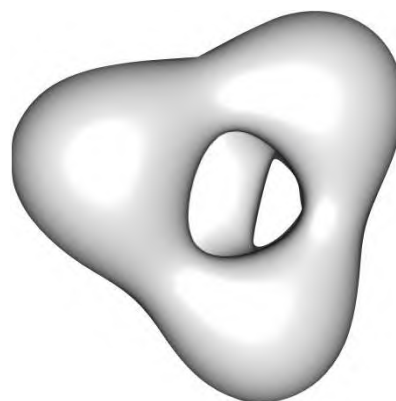


## Surface Representations





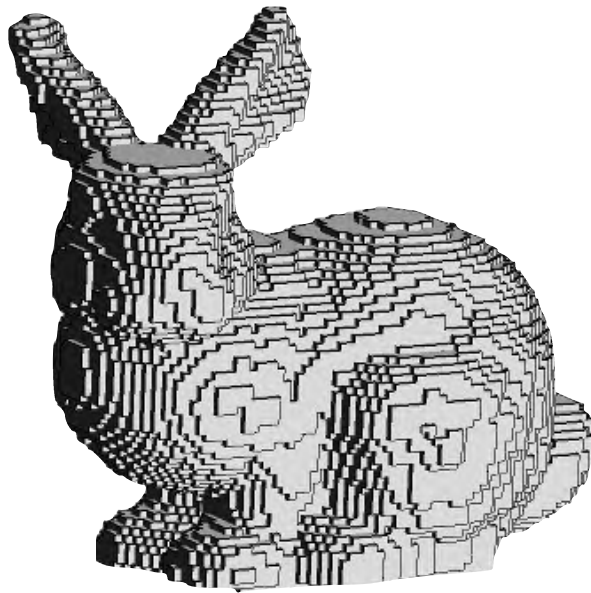
**Surface** Representations



**Volume** Representations

# Volumetric Representations

So far we talked about points, lines, and surfaces  
Volumetric representations model the entire space  
Can be explicit & implicit



# Voxels

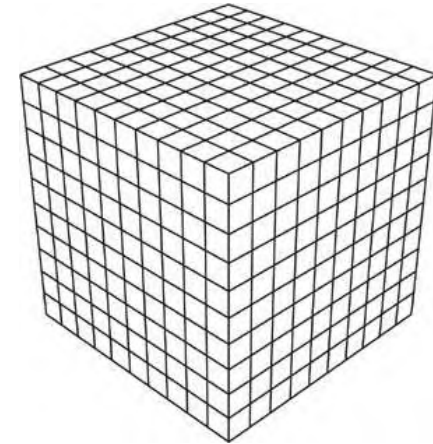
Discretize a 3-D space with some resolution:  $D \times D \times D$

What is stored at each value can be:

- 1 or 0 (occupancy)
- Signed distance
- Color/attributes

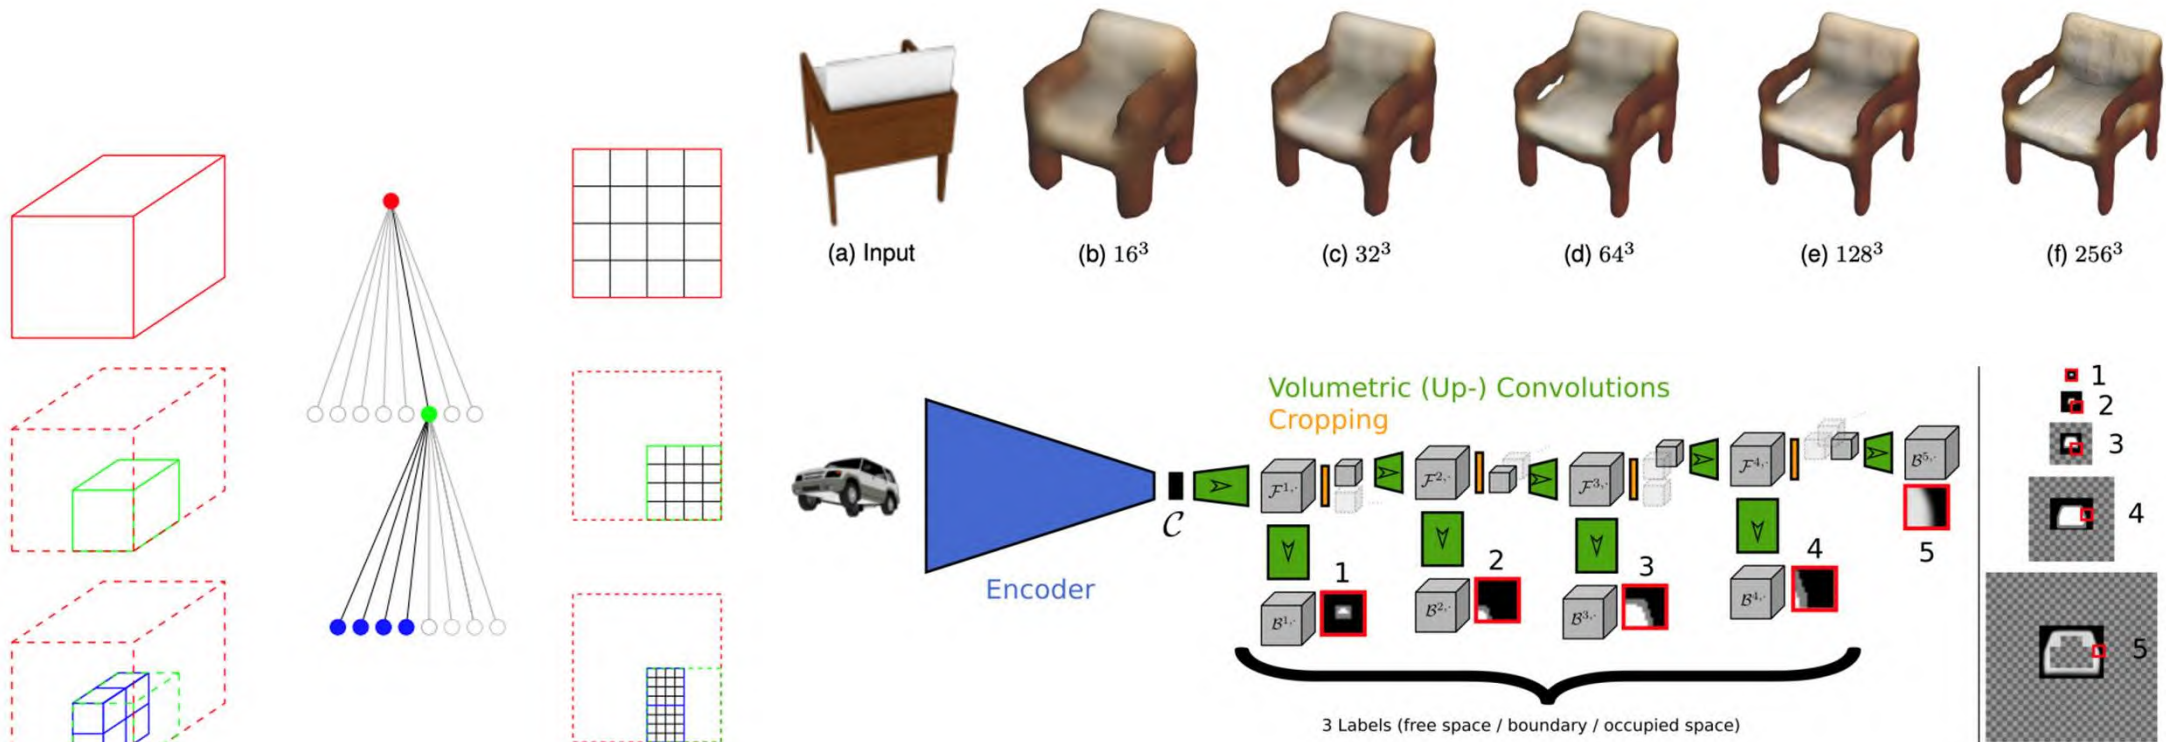
Was the easiest to get adopted into Deep Learning because of its proximity to pixels.

Simple, but very memory expensive!!





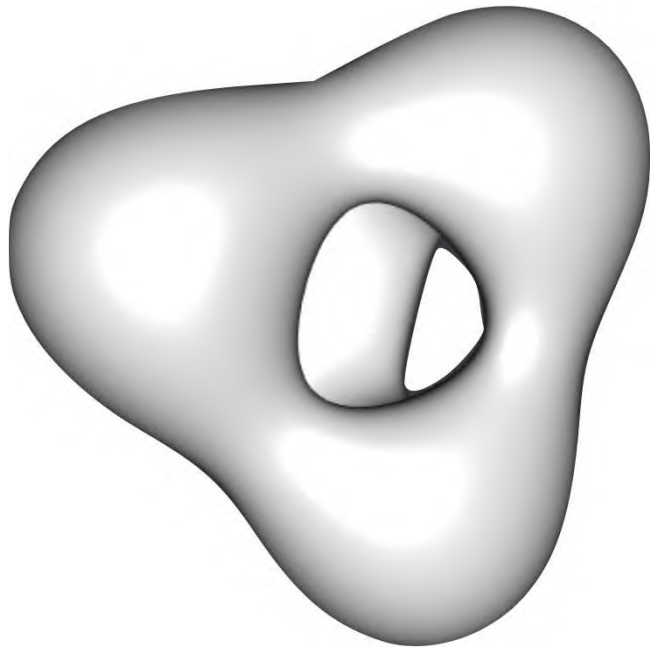
# Oct-tree (八叉树)



OctNet: Learning Deep 3D Representations at High Resolutions, Riegler et al. CVPR2017  
Hierarchical Surface Prediction for 3D Object Reconstruction Häne, 3DV 2017

# Volume for Surfaces

- You can use volumes to represent Surfaces
- the zero-crossing of a continuous function is the surface



$$\{\mathbf{p} \mid f(\mathbf{p}) = 0\}$$

```
def f(p):  
    ...  
    Input:  
        p: a point in R^3  
    Output:  
        v: A scalar value  
    ...
```

Can be anything — an analytic function, or a neural network, or a voxel

# Implicit Volume

Instead of explicitly outputting a discretized volume  $D^3$ , learn a function:

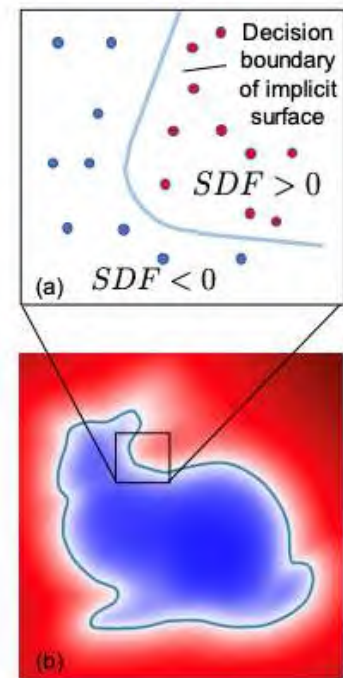
$$f_{\theta} : \mathbb{R}^3 \times \mathcal{X} \rightarrow [0, 1]$$

Output can be an occupancy  $\{0, 1\}$  or a real valued signed distance.

$X$  is some N-dim image embedding

This  $f$  is often just a simple 3~8 layer MLP.

$$f(x, y, z, \text{observation}) = \mathbb{R}$$

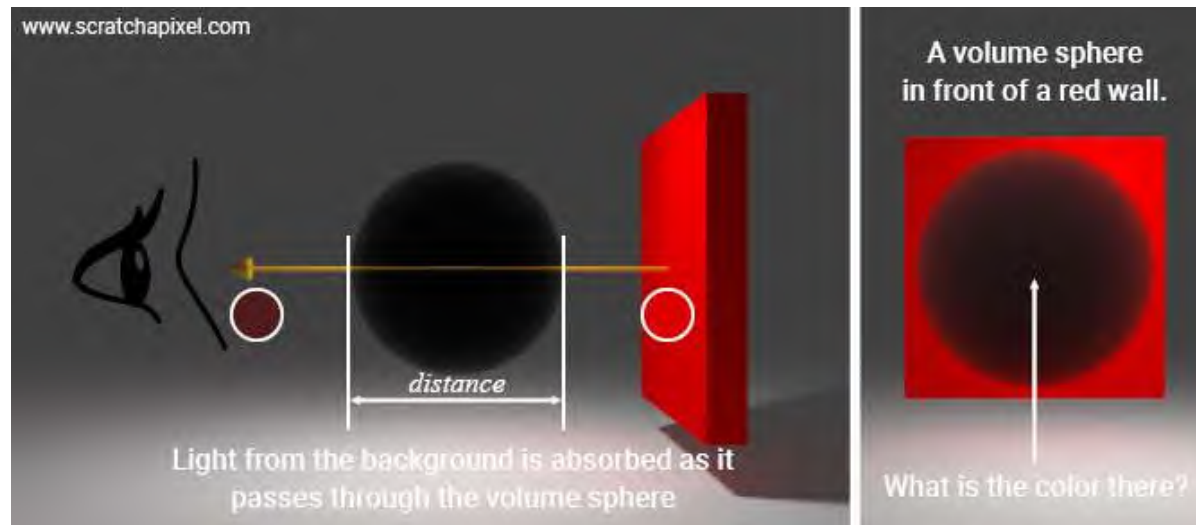


Marching  
Cubes



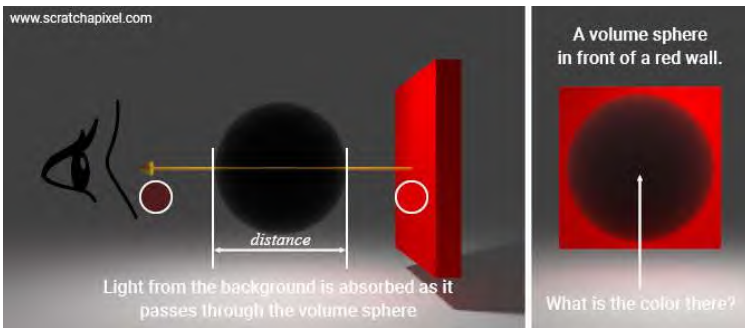
DeepSDF, CVPR'19,  
OccupancyNet, CVPR'19,  
Learning Implicit Fields for Generative Shape Modeling, CVPR'19

# Rendering Volumes





# Rendering Volumes



Given color and density  $(r, g, b, \sigma)$ , we calculate the color of every camera ray using:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t)) dt$$

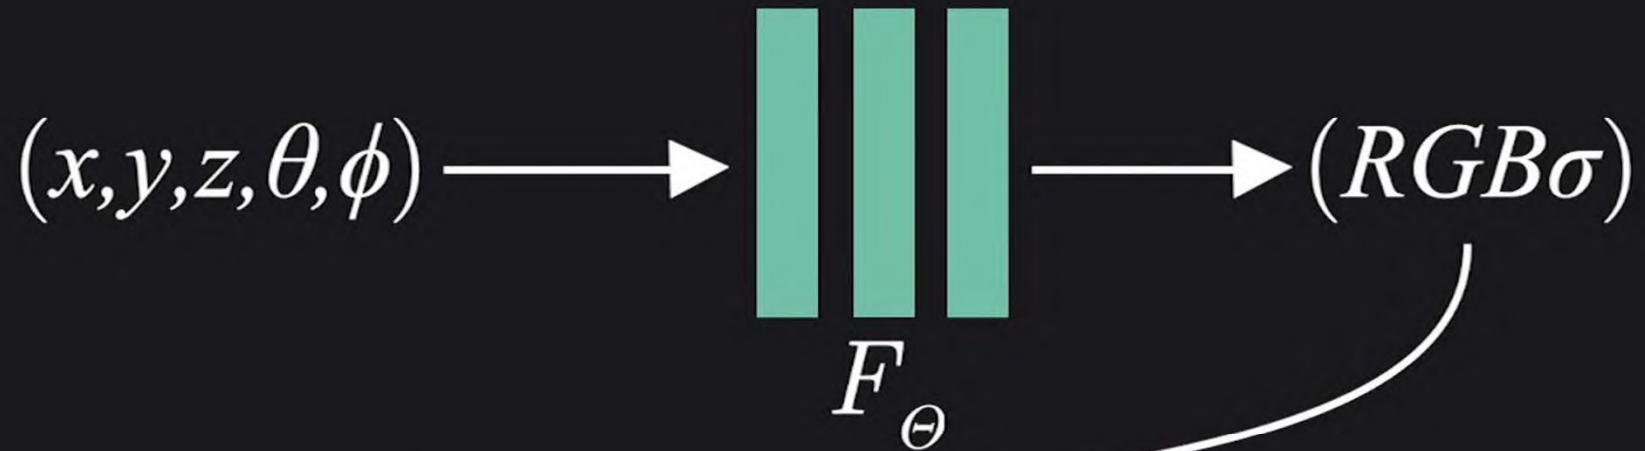
Camera ray:  $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Near and far bounds

Color (r,g,b) at  $\mathbf{r}(t)$

Volume density: Probability of a ray terminating at an infinitesimal particle at location  $\mathbf{r}(t)$

Accumulated transmittance along ray :  
 $T(t) = \exp(-\int_{t_n}^{t_f} \sigma(\mathbf{r}(s)) ds)$



NeRF: Combining VolRendering with DeepLearning

# Advantage: Very Realistic Rendering

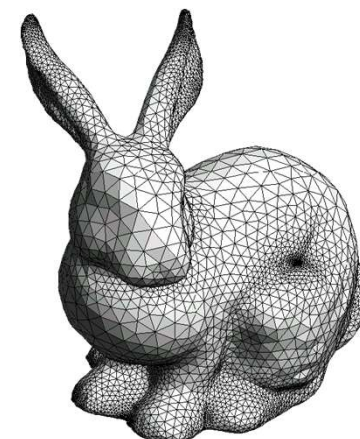
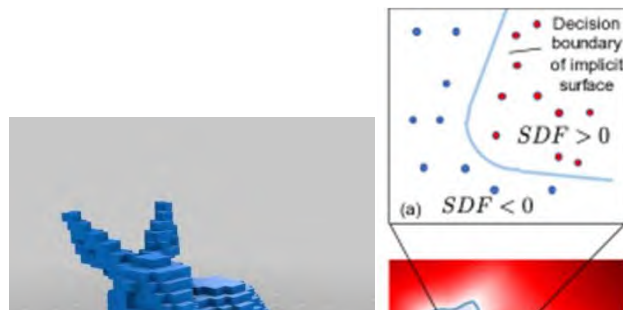


Advantage: Easy for (Topology) Complex Editing





# Summary



合适的才是最好的

Volumetrics

2.5D / Image Based Rendering

Explicit

Implicit

Point clouds

Meshes

- 👍 Can make pretty images
- 👍 Capture higher order lighting effects
- 👍 Robust off-the-shelf models (MiDas, VIT..)
- ✗ often not suitable for large baseline/360 view

- 👍 Easy to train with VolRend
- 👍 Topology free (editing)
- 👍 PhotoRealistic Rendering
- ✗ No surface\*
- ✗ Hard to texture
- ✗ Exp: Memory intensive\*
- ✗ Imp: Expensive to render

- 👍 Better memory
- 👍 Topology free
- ✗ No surface
- ✗ Can't print it
- 🙄 Need to splat / holes
- 🙄 Easy to get stuck in diffrend

- 👍 Better memory
- 👍 Explicit surface
- 👍 Great for texture/light
- 👍 Great for deformation
- 👍 Most common in 3D artists/graphics
- ✗ Topology sensitive
- 🙄 Easy to get stuck in diffrend



中国科学技术大学

University of Science and Technology of China

谢谢观看!